

Getting good performance from your application

Tuning techniques for serial programs on
cache-based computer systems

Overview

- ❑ Introduction
- ❑ Memory Hierarchy
- ❑ General Optimization Techniques
- ❑
- ❑ Compilers
- ❑ Analysis Tools
- ❑ Tuning Guide

Introduction

Introduction

Moore's Law



- ❑ Popular version:
 - ❑ “*CPU speed usually doubles every 18 months.*”
- ❑ More correct version:
 - ❑ “*The number of transistors per integrated circuit will double every 18 months.*”

Introduction


Gordon Moore – co-founder of Intel

"I never said 18 months. I said one year, and then two years ... Moore's Law has been the name given to everything that changes exponentially. ... If Gore invented the Internet, I invented the exponential."


- Gordon Moore in an interview (2000)




Introduction



Introduction



Introduction




Introduction


- ❑ CPU speed usually doubles every 18-24 months (not true any longer – we got other improvements, instead!).
- ❑ Development on the memory side is much slower (~ 4-6 years!).
- ❑ Memory speeds catch up – but also have to serve more cores!
- ❑ New memory technologies do help (e.g. HBM)
- ❑ Something you should have in mind when designing your program!

Motivation for Application Tuning

time flow in a computational task (simplified picture):



new hardware:



ideal: 2x faster computer



reality: 2x faster CPU


code tuning (old hardware):




The Memory Hierarchy

The Memory Hierarchy

Intuitive Performance Graph:



The Memory Hierarchy




*Memory Optimization:
Keep frequently used data close to the processor*


The Memory Hierarchy

Performance is not uniform:


Application Tuning



The Memory Hierarchy



The Memory Hierarchy



The Memory Hierarchy

- ❑ Memory plays a crucial role in performance
- ❑ Not accessing memory in the right way will degrade performance on **all** computer systems
- ❑ The extent of degradation will depend on the system
- ❑ Knowledge about the relevant memory characteristics helps to write code that minimizes those problems

Caches – and all that ...


How do those caches work?

Caches

- ❑ Cache memory or cache for short (from French: cacher – to hide): fast buffers that help to hide the memory latency
- ❑ One distinguishes between
 - ❑ data cache
 - ❑ instruction cache
 - ❑ address cache (also called TLB – Translation Lookaside Buffer) – mapping between virtual and physical addresses

Cache Lines

- ❑ To get good performance, optimal use of the caches is crucial
- ❑ The unit of transfer is a “*cache line*”:
 - ❑ linear structure of fixed length (bytes)
 - ❑ fixed starting address in memory



Cache Organisation

Direct Mapped:

- ❑ Each memory address maps onto exactly one line in cache
- ❑ simple and efficient
- ❑ built-in replacement policy
- ❑ easy to scale to larger sizes
- ❑ downside: no control by usage – danger of replacing data that will be needed again soon

Cache Organisation


Fully Associative:

- ❑ Every memory address can be mapped anywhere in cache
- ❑ Need to track usage of cache lines
- ❑ Requires a replacement policy, e.g.
 - ❑ *least recent used* (LRU),
 - ❑ *least frequent used* (LFU),
 - ❑ random, etc
- ❑ Doesn't scale well to large sizes
- ❑ Costly design

Cache Organisation


N-way Set Associative:

- ❑ Sets of direct mapped caches:



Memory access

- ❑ Memory has a 1-dimensional linear structure
- ❑ Accessing vector elements:




size of a
cache line

Legend:

	vector element
	cache miss
	cache hit
↓	memory access


Memory access

Accessing 2d arrays in C – row wise:




Memory access

Accessing 2d arrays in C – column wise:



Memory access

Access to multi-dimensional arrays depends on how data is stored:





Bad memory access has a huge impact on performance!!!

The TLB cache

- ❑ the Translation Lookaside Buffer (TLB) translates virtual memory addresses (in your application) to physical addresses
- ❑ also called 'address cache'
- ❑ unit: page – typical size 4kB
- ❑ creation of lookup table is an expensive operation
- ❑ cost: 10 – 100 clock cycles/miss
- ❑ modern CPUs are having advanced TLBs
 - ❑ support for variable page sizes

Memory access – TLB misses



- = TLB miss
- = D-cache miss
- = Cached elements
- = Virtual memory page

- If the entire matrix fits in the cache, the access pattern *hardly* matters.
- For large (out-of-cache) matrices, the access pattern **does** matter – both data cache and TLB misses

About cache misses

Some simple rules:

- ❑ You cannot avoid cache misses – they are part of the nature of cache-based systems ...
- ❑ ... but you should try to minimize them to get good performance

Cache Line Utilization

Two key rules: **Maximize ...**


- ❑ **Spatial locality** ⇒ Use all data in one cache line
 - ❑ depends on storage layout
 - ❑ depends on access patterns
 - ❑ stride = 1 is good
 - ❑ random access is really bad
- ❑ **Temporal locality** ⇒ Re-use data in a cache line
 - ❑ depends on algorithm used

Some Terminology

Terminology: Pipelining

stage 1
stage 2
stage 3
stage 4

Let's assume we have an operation that takes 4 stages per iteration



Rule of thumb: keep the pipeline filled for good performance!

Terminology: Superscalar (or ILP)

- ❑ *N-way superscalar:*
 - *Execute N instructions at the same time*
- ❑ *This is also called Instruction Level Parallelism (ILP)*

	slot 1	slot 2	slot 3	slot 4	
cycle 1					4-way superscalar
cycle 2	not used				3-way superscalar
cycle 3		not used	not used		2-way superscalar
cycle 4	not used			not used	2-way superscalar
cycle 5			not used		3-way superscalar

- ❑ *The hardware has to support this, but it is up to the software to take advantage of it*
- ❑ *Often there are restrictions which instructions can be "bundled"*
- ❑ *These are documented in the Architecture Reference Manual for the microprocessor*

Terminology: FMA

One variant of ILP for floating point operations:

- ❑ FMA – Fused Multiply Add
- ❑ special instruction, that calculates
$$a + (b * c)$$
in one operation, with a single rounding
- ❑ faster
- ❑ ... but the result might slightly differ, due to the single rounding!


Terminology: Vectorization (SIMD)

- ❑ Vectorization is the process of transforming a scalar operation that acts on a single data element at a time (**Single Instruction Single Data – SISD**) to an operation that acts on multiple data elements at a time (**Single Instruction Multiple Data – SIMD**).
- ❑ **SIMD** units are hardware arithmetic vector units that can perform the same operation on multiple data points simultaneously by using vector registers.

Terminology: Vectorization (SIMD)

❑ Vector types

- 128 bit: SSE = Streaming SIMD Extension (1999)



2 × double

4 × float


- 256 bit: AVX = Advanced Vector Extension (2011)



4 × double

8 × float

- 512 bit: AVX-512 (2013)



8 × double

16 × float

Latency and Bandwidth

Latency:

- ❑ the time it takes from the initiation of an action till you have the first result
- ❑ unit: time

Bandwidth:

- ❑ how many
 - ❑ actions can be carried out,
 - ❑ results can be obtained,
within a given time
- ❑ unit: #/time

General Optimization Techniques

Optimization Techniques - Overview

- ❑ Most optimization techniques are “loop based”
- ❑ Loop based optimizations:
 - ❑ Interchange
 - ❑ Fission and Fusion
 - ❑ Unrolling
 - ❑ Blocking

Optimization Techniques - Overview

- ❑ Designing your data structures the “right way” can also be important
- ❑ Other techniques:
 - ❑ De-vectorization
 - ❑ Stripmining

Loop based optimizations

Coding style: array indexing

- ❑ To apply safe transformations, the compilers have to analyze data dependencies in a loop
- ❑ Explicit expressions will help the compilers to do a good job in loop optimization

Good

```
for(i=0; i<m; i++)
    for (j=0; j<n; j++)
        .. a[i][j] ..
```

()Reasonable*

```
for(i=0; i<m; i++)
    for (j=0; j<n; j++)
        .. a[i*n+j] ..
```

Bad

```
k = 0;
for(i=0; i<m; i++)
    for (j=0; j<n; j++)
        .. a[k++] ..
```

Caution

```
for(i=0; i<m; i++)
    for (j=0; j<n; j++)
        .. a[ndx[i][j]] ..
```

(*) harder to read for a human, but might be better for the compiler!

Loop Interchange

```
DO I = 1, M
    DO J = 1, N
        A(I,J)=B(I,J)+C(I,J)
    END DO
END DO
```

 *Interchange loops*

```
DO J = 1, N
    DO I = I, M
        A(I,J)=B(I,J)+C(I,J)
    END DO
END DO
```

- ❑ The matrices are accessed over the second dimension first
- ❑ This is the wrong order in Fortran
- ❑ A loop interchange solves the problem
- ❑ In C, the situation is reversed:
 - ❑ row access is okay
 - ❑ column access is bad

Loop Fission

```
for (j=0; j<n; j++)  
{  
    c[j] = exp(j/n);  
    for (i=0; i<m; i++)  
        a[i][j]=b[i][j]+d[i]*e[j];  
}
```

- ◆ Access on arrays 'a' and 'b' is bad
- ◆ We can not simply interchange the loops
- ◆ Fission/splitting is the solution

Fission

This loop can now also be vectorized

Interchange loops for better performance

```
for (j=0; j<n; j++)  
    c[j] = exp(j/n); New loop created
```

```
for (j=0; j<n; j++)  
    for (i=0; i<m; i++)  
        a[i][j]=b[i][j]+d[i]*e[j];
```

Loop Fusion

```
for (i=0; i<n; i++)
    a[i] = 2 * b[i];
```

```
for (i=0; i<n; i++)
    c[i] = a[i] + d[i];
```

Fusion


- ◆ Assume that 'n' is large
- ◆ In the second loop, a[i] will no longer be in the cache
- ◆ Fusing the loops will ensure a[i] is still in the cache when needed

Note that it is possible to apply fusion to loops with (slightly) different boundaries

In such a case, some iterations will have to be 'peeled' off

```
for (i=0; i<n; i++)
{
    a[i] = 2 * b[i];
    c[i] = a[i] + d[i];
}
```

Fission and Fusion – Summary



Fission

- ✓ *Reduce register pressure*
- ✓ *Enable loop interchange*
- ✓ *Isolate dependencies*
- ✓ *Increase opportunities for optimization (e.g. vectorization of intrinsics)*

Fusion

- ✓ *Reduce cache reloads*
- ✓ *Increase Instruction Level Parallelism (ILP)*
- ✓ *Reduce loop overhead*

(Inner) Loop Unrolling

Through unrolling, the loop overhead ('book keeping') is reduced

```
for (i=0; i<n; i++)
    a[i] = b[i] + c[i];
```

*Loop is unrolled
with a factor of 4*

```
for (i=0; i<n; i+=4)
{
    a[i] = b[i] + c[i];
    a[i+1] = b[i+1] + c[i+1];
    a[i+2] = b[i+2] + c[i+2];
    a[i+3] = b[i+3] + c[i+3];
}
<clean-up loop>
```

Loads : 2
Stores : 1
FP Adds : 1
I=I+1
Test I < N ?
Branch
Addr. incr: 3


*Work: 4
Overhead: 6*

Loads : 8
Stores : 4
FP Adds : 4
I=I+4
Test I < N ?
Branch
Addr. incr: 3

*Work: 16
Overhead: 6*

Note: the amount of addressing needed in reality is less

Outer Loop Unrolling



```
for (i=0; i<m; i++)
  for(j=0; j<n; j++)
  {
    a[i] += b[i][j] * c[j];
  }
```

```
for (i=0; i<m; i+=4)
  for(j=0; j<n; j++)
  {
    a[i ] += b[i ][j] * c[j];
    a[i+1] += b[i+1][j] * c[j];
    a[i+2] += b[i+2][j] * c[j];
    a[i+3] += b[i+3][j] * c[j];
  }
<clean-up loop>
```

- ♦ **Advantage:**
 - $c[j]$ is re-used 3 more times (*temporal locality*)
- ♦ **Deeper unrolling, say 8, requires more fp registers (17 instead of 9), but improves re-use of $c[j]$**

Outer Loop Unrolling – how to

```
for (i=0; i<m; i++)
    for(j=0; j<n; j++)
        a[i] += b[i][j] * c[j];
```

Unroll and Jam

*Outer loop
unrolling*

```
for (i=0; i<m-m%4; i+=4)
{
    for(j=0; j<n; j++)
        a[i] += b[i][j] * c[j];
    for(j=0; j<n; j++)
        a[i+1] += b[i+1][j] * c[j];
    for(j=0; j<n; j++)
        a[i+2] += b[i+2][j] * c[j];
    for(j=0; j<n; j++)
        a[i+3] += b[i+3][j] * c[j];
}
for (i=m-m%4; i<m; i++)
    for(j=0; j<n; j++)
```

clean-up loop

```

{
    a[i] += b[i][j] * c[j];
    a[i+1] += b[i+1][j] * c[j];
    a[i+2] += b[i+2][j] * c[j];
    a[i+3] += b[i+3][j] * c[j];
}
for (i=m-m%4; i<m; i++)
    for(j=0; j<n; j++)
        a[i] += b[i][j] * c[j];
```

clean-up loop

*Jam the loops
together again*

Loop unrolling – structure

```
for (i=0; i<n; i++)  
{  
    ... [i] ...  
}
```

```
DO I = 1, N  
    ... (I) ...  
END DO
```

*Loop unroll factor
is "unroll"*

```
for(i=0;i<n-n%unroll;i+=unroll)  
{  
    ... [i] ...  
    ... [i+1] ...  
    ... [i+2] ...  
    ... [i+unroll-1] ...  
}
```

```
DO I = 1, N-mod(N,unroll), unroll  
    ... (I) ...  
    ... (I+1) ...  
    ... (I+2) ...  
    ... (I+unroll-1) ...  
END DO
```

Unrolled Loop

Cleanup Loop

```
for(i=n-n%unroll;i<n;i++)  
{  
    ... [i] ...  
}
```

```
DO I = N-mod(N,unroll)+1, N  
    ... (I) ...  
END DO
```

Loop Unrolling – Summary

- ❑ More than one iteration per loop pass
- ❑ Inner loop unrolling:
 - ❑ reduce loop overhead
 - ❑ better instruction scheduling
- ❑ Outer loop unrolling:
 - ❑ improve cache line usage (spatial locality)
 - ❑ re-use data (temporal locality)
- ❑ Disadvantages:
 - ❑ more registers needed, clean-up code required


Loop Unrolling – Compilers

- ❑ compilers do (usually) a good job in loop unrolling – leave it to the compiler
- ❑ some compiler need to be told to do loop unrolling – not part of standard optimization, always check if in doubt
- ❑ there are options to control the unroll depth
 - ❑ gcc: -funroll-loops --params max-unroll-times=n
 - ❑ clang: -funroll-loops --unroll-count=n
 - ❑ Intel: -unroll[n] (0: disable loop unrolling)
- ❑ or compiler specific pragmas (see docs)

Loop Blocking – 1

Transposing a matrix

```
for (j=0; j<n; j++)
    for (i=0; i<n; i++)
        b[j][i] = a[i][j];
```




- ◆ *Loop interchange will not help here:*
 - Role of 'a' and 'b' will only be interchanged
- ◆ *Change of programming language won't help either*
- ◆ *Unrolling the i-loop can be beneficial, but requires more registers and doesn't address TLB-misses*
- ◆ *Loop blocking achieves good memory performance, without the need for additional registers*


Loop Blocking – 2

Blocking and interchanging the I-loop

```
for(il=0; il<n; il+=nbi)
    for (j=0; j<n; j++)
        for (i2=0;i2<MIN(n-il,nbi);i2++)
            b[j][il+i2] = a[il+i2][j];
```



- ◆ Parameter '*nbi*' is the *blocking size*
- ◆ Should be chosen as large as possible
- ◆ Actual value depends on the cache to block for:
 - ✓ L1-cache
 - ✓ L2-cache
 - ✓ TLB
 - ✓



Loop Blocking – Summary

- ❑ Powerful technique to improve:
 - ❑ memory access (spatial locality)
 - ❑ data re-use (temporal locality)
- ❑ Preserves portability – but blocking size depends on:
 - ❑ cache type/level/capacity (hardware)
 - ❑ data requirements

Loop Blocking – Summary

Recommendations:

- ❑ choose blocking size as large as possible
 - ❑ leave space for other data
 - ❑ parameterize cache characteristics, especially size
-
- ❑ How many loops should I block?
 - ❑ depends on the data access pattern, i.e. look for load and store operations

Tricked by the compiler

Fortran code example: long and bulky loop

```
DO I1=1,NAT(IMT)
    IA1=IM1+(I1-1)*3
    DO I2=1,NAT(JMT)
        IA2=IM2+(I2-1)*3
        ... statements removed ...
        DX(1)=XNOW(IMT,IA1+1)-XNOW(JMT,IA2+1)
        DX(2)=XNOW(IMT,IA1+2)-XNOW(JMT,IA2+2)
        DX(3)=XNOW(IMT,IA1+3)-XNOW(JMT,IA2+3)
        CX(1)=CM1(1)-CM2(1)
        CX(2)=CM1(2)-CM2(2)
        CX(3)=CM1(3)-CM2(3)
        ... statements removed ...
    ENDDO
    ENDDO
```

Independent of loop indices!!!

Moved the 3 lines above the DO loops – no improvement!
Compiler had done this already!

Only the programmer knows ...

```
subroutine do_calc(...)  
  
real(8),dimension(N,M,O,P):: r,s,t  
  
!---- data initialization  
r = 0.0d0; s = 0.0d0; t = 0  
  
select case(calc_type)  
  case(most_of_the_time)  
    ...  
    r(i,j,k,l) = r(i,j,k,l)  
    s(i,j,k,l) = s(i,j,k,l)  
    ...  
  
  case(rare_event)  
    ...  
    r(i,j,k,l) = r(i,j,k,l)  
    s(i,j,k,l) = s(i,j,k,l)  
    t(i,j,k,l) = t(i,j,k,l)  
    ...  
end select
```

```
subroutine do_calc(...)  
  
real(8),dimension(N,M,O,P):: r,s,t  
  
!---- data initialization  
r = 0.0d0; s = 0.0d0  
  
select case(calc_type)  
  case(most_of_the_time)  
    ...  
    r(i,j,k,l) = r(i,j,k,l) + ...  
    s(i,j,k,l) = s(i,j,k,l) + ...  
    ...  
  
  case(rare_event)  
    t = 0.0d0  
    ...  
    r(i,j,k,l) = r(i,j,k,l) + ...  
    s(i,j,k,l) = s(i,j,k,l) + ...  
    t(i,j,k,l) = t(i,j,k,l) + ...  
    ...  
end select
```

Data structure design

Access your data in the right way

Data structure design

- ❑ “Good advice” from a HPC tutorial: Use data structures to avoid too many index calculations
- ❑ Example: particle simulation in 3D with x, y, z coordinates and some other information about each particle, e.g. distance to origin, particle type
 - ❑ `x[i]`, `y[i]`, `z[i]`, `dist[i]`, `ptype[i]`
 - ❑ turn this into a data structure ...

Data structure design

- ❑ Particle data structure

```
typedef struct particle {  
    double x, y, z;  
    char ptype;  
    double dist;  
} particle_t;
```

- ❑ Is this a good idea?

Data structure design

- ❑ Answer: It depends ...
 - ❑ ... on problem size
 - ❑ ... how you access the data
 - ❑ ... and how often
 - ❑ ... cache, CPU, etc.
- ❑ Example: program with 2 functions/routines
 - ❑ calc() - accesses all parts of particle_t
 - ❑ re-use() - accesses particle.dist only
 - ❑ usage ratio of both functions is 1:1

Data structure design

Advantages of the data structure:

- ❑ `particle_t *p;` (and then allocate N)
- ❑ easy to pass data in function calls
 - ❑ `func(p, N);`
 - ❑ ... vs `func(x, y, z, ptype, dist, N);`
- ❑ flexible: I can add new elements to `particle_t`, without having to change the function interfaces/prototypes

Data structure design

Downside of the data structure:

- ❑ memory access is no longer optimal
 - ❑ neighbouring elements of the same type, e.g. ‘x’ are no longer ‘stride 1’:
 - ❑ $x[i] \rightarrow x[i+1]$: stride 1
 - ❑ $p[i].x \rightarrow p[i+1].x$: ~stride 5 (in this example)
 - ❑ no good cache line usage
 - ❑ compiler cannot optimize loops

Data structure design

Get the best of both worlds!

- ❑ instead of using an “array of struct” (AOS),
- ❑ we can create a “struct of arrays” (SOA):

```
typedef struct particle {  
    double *x, *y, *z;  
    char   *ptype;  
    double *dist;  
} particle_t;
```

Data structure design

Downside of the SOA:

- ❑ memory allocation is more complicated
 - ❑ one call to malloc() for each element
- ❑ need to change code:
 - ❑ $p[i].x \rightarrow p.x[i]$
 - ❑ ... and function prototypes: $f(*p, N) \rightarrow f(p, N)$
- ❑ but that's a 'one time effort'

Data structure design

Advantages of SOA outweigh this by far:

- ❑ memory access is optimal again
- ❑ neighbouring elements of the same type, e.g. ‘x’ are again ‘stride 1’:
 - ❑ $p.x[i] \rightarrow p.x[i+1]$: stride 1
 - ❑ better cache line usage
 - ❑ access to single components, e.g. $p.dist$, do not need to load ‘all the rest’
- ❑ compiler can optimize loops better

Data structure design


- ❑ Example: program with 2 functions/routines
 - ❑ calc() - accesses all parts of particle_t
 - ❑ re-use() - accesses particle.dist only
 - ❑ usage ratio of both functions is 1:1

Data structure design

Comparison – two versions:


- ❑ aos – use ‘array of struct’
- ❑ soa – use ‘struct of arrays’
- ❑ different problem sizes:
 - ❑ number of particles

Data structure design




XeonE5-2660v3 25MB L3 cache

Data structure design



XeonE5-2660v3 25MB L3 cache


Data structure design



Xeon X5550 8MB L3 cache

Data structure design

Data structure comparison: AoS vs SoA



Xeon X5550 8MB L3 cache

Data structure lab (aka tune_labs)

- ❑ download the ZIP file from DTU Learn
- ❑ read the instructions
 - ❑ do parts I and II now
 - ❑ part III needs more information (after lunch)
- ❑ Goals:
 - ❑ design and do performance experiments
 - ❑ extract the relevant data – and plot it

Performance experiments – how to

Think about the ‘design’ of your experiment!


- ❑ What can I measure directly?
 - ❑ Almost always: time!
 - ❑ Time is not always a “good” quantity
- ❑ Adding ‘a priori’ knowledge: for a given problem size we typically know the ...
 - ❑ number of iterations
 - ❑ number of operations (flops, loads, stores)
- ❑ allows us to calculate Mflop/s, iter/s, bandwidth, ...

Performance experiments – how to


Think about the ‘design’ of your experiment!

- ❑ “Where” should I measure?
 - ❑ your measurement should support a model (or an ‘idea’), e.g. ‘performance vs memory footprint’
 - ❑ make use of prior knowledge, e.g. the memory layout of the hardware (L1,...,L3 cache sizes, and latencies/bandwidths)
 - ❑ choose your measuring points, e.g. memory footprint, to map the “interesting areas” in the “performance landscape”

Performance experiments – how to



Performance experiments – how to



Performance experiments – how to

Best practice:

- ❑ isolate the “kernel” you want to benchmark (e.g. in a function)
- ❑ do a pre-measurement call (“warm up”)
- ❑ repeat the measurements to get a “reasonable” run time (e.g. 2-3 secs), take the average
- ❑ make sure to use the same system (CPU type), to get comparable results
- ❑ reduce the “system noise”, i.e. other processes, users, etc – use the batch system

End of lecture 1