

# Getting good performance from your application

Tuning techniques for serial programs on  
cache-based computer systems

# Application Tuning


Selected Topics

# Application Tuning


- ❑ Selected Topics:
  - ❑ 32- vs 64-bit
  - ❑ binary data portability
  - ❑ floating point numbers and IEEE 754
    - ❑ compiler options
  - ❑ a case study
  
- ❑ Summary

# 32-bit vs 64-bit issues

- ❑ 64-bit operating systems
- ❑ Implication: The address space of a single application can be larger than 4 GB



# 32-bit vs 64-bit issues



- ❑ Addresses ≠ Data
- ❑ An 'n'-byte data type fills always n bytes in memory (byte addressable)
- ❑ I.e. the next element is n bytes further in memory
- ❑ This increment is not related to the size of the addresses (32-bit or 64-bit)

# 32-bit vs 64-bit issues

<u>C data type</u>	<u>ILP32 (bits)</u>	<u>LP64 (bits)</u>
<i>char</i>	8	<i>same</i>
<i>short</i>	16	<i>same</i>
<i>int</i>	32	<i>same</i>
<i>long</i>	32	<b>64</b>
<i>long long</i>	64	<i>same</i>
<i>pointer</i>	32	<b>64</b>
<i>enum</i>	32	<i>same</i>
<i>float</i>	32	<i>same</i>
<i>double</i>	64	<i>same</i>
<i>long double</i>	128	<i>same</i>

UNIX and Linux support LP64; Windows 64-bit uses LLP64, where long stays 32 bits

# (p)l~~d~~ and LD\_LIBRARY\_PATH

How to check which shared-libraries are loaded?

- ❑ Static check: use the ldd command
  - ❑ \$ ldd executable
- ❑ Dynamic check: use pl~~d~~ on the PID
  - ❑ \$ pl~~d~~ pid
  - ❑ Solaris only
  - ❑ there are scripts available for Linux as well
  - ❑ we have installed pl~~d~~ on the DTU HPC cluster

# (p)ldd and LD\_LIBRARY\_PATH

- ❑ How to change the search path for dynamic libraries?
  - ❑ Use LD\_LIBRARY\_PATH – but use it with care!
- ❑ Best practice:
  - ❑ Compile the path into your application:
    - ❑ GCC: -Wl,-rpath <path\_to\_lib>
    - ❑ ld.so will then use this path
  - ❑ Avoid LD\_LIBRARY\_PATH in your shell environment – use a wrapper script for the application
  - ❑ Check out this blog note ([www.hpc.dtu.dk/?page\\_id=1180](http://www.hpc.dtu.dk/?page_id=1180)), too!

# Binary data storage

- ❑ Storing your data in binary format
- ❑ Advantages:
  - ❑ compact
  - ❑ fast
  - ❑ no loss of precision
- ❑ Drawbacks:
  - ❑ not “human readable”
  - ❑ data analysis more complicated
  - ❑ and ...

# Binary data storage

- ❑ Example: integer 0x12345678 (hexadecimal)

```
value = 0x12345678;           // 305419896
printf("%d\n", value);
fwrite(&value, sizeof(value), 1, fptr);
```

- ❑ Write it ...

- ❑ ... on i386:

```
305419896
Architecture: i386
Value written to endian_i386.dat.
```

- ❑ ... on SPARC:

```
305419896
Architecture: sparc
Value written to endian_sparc.dat.
```

# Binary data storage

## ❑ Read it:

```
fread(&value, sizeof(value), 1, fptr);  
printf("%d\n", value);
```

## ❑ on i386 data from i386:

Architecture: i386

Read from endian\_i386.dat: 305419896

## ❑ on i386 data from SPARC:

Architecture: i386

Read from endian\_sparc.dat: 2018915346




# Little Endian vs Big Endian

- ❑ The order in which the bits are interpreted has not been standardized!
- ❑ Two 'popular' formats in use
  - ❑ Big Endian – SPARC, PowerPC, ...
  - ❑ Little Endian – Intel x86, AMD64, ...
- ❑ This is an issue when using the same binary data file on both platforms ...

# Little Endian vs Big Endian

- ❑ Example: integer 0x12345678 (hexadecimal)



- ❑ Check with 'od' command:

```
$ od -x endian_sparc.dat
0000000 1234 5678
0000004
```

```
$ od -x endian_i386.dat
0000000 7856 3412
0000004
```

# Floating point numbers & IEEE 754

## Lesser known side effects of IEEE 754:

- ❑ Will this code run or fail?

```
#include <stdio.h>
#include <math.h>

int
main(int argc, char *argv[]) {

    double x;

    for(int i = 0; i < 10; i++) {
        x = sqrt(5.0 - i);
        printf("%lf\n", x);
    }
}
```

# Floating point numbers & IEEE 754

## Lesser known side effects of IEEE 754:

- ❑ What do you prefer?

```
$ gcc -ftrapex trapex.c -lm  
$ ./trapex  
2.236068  
2.000000  
1.732051  
1.414214  
1.000000  
0.000000  
-nan  
-nan  
-nan  
-nan  
$
```

IEEE 754 compliant!

```
$ suncc -ftrap=common -o trapex  
$ ./trapex  
2.236068  
2.000000  
1.732051  
1.414214  
1.000000  
0.000000  
Floating point exception!  
$
```

IEEE 754 compliant!

# Floating point numbers & IEEE 754

## Lesser known side effects of IEEE 754:

- ❑ The IEEE 754 standard doesn't "allow" traps on floating point exceptions, like invalid arguments, division by zero, over- and underflows
- ❑ Some compilers provide options to change that.
  - ❑ Studio: -ftrap=<exception\_list>, e.g. common
  - ❑ Intel: -fp-trap=<exception\_list>, e.g. common
- ❑ However: GCC has no such option, needs to be implemented by the programmer via library calls (see 'man fenv' – and next slide)

# Floating point numbers & IEEE 754

## Enabling floating point traps with GCC

```
#include <stdio.h>
#include <math.h>

#define __USE_GNU
#include "fenv.h"

void enable_ftraps(void) {
    feenableexcept (FE_DIVBYZERO);
    feenableexcept (FE_INVALID);
    feenableexcept (FE_OVERFLOW); }

int main(int argc, char *argv[]) {
    double x;
    enable_ftraps();
    for(int i = 0; i < 10; i++) {
        x = sqrt(5.0 - i);
        printf("%lf\n", x);
    }
}
```

# Floating point numbers & IEEE 754

Now it works with GCC ...

```
$ gcc -o trapex2 trapex2.c -lm
$ ./trapex2
2.236068
2.000000
1.732051
1.414214
1.000000
0.000000
Floating point exception
$
```

IEEE 754 compliant!  
not(!)

# Floating point numbers & IEEE 754

- ❑ Compilers do not re-arrange your arithmetic expressions – unless you ask for it!
- ❑ This is part of the optimization flags (and varies from compiler to compiler!)
- ❑ Example on the next slides shows this for a division inside a for-loop
  - ❑ this is done with an old Studio compiler, but illustrates nicely what happens
  - ❑ the optimization option is ‘-fsimple=[0|1|2]’, from none (0) to aggressive arithmetic optimizations (2).

# Floating point numbers & IEEE 754

## Effects of -fsimple:

- ❑ compiled with -fast -xrestrict -fsimple=0:

```
1. void
2. divvec(int n, double div, double *a, double *b) {
3.
4.     int i;
5.

Source loop below has tag L1
L-1 scheduled with steady-state cycle count = 17
L-1 has 1 loads, 1 stores, 2 prefetches, 0 FPadds,
        0 FPmuls, and 1 FPdivs per iteration
6.     for(i = 0; i < n; i++)
7.         b[i] = a[i]/div;
8. }
```

UltraSPARC

# Floating point numbers & IEEE 754

## Effects of -fsimple:

- ❑ compiled with -fast -xrestrict -fsimple=2:

```
1. void
2. divvec(int n, double div, double *a, double *b) {
3.
4.     int i;
5.

Source loop below has tag L1
L-1 scheduled with steady-state cycle count = 1
L-1 unrolled 8 times
L-1 has 1 loads, 1 stores, 4 prefetches, 0 FPadds,
        1 FPMuls, and 0 FPdivs per iteration
6.     for(i = 0; i < n; i++)
7.         b[i] = a[i]/div;
8. }
```

UltraSPARC

# A closer look ...

- ❑ on x86\_64 Linux, we have no compiler that does give us the information we want ... i.e. the number of fpmul and fpdiv
- ❑ What now???
- ❑ We need some other tools, to get this information.
  - ❑ objdump – a Linux tool “to look into” executables
  - ❑ objdump has a lot of options – we need just one
    - ❑ objdump -S file.o : show disassembly and source intermixed
  - ❑ ... but how to interpret the assembly code???

# A closer look ...

The really quick guide to x86 assembly

- ❑ instructions of interest: double precision floating point multiplication and division
- ❑ multiplication: \*mul\*d
  - ❑ mulsd, mulpd, vmulsd, vmulpd
- ❑ division: \*div\*d
  - ❑ divsd, divpd, vdivsd, vdivpd
- ❑ ‘s’ is for single data, ‘p’ for packed data (vector)
- ❑ the leading letter indicates vector type

# A closer look ...

compiled with ‘suncc -g -fast -fsimple=0’:

```
$ objdump -S divvec.o | egrep "mul.d|div.d"
 9d:    c5 e3 5e e0          vdivsd %xmm0,%xmm3,%xmm4
 b8:    c5 d3 5e f0          vdivsd %xmm0,%xmm5,%xmm6
 c8:    c5 43 5e c0          vdivsd %xmm0,%xmm7,%xmm8
 d8:    c5 33 5e d0          vdivsd %xmm0,%xmm9,%xmm10
 e8:    c5 23 5e e0          vdivsd
%xmm0,%xmm11,%xmm12
 f8:    c5 13 5e f0          vdivsd
%xmm0,%xmm13,%xmm14
 108:   c5 83 5e c8         vdivsd %xmm0,%xmm15,%xmm1
 118:   c5 eb 5e d8         vdivsd %xmm0,%xmm2,%xmm3
 149:   c5 f3 5e d0         vdivsd %xmm0,%xmm1,%xmm2
 1cf:   c5 35 5e d6         vdivpd %ymm6,%ymm9,%ymm10
 1ed:   c5 25 5e e6         vdivpd
%ymm6,%ymm11,%ymm12
 21a:   c5 45 5e c6         vdivpd %ymm6,%ymm7,%ymm8
 259:   c5 db 5e e8         vdivsd %xmm0,%xmm4,%xmm5
```

■ Single data and packed: multi-versioning of loop

# A closer look ...

compiled with ‘suncc -g -fast -fsimple=2’:

```
$ objdump -S divvec.o | egrep "mul.d|div.d"
82:    c5 d3 5e f0          vdivsd %xmm0,%xmm5,%xmm6
a8:    c5 4b 59 04 06      vmulsd (%rsi,%rax,1),%xmm6,%xmm8
ba:    c5 4b 59 4c 06 08    vmulsd 0x8(%rsi,%rax,1),%xmm6,%xmm9
ca:    c5 4b 59 54 06 10    vmulsd 0x10(%rsi,%rax,1),%xmm6,%xmm10
d6:    c5 4b 59 5c 06 18    vmulsd 0x18(%rsi,%rax,1),%xmm6,%xmm11
e2:    c5 4b 59 64 06 20    vmulsd 0x20(%rsi,%rax,1),%xmm6,%xmm12
ee:    c5 4b 59 6c 06 28    vmulsd 0x28(%rsi,%rax,1),%xmm6,%xmm13
fa:    c5 4b 59 74 06 30    vmulsd 0x30(%rsi,%rax,1),%xmm6,%xmm14
106:   c5 4b 59 7c 06 38    vmulsd 0x38(%rsi,%rax,1),%xmm6,%xmm15
130:   c5 cb 59 3c 06      vmulsd (%rsi,%rax,1),%xmm6,%xmm7
19b:   c4 41 45 5e c8      vdivpd %ymm8,%ymm7,%ymm9
1cf:   c4 41 1d 59 e9      vmulpd %ymm9,%ymm12,%ymm13
1ee:   c4 41 0d 59 f9      vmulpd %ymm9,%ymm14,%ymm15
21a:   c4 41 2d 59 d9      vmulpd %ymm9,%ymm10,%ymm11
25e:   c5 e3 5e c0          vdivsd %xmm0,%xmm3,%xmm0
270:   c5 fb 59 24 06      vmulsd (%rsi,%rax,1),%xmm0,%xmm4
```

# A closer look ...

compiled with ‘suncc -g -fast -fsimple=2’

- ❑ ... and using the ‘restrict’ keyword:

```
$ objdump -S divvec.o | egrep "mul.d|div.d"
7d:    c4 41 35 5e da          vdivpd %ymm10,%ymm9,%ymm11
af:    c4 41 0d 59 fb          vmulpd %ymm11,%ymm14,%ymm15
ce:    c4 c1 75 59 d3          vmulpd %ymm11,%ymm1,%ymm2
fa:    c4 41 1d 59 eb          vmulpd %ymm11,%ymm12,%ymm13
140:   c5 db 5e e8          vdivsd %xmm0,%xmm4,%xmm5
150:   c4 a1 53 59 34 0e      vmulsd (%rsi,%r9,1),%xmm5,%xmm6
```

- ❑ mostly multiplications – it works!
- ❑ the multi-versioning is gone – only the vectorized code (and a clean-up loop) left!

# A closer look ... now with GCC

- ❑ -O3 does not enable math optimizations, like replacing a division with a constant by a multiplication with the inverse
- ❑ to enable this, we need -ffast-math option
- ❑ to get some optimization information (at compile time), we can use -fopt-info

# A closer look ... now with GCC

compiled with ‘gcc -g -O3 -fopt-info’

```
$ gcc -g -O3 -fopt-info -c divvec.c
divvec.c:11:5: optimized: loop vectorized using 16 byte
vectors
divvec.c:11:5: optimized: loop versioned for vectorization
because of possible aliasing
```

```
$ objdump -S divvec.o | egrep "mul.d|div.d"
35: 66 0f 5e c8           divpd  %xmm0,%xmm1
55: f2 0f 5e c2           divsd  %xmm2,%xmm0
70: f2 0f 5e c2           divsd  %xmm2,%xmm0
```

- ❑ GCC does multi-versioning, too!
- ❑ no multiplications – only divisions!

# A closer look ... now with GCC

compiled with ‘gcc -g -O3 -ffast-math -fopt-info’

```
$ gcc -g -O3 -ffast-math -fopt-info -c divvec.c
divvec.c:11:5: optimized: loop vectorized using 16 byte
vectors
divvec.c:11:5: optimized: loop versioned for vectorization
because of possible aliasing
```

```
$ objdump -S divvec.o | egrep "mul.d|div.d"
1e: f2 0f 5e c1           divsd    %xmm1,%xmm0
45: 66 0f 59 ca          mulpd    %xmm2,%xmm1
60: f2 0f 59 04 c6        mulsd    (%rsi,%rax,8),%xmm0
80: f2 0f 59 c8          mulsd    %xmm0,%xmm1
```

- ❑ GCC does multi-versioning, too!
- ❑ mostly multiplications – it works!

# A closer look ... now with GCC

compiled with ‘gcc -g -O3 -ffast-math -fopt-info’

- ❑ ... and using the ‘restrict’ keyword

```
$ gcc -g -O3 -ffast-math -fopt-info -c divvec.c
divvec.c:11:5: optimized: loop vectorized using 16 byte
vectors
```

```
$ objdump -S divvec.o | egrep "mul.d|div.d"
10: f2 0f 5e c1           divsd    %xmm1,%xmm0
35: 66 0f 59 ca          mulpd    %xmm2,%xmm1
52: f2 0f 59 04 c6        mulsd    (%rsi,%rax,8),%xmm0
```

- ❑ multi-versioning is gone!
- ❑ less instructions than with the suncc compiler!?

# A closer look ... now with GCC

GCC does not unroll loops with -O3!

- ❑ we need to add -funroll-loops, too!

```
gcc -g -O3 -ffast-math -funroll-loops -fopt-info -c divvec.c
divvec.c:11:5: optimized: loop vectorized using 16 byte
vectors
divvec.c:11:16: optimized: loop unrolled 7 times
```

```
$ objdump -S divvec.o | egrep "mul.d|div.d" | grep -c mulpd
15
```

- ❑ now we have more mulpd instructions (14 more from loop unrolling)
- ❑ Lesson learned: check what ‘-O3’ does!
  - ❑ Reminder: ‘gcc -g -Q --help=optimizers ...’

# A closer look ... summary

- ❑ Even with no compiler commentary, we can still get useful information, using
  - ❑ extra compiler options like -fopt-info (gcc)
  - ❑ tools like objdump – plus some basic knowledge about assembly code
- ❑ Caveat: this kind of analysis is feasible on small code kernels, only!
- ❑ Best practice: extract a small kernel from larger application, do the tests/tuning (- and reinsert).

# Summary

- ❑ You have now heard about
  - ❑ tuning techniques
  - ❑ tools: compilers, analysis tools
  - ❑ libraries (see extra slide deck – self study)
  - ❑ other performance parameters
  - ❑ debuggers: try Totalview
- ❑ Now you have to apply that and get experience!
- ❑ But never forget:

Correct code has the highest priority – not speed!