

Introduction to Parallel Computing

Overview

- ❑ Parallel Computing
 - ❑ The name of the game
 - ❑ Programming models
 - ❑ Caches revisited
 - ❑ Parallel architectures
 - ❑ Multi-core everywhere
 - ❑ Hardware examples
 - ❑ Memory and multi-core

Parallelism is everywhere

- ❑ In today's computer installations one has many levels of parallelism:
 - ❑ Instruction level (ILP)
 - ❑ Chip level (multi-core, multi-threading)
 - ❑ System level (SMP)
 - ❑ GP-GPUs
 - ❑ Clusters

What is Parallelization?

The Name of the Game

What is Parallelization?

An attempt of a definition:

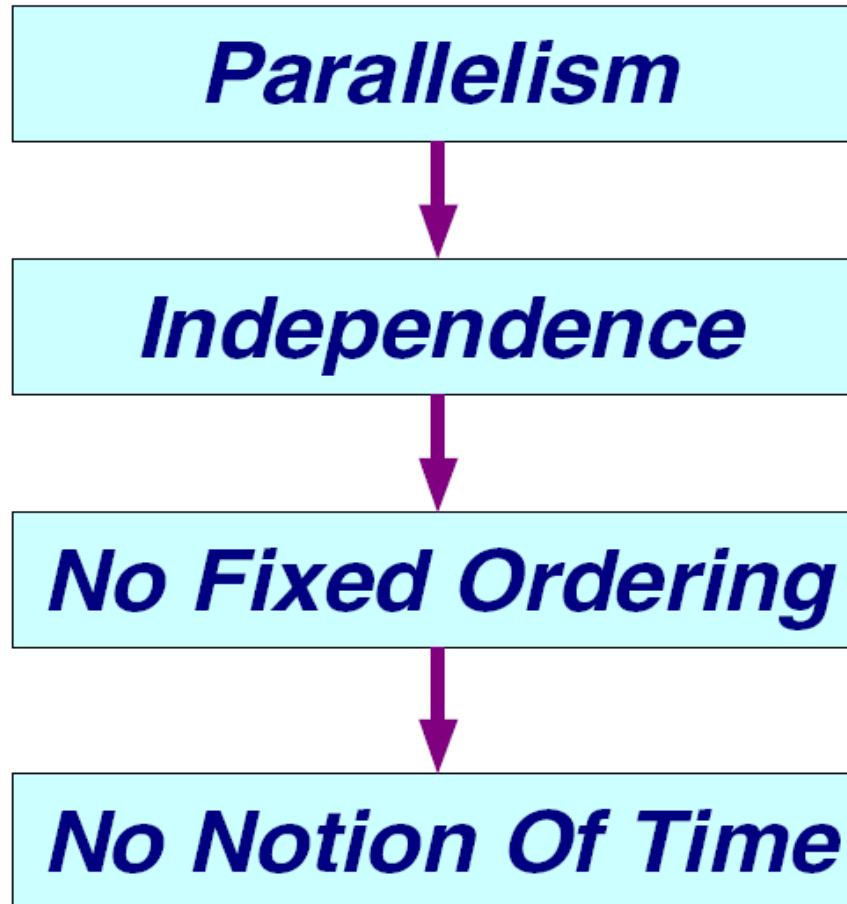
“Something” is parallel, if there is a certain level of independence in the order of operations

“Something” can be:

- ▶ A collection of program statements
- ▶ An algorithm
- ▶ A part of your program
- ▶ The problem you are trying to solve



Parallelism – when?

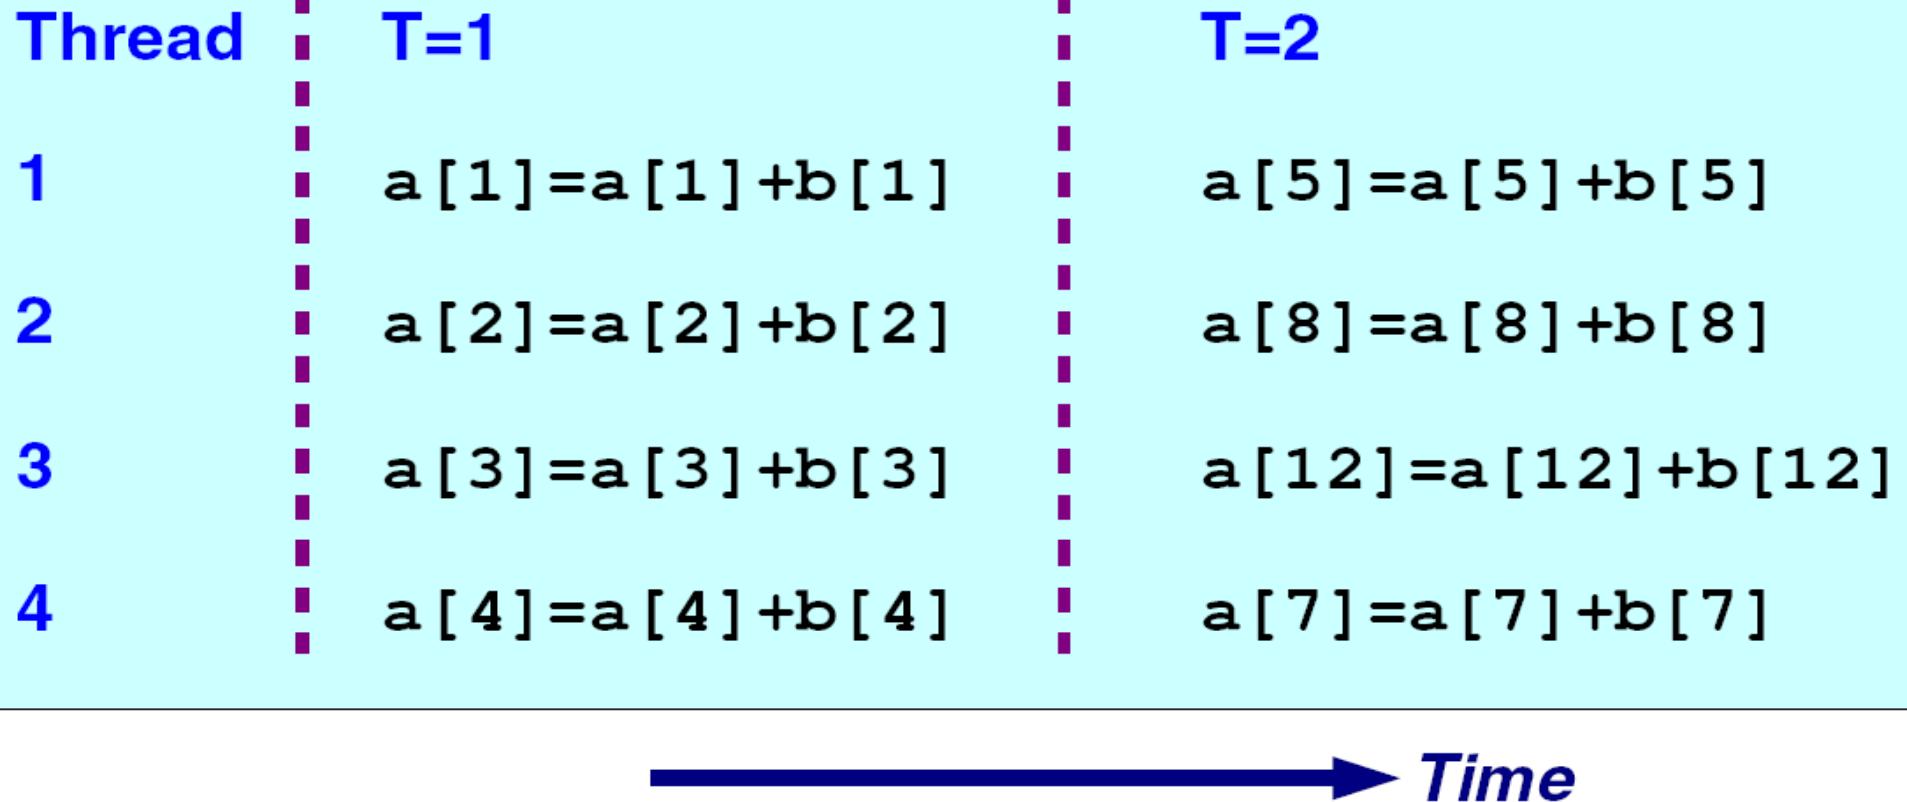


Something that does not follow this rule is **not parallel !!!**

Parallelism – Example 1

```
for (i=0; i<n; i++)  
    a[i] = a[i] + b[i];
```

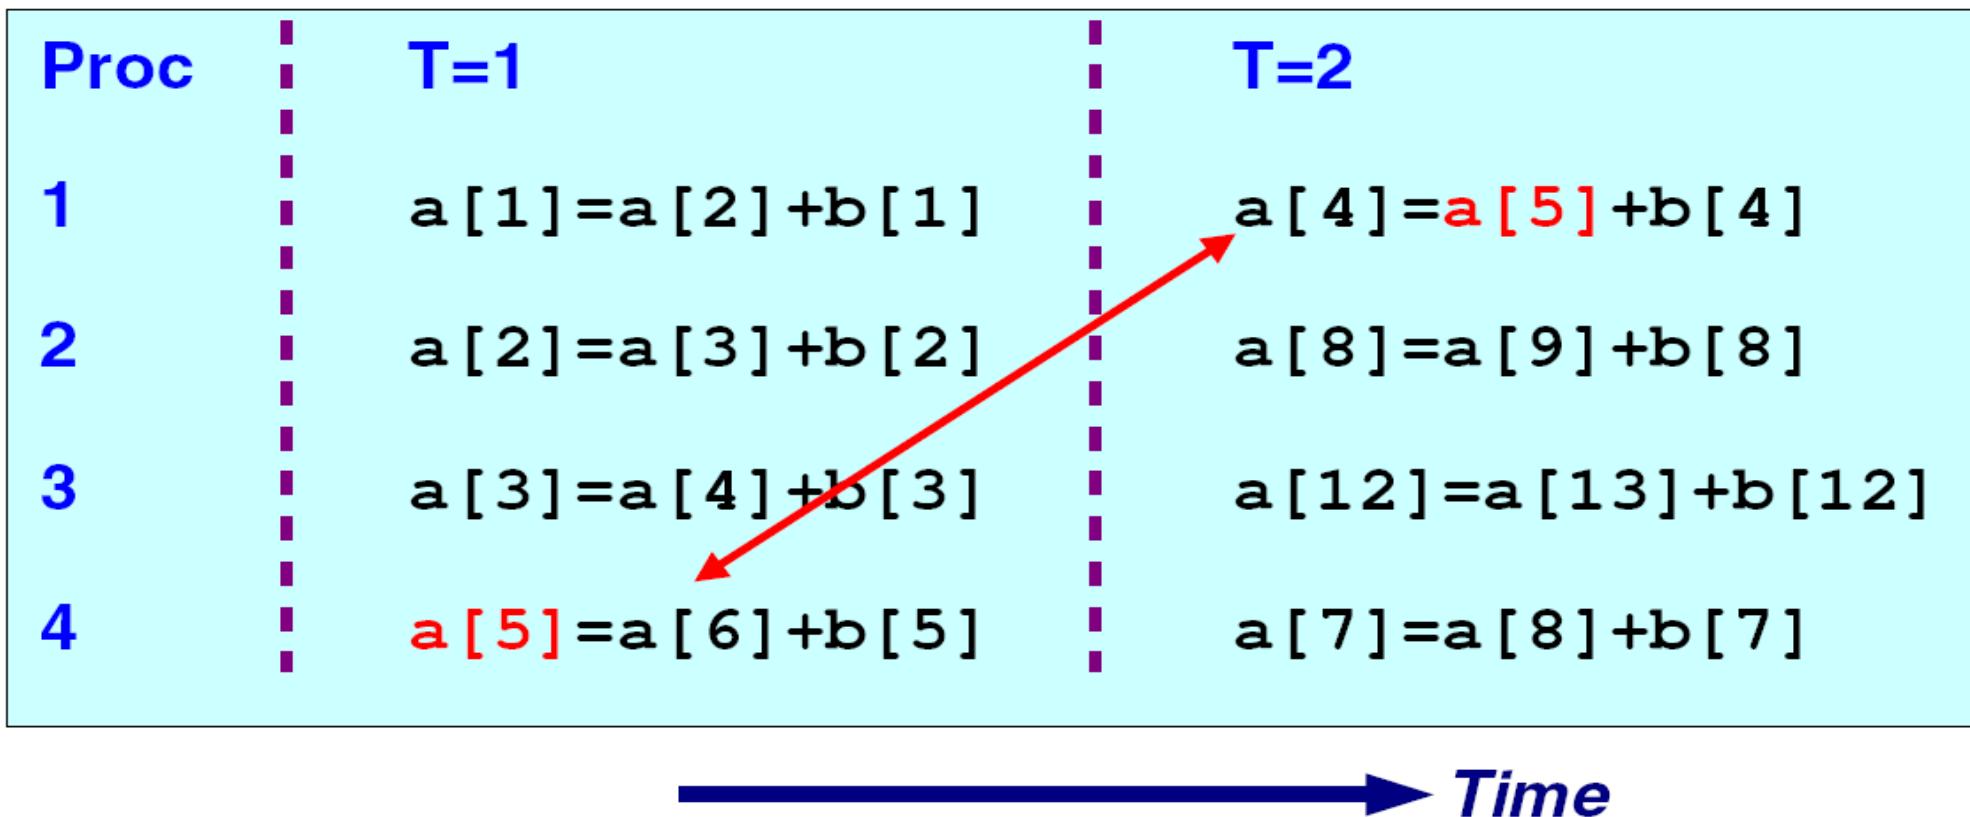
Every iteration in this loop is independent of the other iterations



Parallelism – Example 2

```
for (i=0; i<n; i++)  
    a[i] = a[i+1] + b[i];
```

*This operation is
not parallel !*



Parallelism – Results example 2

Results for P=1

```
12512501.0  
12512501.0  
12512501.0  
12512501.0
```

Results for P=8

```
12512508.0  
12512508.0  
12512508.0  
12512508.0
```

Results for P=32

```
12512526.0  
12512530.0  
12512528.0  
12512527.0
```

Results for P=64

```
12512548.0  
12512545.0  
12512549.0  
12512547.0
```

- ❑ parallel version of example 2 was run 4 times each on 1, 8, 32 and 64 threads/processors
- ❑ Output: sum over all elements of vector a
- ❑ Except for P=1, the results are:
 - ❑ Wrong
 - ❑ Inconsistent
 - ❑ NOT reproducible
- ❑ This is called a '**Data Race**'

Parallelism

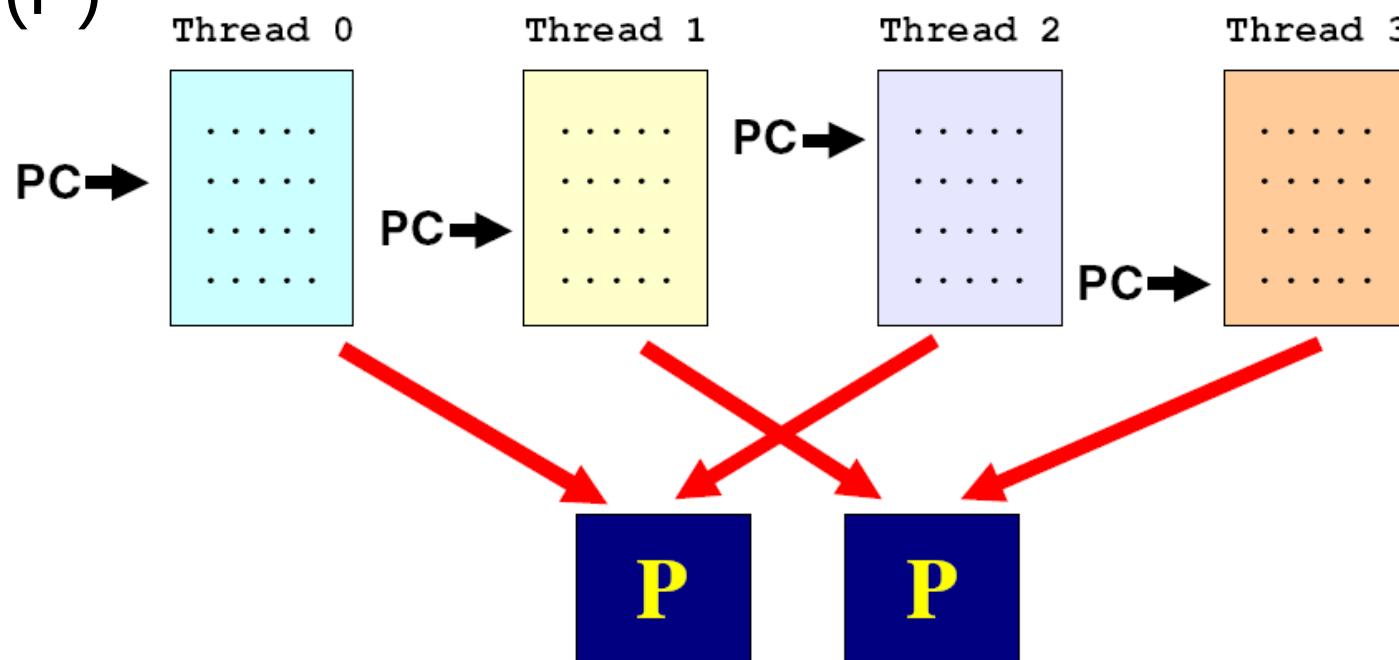
Fundamental problem:

```
for (i = 0; i < n; i++ )  
    a[i] = a[i+M] + b[i];
```

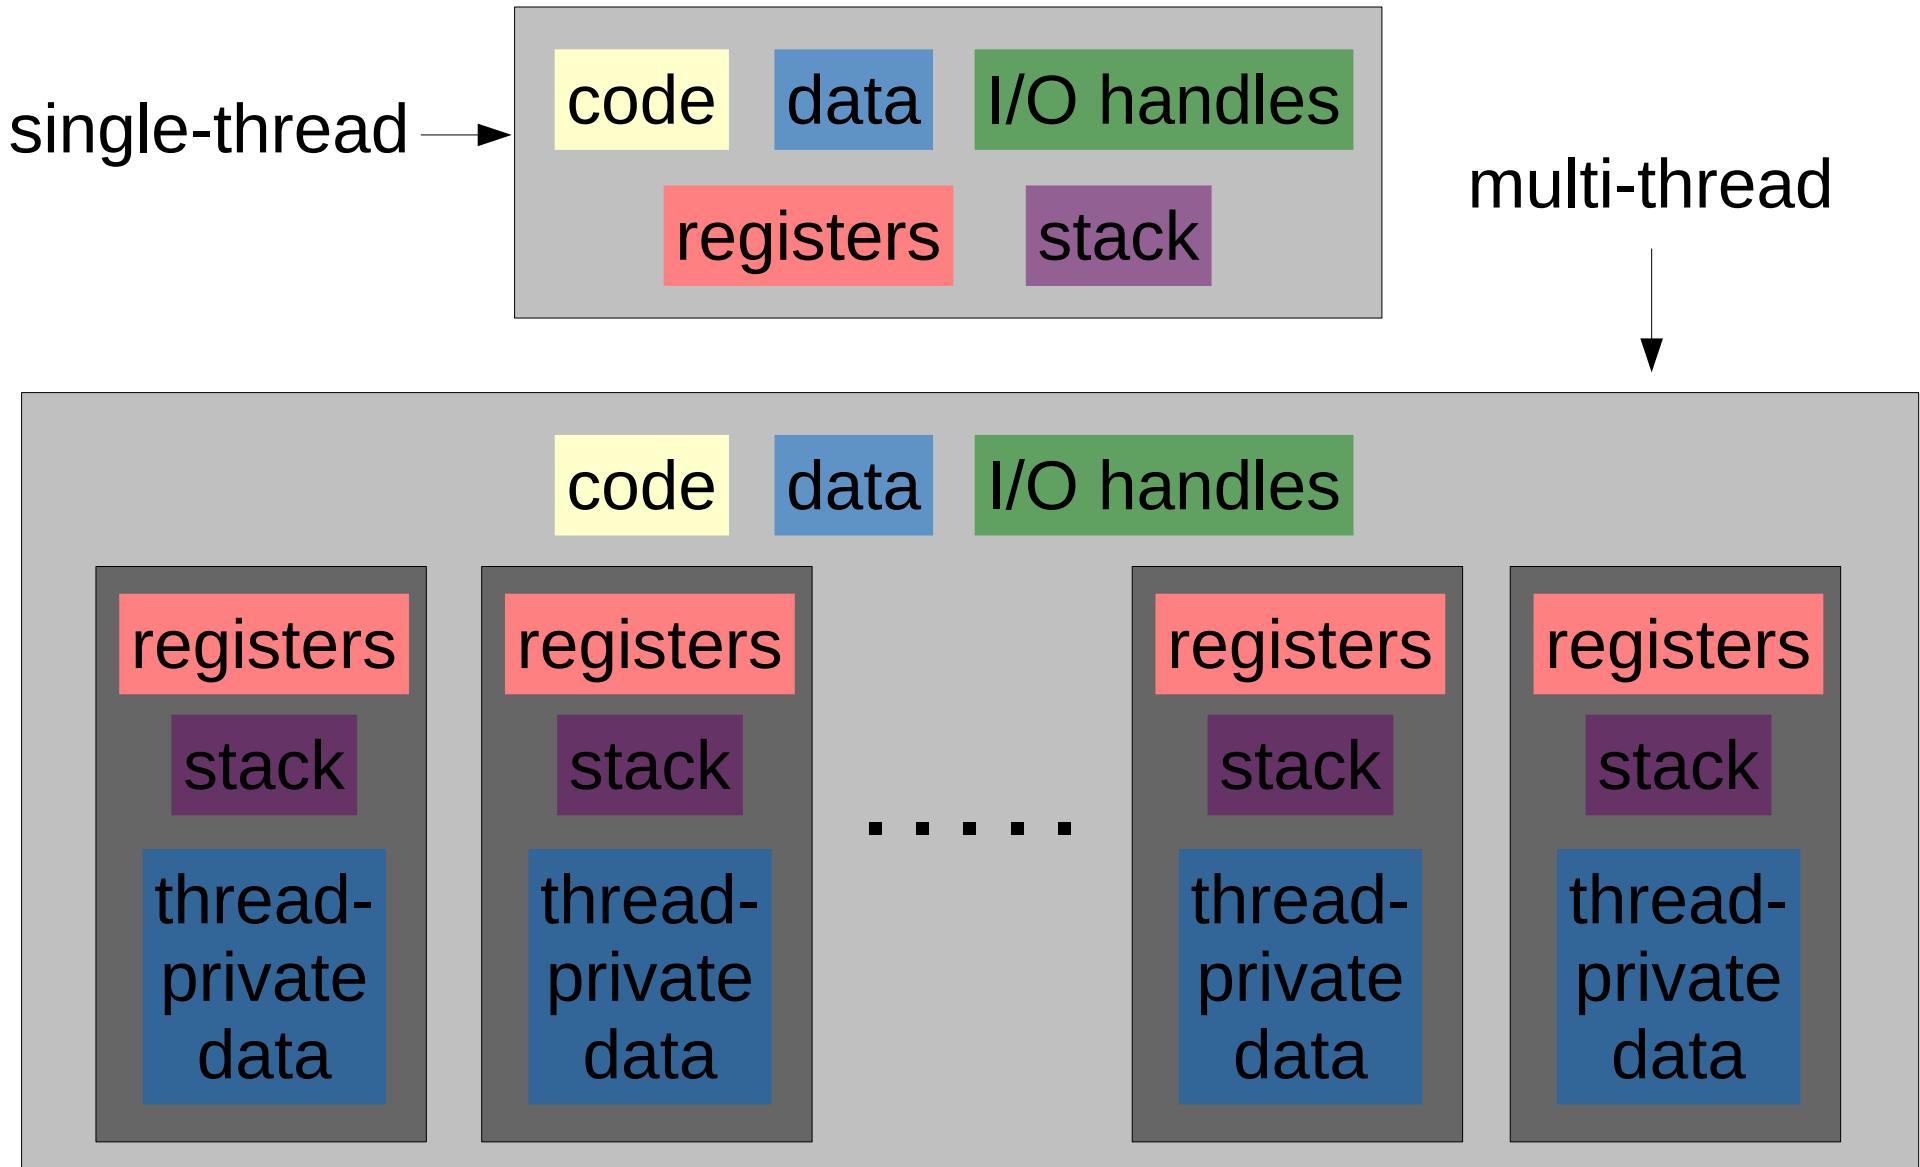
M = 0 : parallel
M >= 1 : not parallel

What is a Thread?

- ❑ Loosely said, a thread consists of a series of instructions with it's own program counter ("PC") and state
- ❑ A parallel program will execute threads in parallel
- ❑ These threads are then scheduled onto processing units (P)

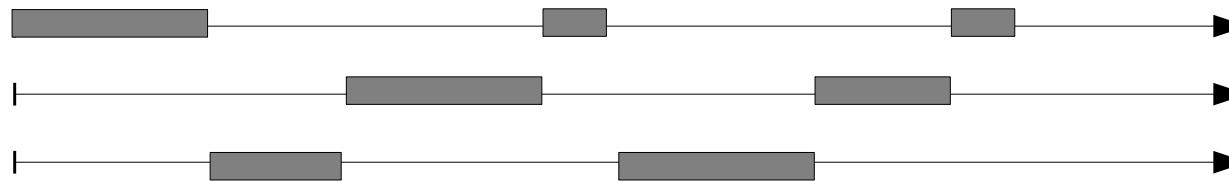


Single- vs. multi-threaded



Parallelism vs Concurrency

Concurrent, non-parallel execution:

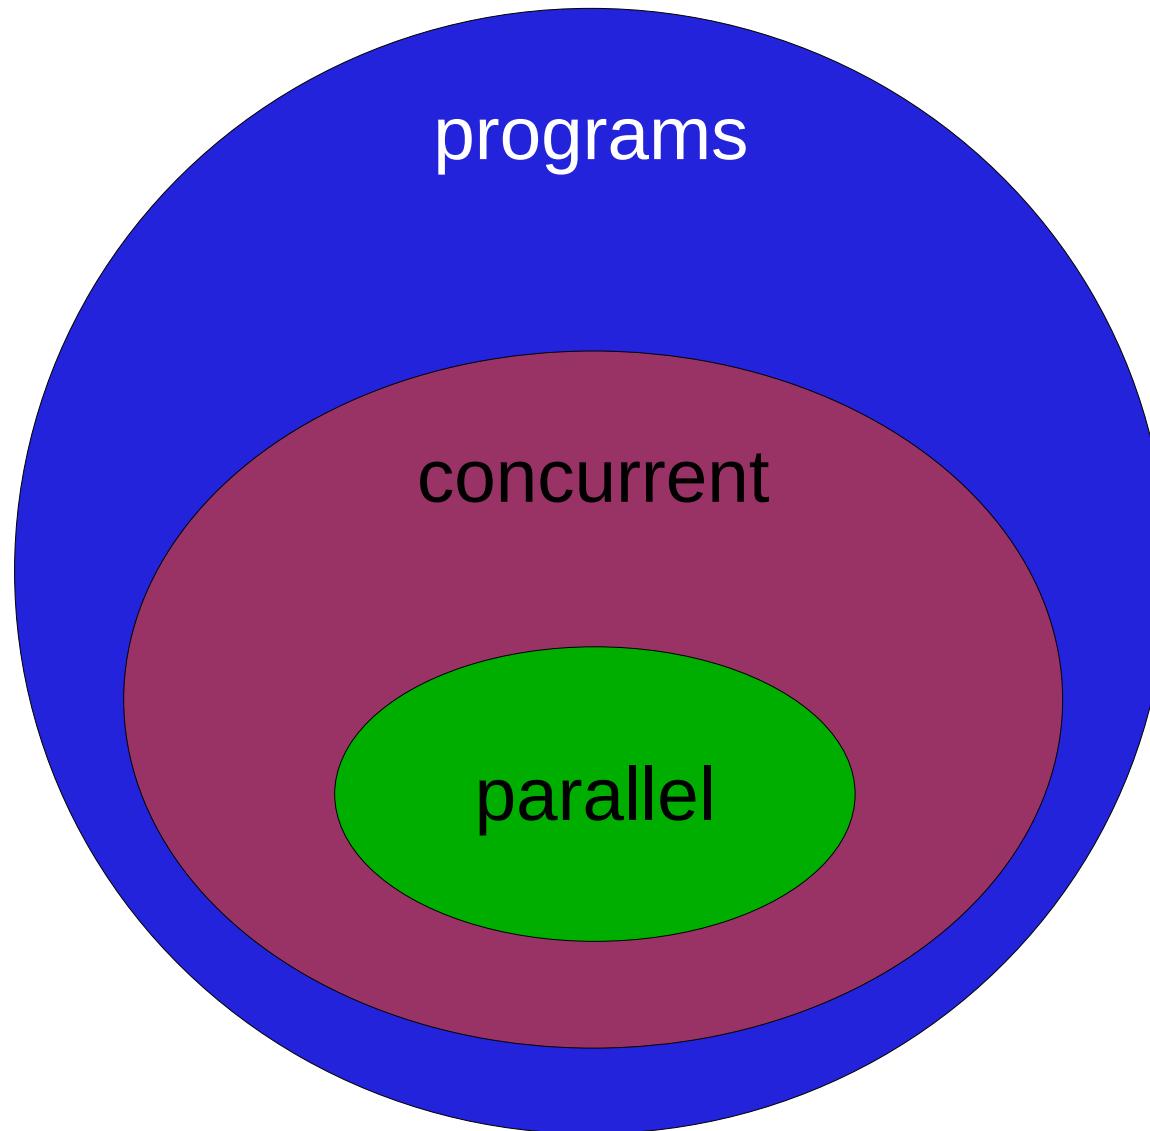


e.g. multiple threads on a single core CPU

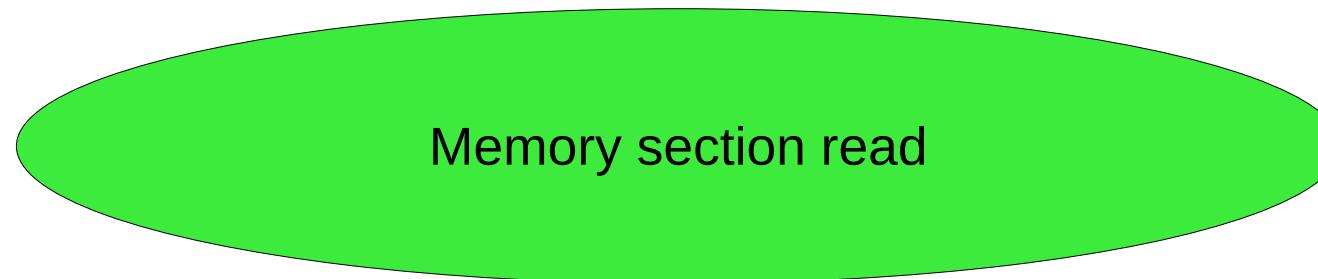
Concurrent, and parallel execution



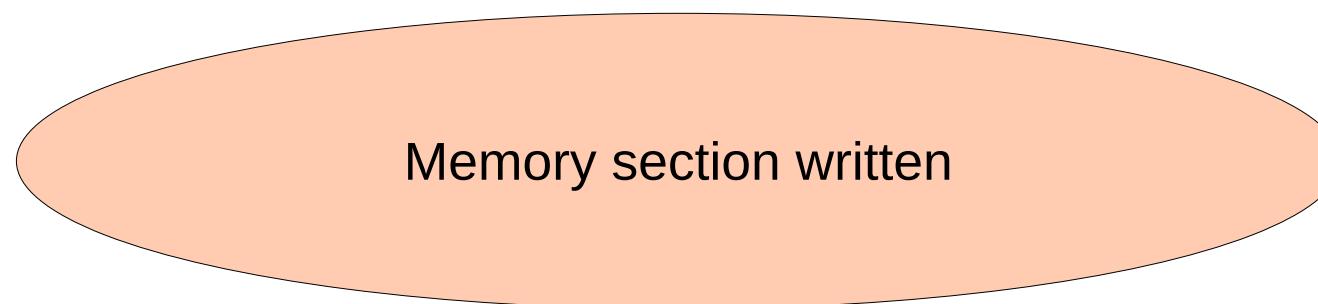
Parallelism vs Concurrency



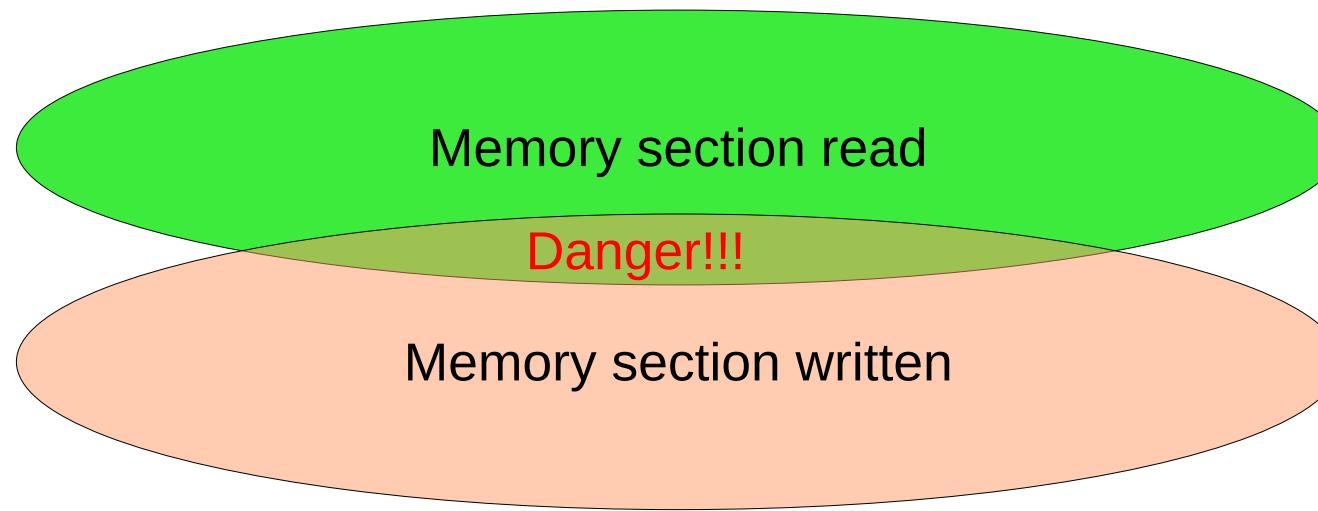
Parallelism – Memory Access



parallel



not parallel! -
... unless one can
protect the
overlap area



Parallelism – Data Races

- ❑ Race conditions can be nasty – and difficult to detect:
 - ❑ Numerical results differ (slightly) from run to run
 - ❑ Difficult to distinguish from numerical side effects
 - ❑ Changing the number of threads can make the problem disappear – or appear again
 - ❑ Shows very often first when using many threads, i.e. late in development

Parallelism – Data Race Detection

- ❑ A few data race detection tools are available:
 - ❑ Intel: ThreadChecker (now: Inspector XE)
 - ❑ part of the Intel OneAPI HPC edition
 - ❑ not so easy to use
 - ❑ GCC/LLVM: come with a thread checker library
 - ❑ compiler option: `-fsanitize-thread`
 - ❑ problem: doesn't work well with OpenMP 'out-of-the-box'
 - ❑ solution: Archer (more in a later lecture)
- ❑ Those tools instrument your code, i.e. a detection run takes substantially longer

Numerical Results

Consider:

$$A = B + C + D + E$$

Serial Processing

$$A = B + C$$

$$A = A + D$$

$$A = A + E$$

Parallel Processing

Thread 0

$$T1 = B + C$$

Thread 1

$$T1 = T1 + T2$$

- ☞ *The roundoff behaviour is different and so the numerical results may be different too*
- ☞ *This is natural for parallel programs, but it may be hard to differentiate it from an ordinary bug*

Basic concepts

- ❑ Consider the following code with two loops

```
for (i = 0; i < n; i++)
    a[i] = b[i] + c[i];
```

```
for (i = 0; i < n; i++)
    d[i] = a[i] + e[i];
```

- ❑ Running this in parallel over i might give the wrong answer.

Basic concepts – the barrier

- ❑ The problem can be fixed:

```
for (i = 0; i < n; i++)  
    a[i] = b[i] + c[i];
```

wait!

```
for (i = 0; i < n; i++)  
    d[i] = a[i] + e[i];
```

- ❑ The barrier assures that no thread starts working on the second loop before the work on loop one is finished.

Basic concepts – the barrier

When to use barriers?

- ❑ To assure data integrity, e.g.
 - ❑ after one iteration in a solver
 - ❑ between parts of the code that read and write the same variables
- ❑ Barriers are expensive and don't scale to a large number of threads

Basic concepts – reduction

- ❑ A typical code fragment:

```
for( i = 0; i < n; i++ ) {  
    ...  
    sum += a[i];  
    ...  
}
```

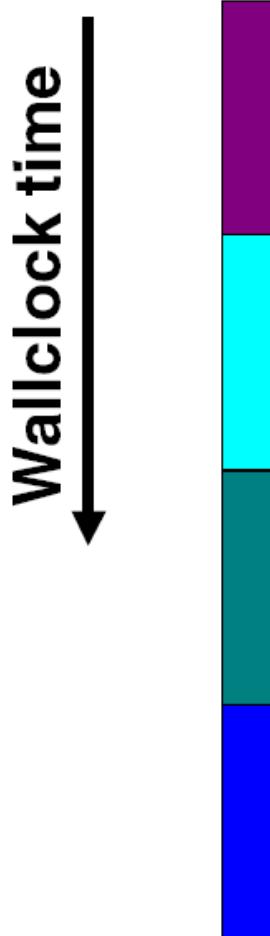
- ❑ This loop can not run in parallel, unless the update of sum is protected.  **serial code**
- ❑ An operation like the above is called a “reduction” operation, and there are ways to handle this issue (more later...).

Parallel Overhead

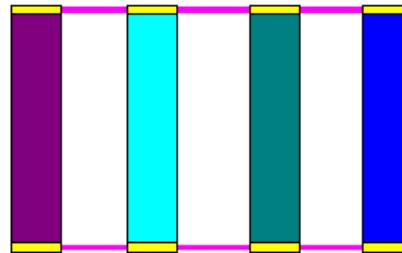
- ❑ The total CPU time may exceed the serial CPU time:
 - ✓ The newly introduced parallel portions in your program need to be executed
 - ✓ Threads need time for sending data to each other and for synchronizing (“communication”)
- ❑ Typically, things also get worse when increasing the number of threads
- ❑ Efficient parallelization is about minimizing the communication overhead

Communication

Serial Execution

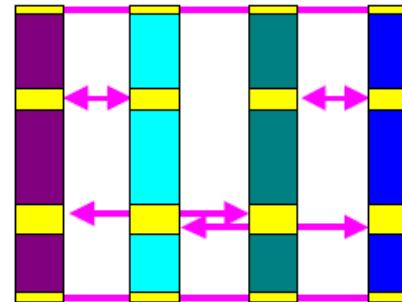


Parallel - Without communication



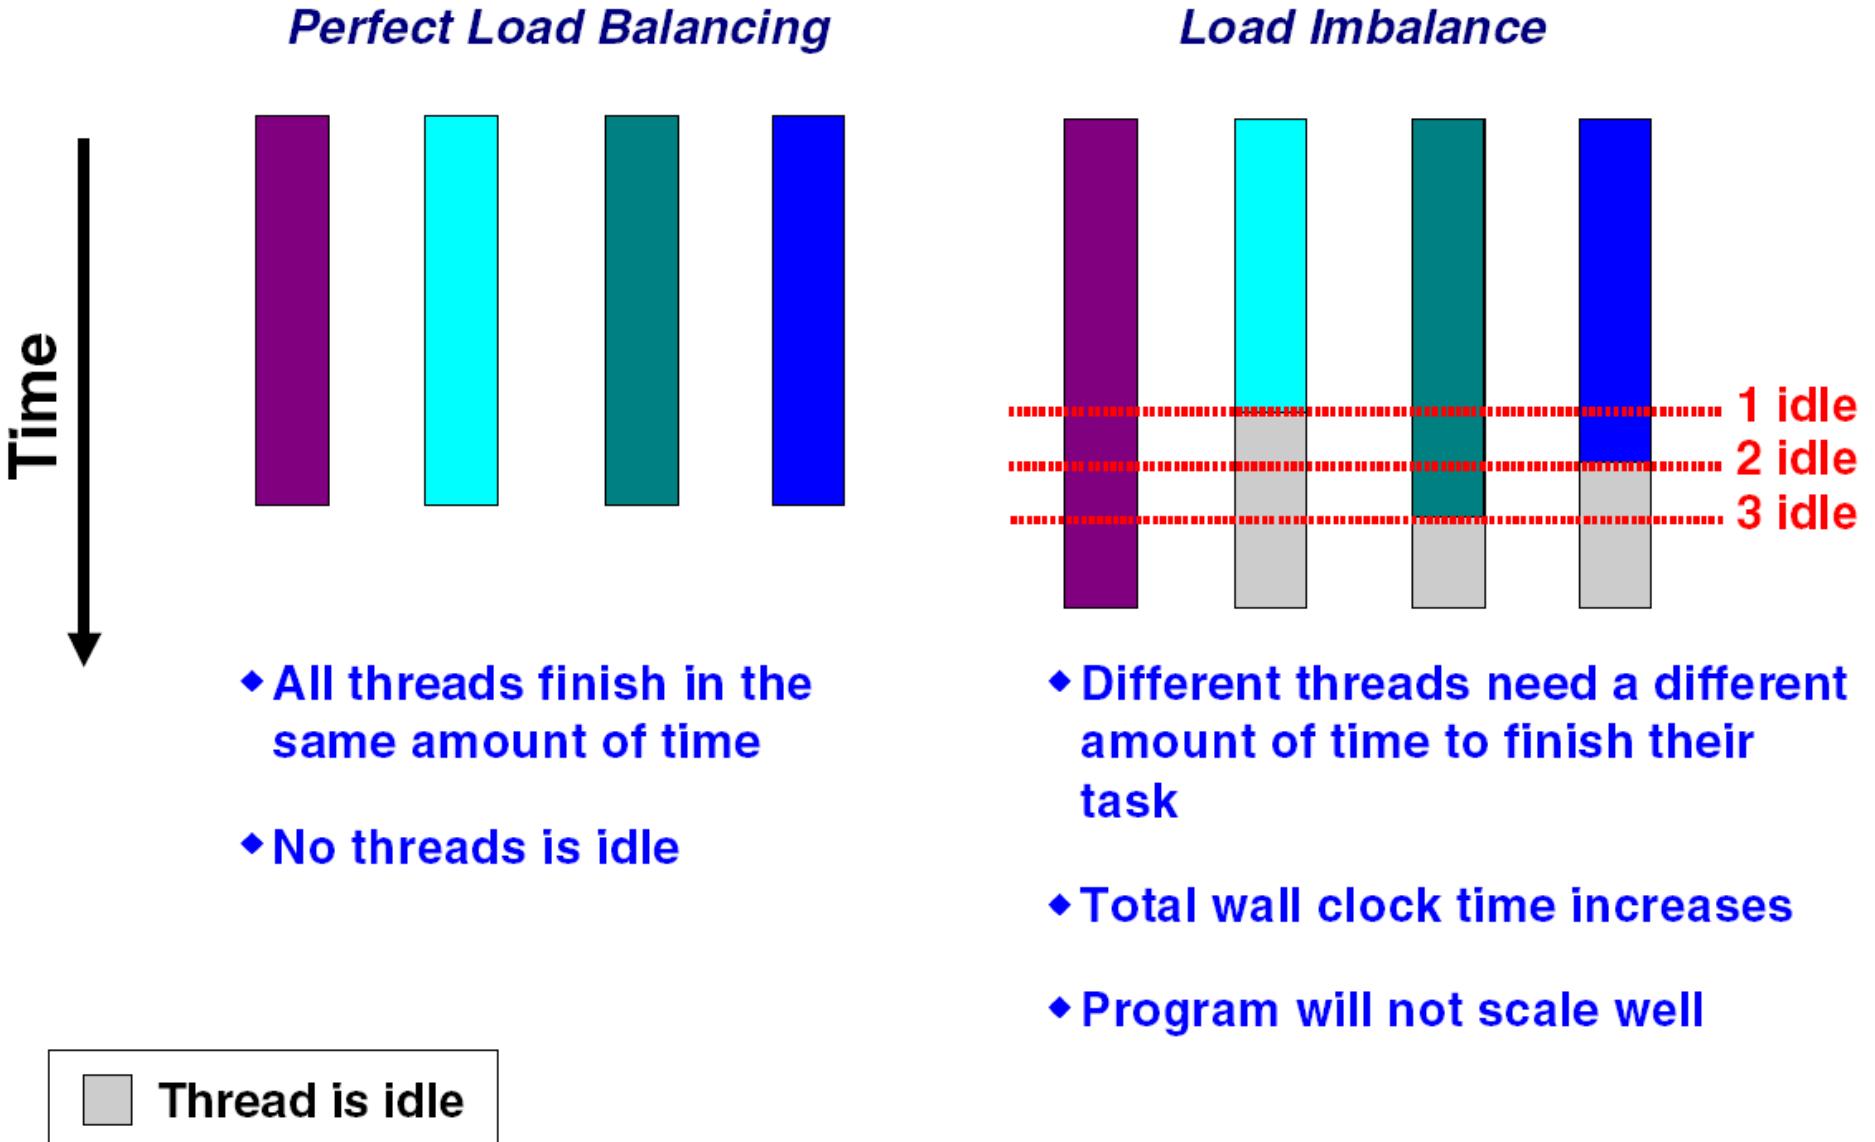
- ◆ Embarrassingly parallel: 4x faster
- ◆ Wallclock time is $\frac{1}{4}$ of serial wallclock time

Parallel - With communication

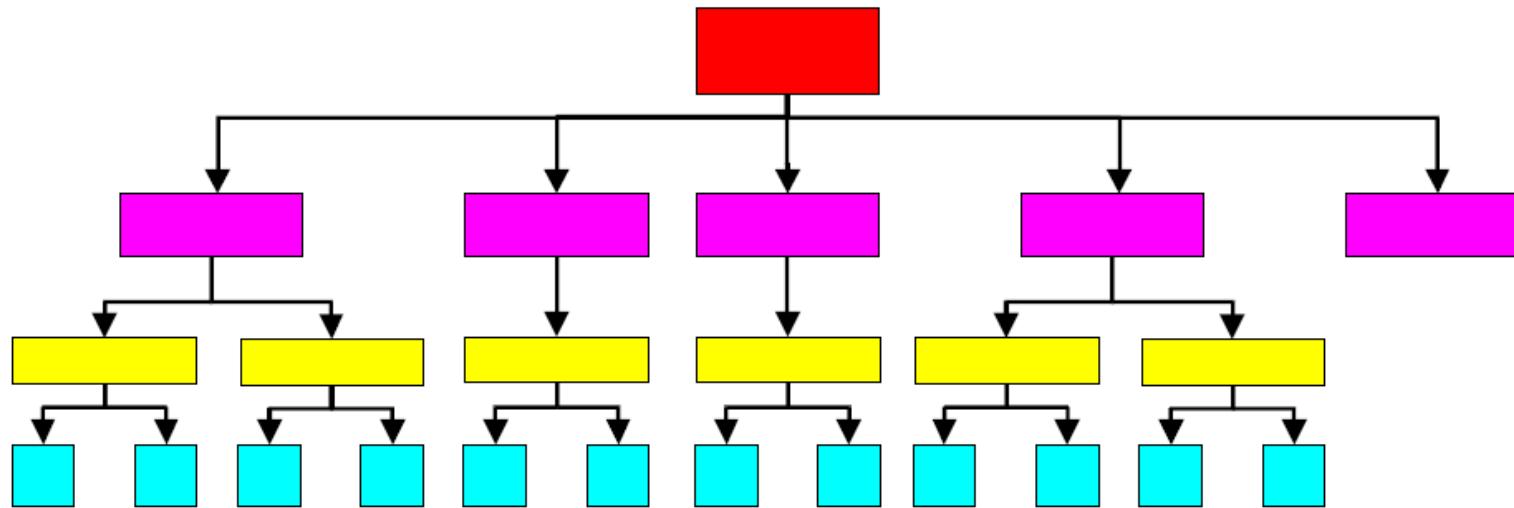


- ◆ Additional communication
- ◆ Less than 4x faster
- ◆ Consumes additional resources
- ◆ Wallclock time is more than $\frac{1}{4}$ of serial wallclock time
- ◆ Total CPU time increases

Load Balancing



Dilemma – Where to parallelize?



- ◆ **Parallelization at the highest (█) level:**

- ✓ *Low communication cost*
- ✓ *Limited to 5 processors only*
- ✓ *Potential load balancing issue*

- ◆ **Parallelization at the lowest (█) level:**

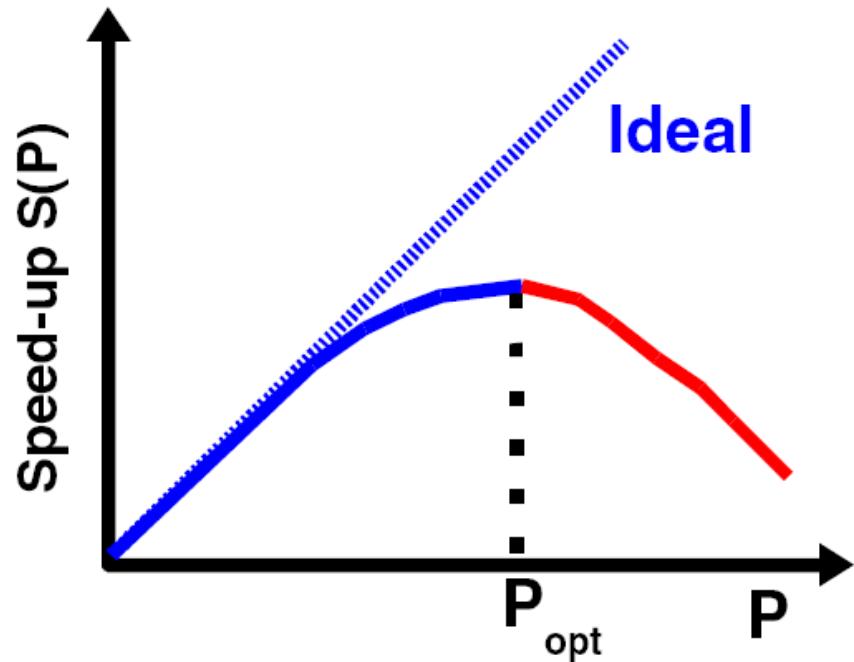
- ✓ *Higher communication cost*
- ✓ *Not limited to a certain number of processors*
- ✓ *Load balancing probably less of an issue*

Scalability

We distinguish ...

- ❑ ... how well a solution to some problem will work when the size of the problem increases.
 - ❑ typically associated with algorithmic complexity
- ❑ ... how well a parallel solution to some problem will work when the number of processing units (PUs) increases.
 - ❑ Strong scaling (speed-up) or weak scaling

Scalability – speed-up & efficiency



- ◆ Define the speed-up $S(P)$ as $S(P) := T(1)/T(P)$
- ◆ The efficiency $E(P)$ is defined as $E(P) := S(P)/P$
- ◆ In the ideal case, $S(P)=P$ and $E(P)=P/P=1=100\%$
- ◆ Unless the application is embarrassingly parallel, $S(P)$ will start to deviate from the ideal curve
- ◆ Past this point P_{opt} , the application will get less and less benefit from adding processors
- ◆ Note that both metrics give no information on the actual run-time
- ◆ As such, they can be dangerous to use

Amdahl's Law

Assume our program has a parallel fraction “f”

This implies the execution time $T(1) := f*T(1) + (1-f)*T(1)$

On P processors: $T(P) = (f/P)*T(1) + (1-f)*T(1)$

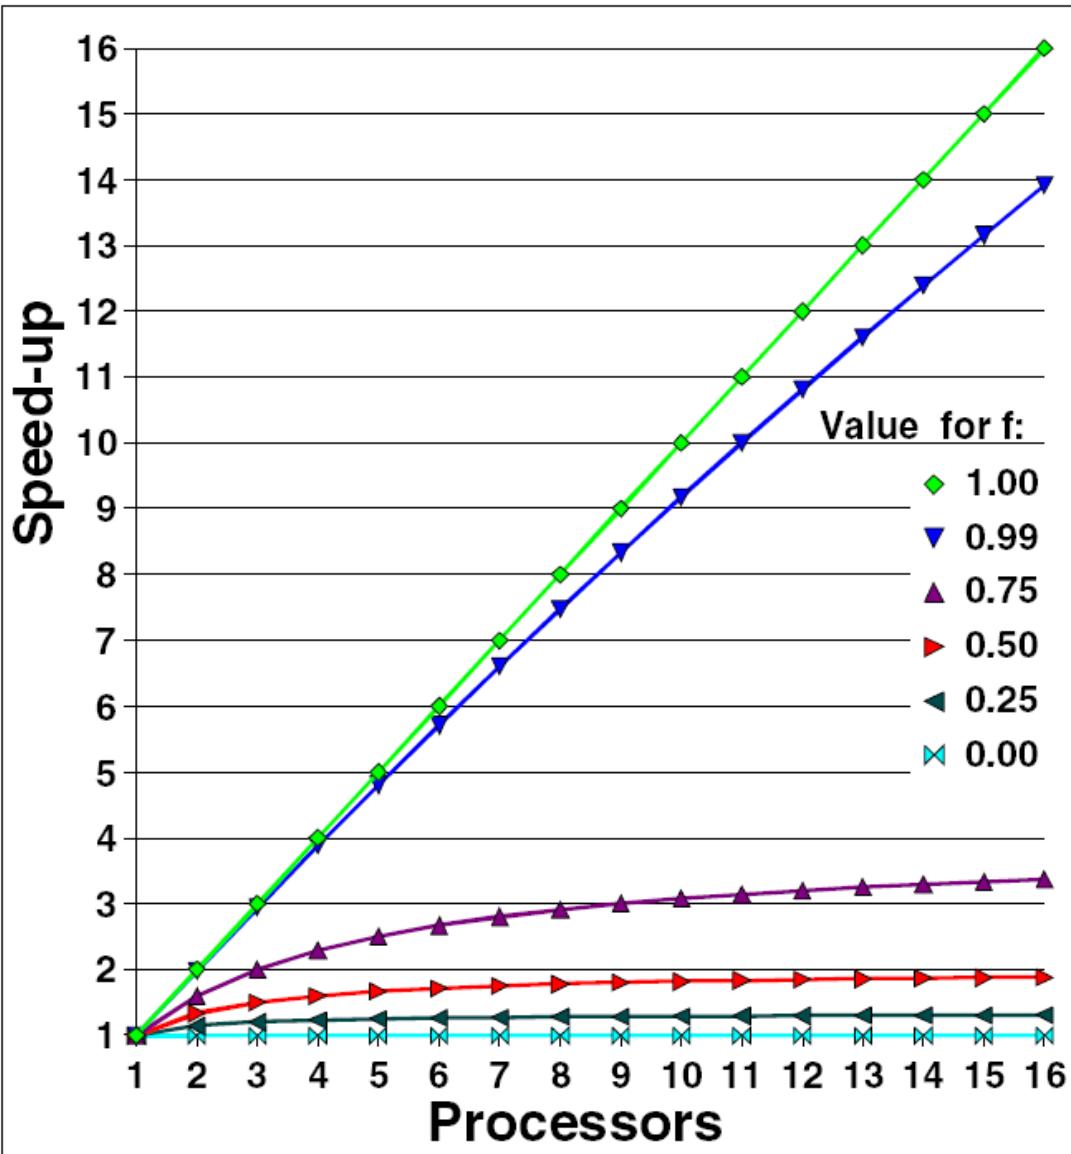
Amdahl's law:

$$S(P) := T(1) / T(P) = 1 / (f/P + 1-f)$$

Comments:

- ☞ This "law" describes the effect that the non-parallelizable part of a program has on scalability
- ☞ Note that the additional overhead caused by parallelization and speed-up because of cache effects are not taken into account

Amdahl's Law



- ◆ *It is easy to scale on a small number of processors*
- ◆ *Scalable performance however requires a high degree of parallelization i.e. f is very close to 1*
- ◆ *This implies that you need to parallelize that part of the code where the majority of the time is spent*

Amdahl's Law in Practice

We can estimate the parallel fraction “f”

Recall: $T(P) = (f/P)*T(1) + (1-f)*T(1)$

It is trivial to solve this equation for “f”:

$$f = (1 - T(P)/T(1))/(1 - (1/P))$$

Example:

$$\begin{aligned} T(1) &= 100 \text{ and } T(4)=37 \Rightarrow S(4) = T(1)/T(4) = 2.70 \\ f &= (1-37/100) / (1-(1/4)) = 0.63/0.75 = 0.84 = 84\% \end{aligned}$$

Estimated performance on 8 processors is then:

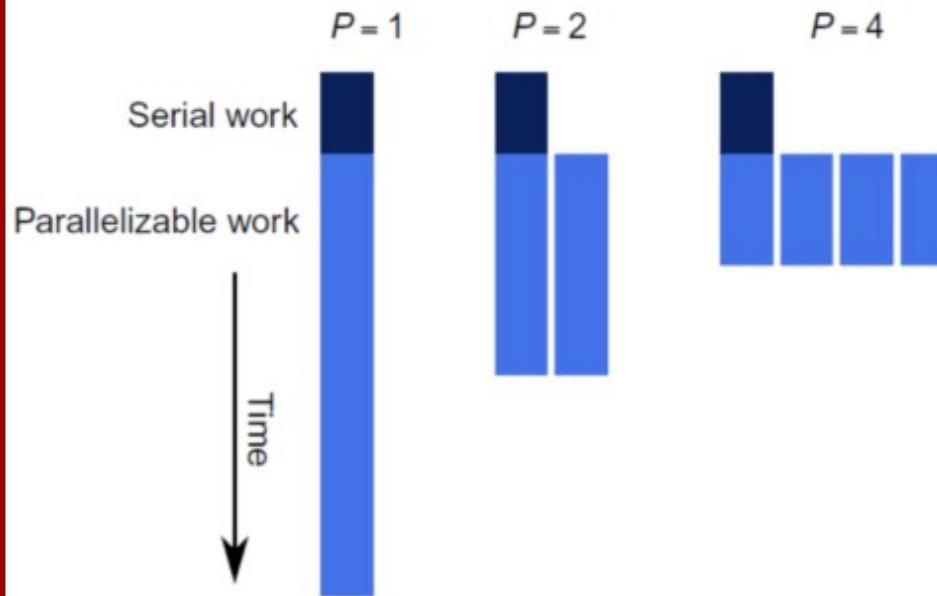
$$\begin{aligned} T(8) &= (0.84/8)*100 + (1-0.84)*100 = 26.5 \\ S(8) &= T/T(8) = 3.78 \end{aligned}$$

Scaling: strong vs. weak

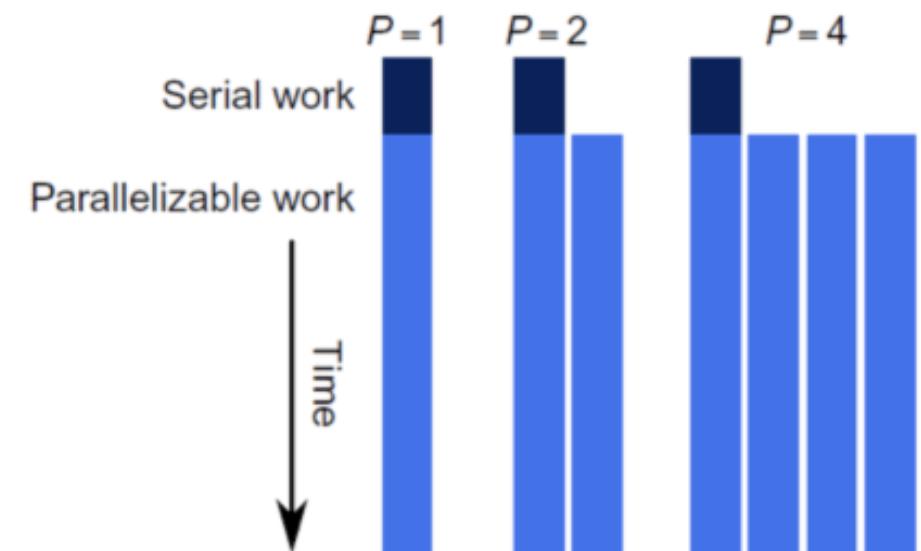
- ❑ How does the execution time go down for a fixed problem size by increasing the number of PUs?
 - ❑ Amdahl's law \Rightarrow speed-up, i.e. reduce time
 - ❑ also known as “strong scaling”
- ❑ How much can we increase the problem size by adding more PUs, keeping the execution time approx. constant?
 - ❑ Gustafson's law \Rightarrow scale-up, i.e. increase work
 - ❑ also known as “weak scaling”

Amdahl's vs Gustafson's law

Amdahl: fixed work



Gustafson: fixed work/PU



Amdahl's vs Gustafson's Law

- ❑ Amdahl's law
 - ❑ Theoretical performance of an application with a **fixed amount of parallel work** given a particular number of Processing Units (PUs)
- ❑ Gustafson's Law:
 - ❑ Theoretical performance of an application with a **fixed amount of parallel work per PU** given a particular number of PUs

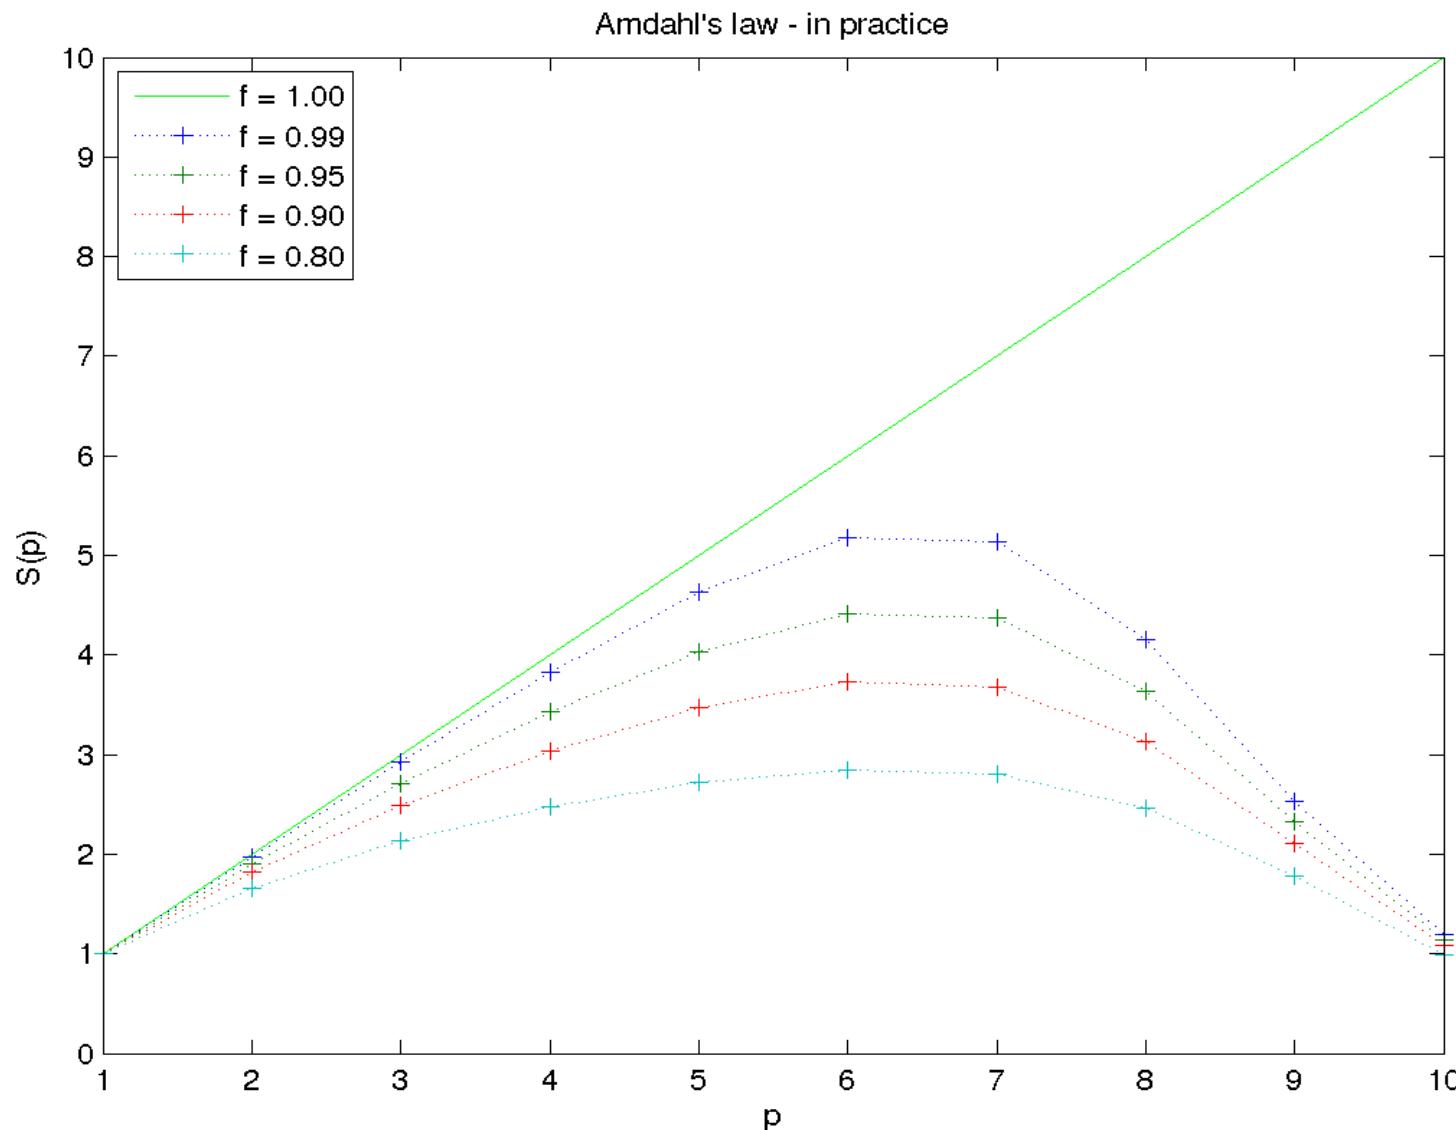
Code scalability in practice – I

- ❑ Although Amdahl and Gustafson provide theoretical upper bounds, eventually real data are necessary for analysis
- ❑ Inconsistencies in performance – especially on shared systems – often appear in singular runs
- ❑ Best practice: Monitor codes several times and average the results to filter out periods of heavy usage due to other users

Code scalability in practice – II

- ❑ Ideally, HPC codes would be able to scale to the theoretical limit, but ...
 - ❑ Never the case in reality
 - ❑ All codes eventually reach a real upper limit on speedup
 - ❑ At some point codes become “bound” to one or more limiting hardware factors (memory, network, I/O)

Code scalability in practice – III



What is Parallelization? - Summary

- ❑ Parallelization is simply another optimization technique to get your results sooner
- ❑ To this end, more than one processor is used to solve the problem
- ❑ The “Elapsed Time” (also called wallclock time) will come down, but total CPU time will probably go up
- ❑ The latter is a difference with serial optimization, where one makes better use of existing resources, i.e. the cost will come down

Parallel Programming Models

Parallel Programming Models

Two “classic” parallel programming models:

- ❑ Distributed memory
 - ❑ PVM (standardized)
 - ❑ MPI (de-facto standard, widely used)
 - ❑ <http://mpi-forum.org> or <http://open-mpi.org/>
- ❑ Shared memory
 - ❑ Pthreads (standardized)
 - ❑ C++11 threads
 - ❑ OpenMP (de-facto standard) <http://openmp.org/>
 - ❑ Automatic parallelization (depends on compiler)

Clusters,
SMPs

SMP
only

Parallel Programming Models

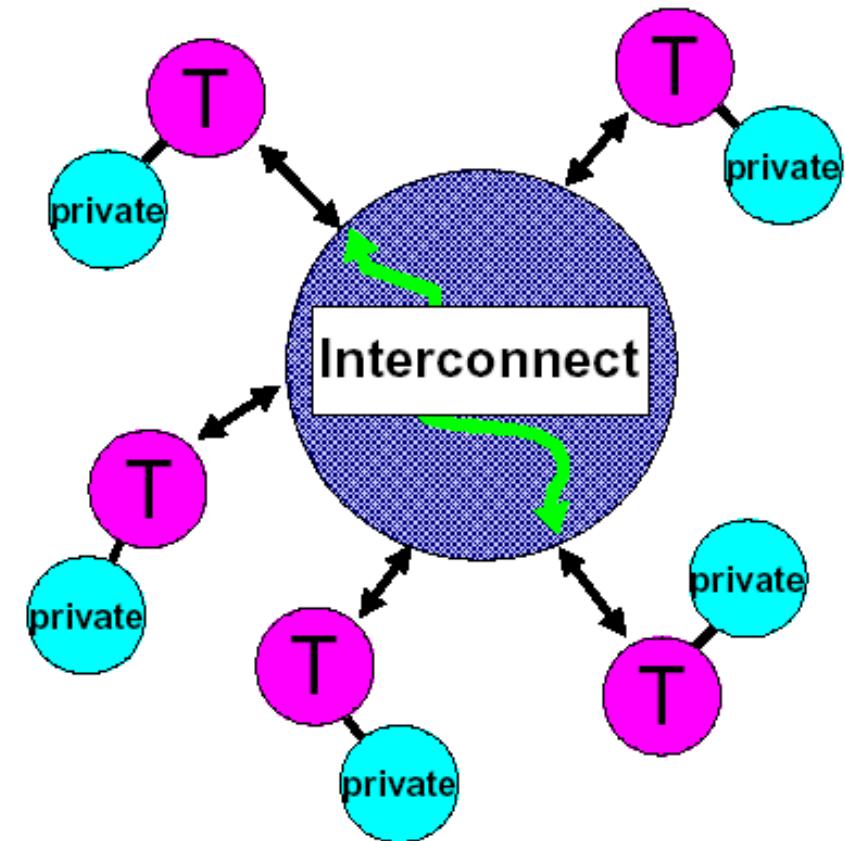
Other programming models

- ❑ PGAS (Partitioned Global Address Space):
 - ❑ UPC (Unified Parallel C)
 - ❑ Co-Array Fortran
- ❑ GPUs: massively parallel & shared memory
 - ❑ CUDA
 - ❑ OpenCL
 - ❑ Shader languages
 - ❑ OpenMP (target offloading)

Parallel Programming Models

Distributed memory
programming model,
e.g. MPI:

- ❑ all data is private to the threads
- ❑ data is shared by exchanging buffers
- ❑ explicit synchronization



Parallel Programming Models

MPI:

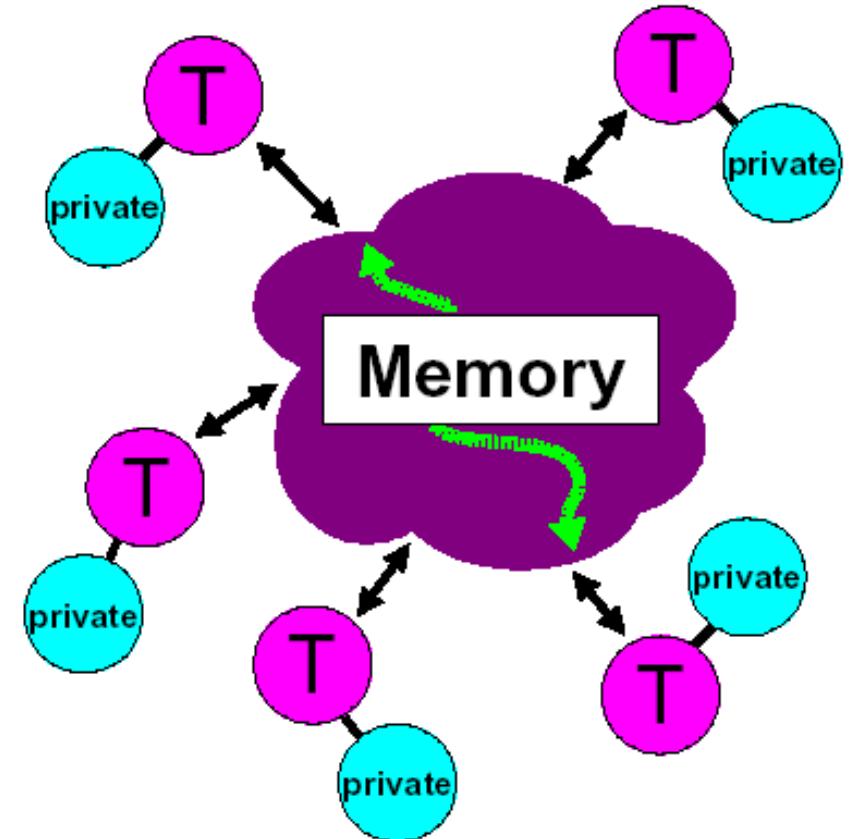
- ❑ An MPI application is a set of independent processes (aka ranks)
 - ❑ on different machines
 - ❑ on the same machine
- ❑ communication over the interconnect
 - ❑ network (network of workstations, cluster, grid)
 - ❑ memory (SMP)
- ❑ communication is under control of the programmer

Parallel Programming Models

Shared memory model, e.g.

OpenMP:

- ❑ all threads have access to the same global memory
- ❑ data transfer is transparent to the programmer
- ❑ synchronization is (mostly) implicit
- ❑ there is private data as well



Parallel Programming Models

OpenMP:

- ❑ needs an SMP
- ❑ but ... with current CPU designs, there is an SMP in every computer
 - ❑ multi-core CPUs (CMP)
 - ❑ chip multi-threading (CMT)
 - ❑ or a combination of both
 - ❑ or ... (whatever we'll see in the future)

MPI vs OpenMP

OpenMP version of “Hello world”:

```
#include <stdio.h>

int main(int argc, char *argv[] ) {
    #pragma omp parallel
    {
        printf("Hello world!\n");
    } /* end parallel */
    return(0);
}
```

MPI vs OpenMP

```
% gcc -o hello -fopenmp hello.c
```

```
% ./hello
```

```
Hello world!
```

```
% OMP_NUM_THREADS=2 ./hello
```

```
Hello world!
```

```
Hello world!
```

```
% OMP_NUM_THREADS=8 ./hello
```

```
Hello world!
```

```
Hello world!
```

```
Hello world!
```

```
Hello world!
```

This behaviour is
implementation dependent!

no. of threads: `OMP_NUM_THREADS`

MPI vs OpenMP

MPI version of “Hello world”:

```
#include <stdio.h>
#include <stdlib.h>
#include "mpi.h"

int main(int argc, char **argv) {
    int myrank, p;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
    MPI_Comm_size(MPI_COMM_WORLD, &p);
    printf("Hello world from %d!\n", myrank);
    MPI_Finalize();
    return 0;
}
```

MPI vs OpenMP

MPI version: compile and run

```
$ module load mpi  
$ mpicc -o hello_mpi hello_mpi.c
```

```
$ ./hello_mpi  
Hello world from 0!
```

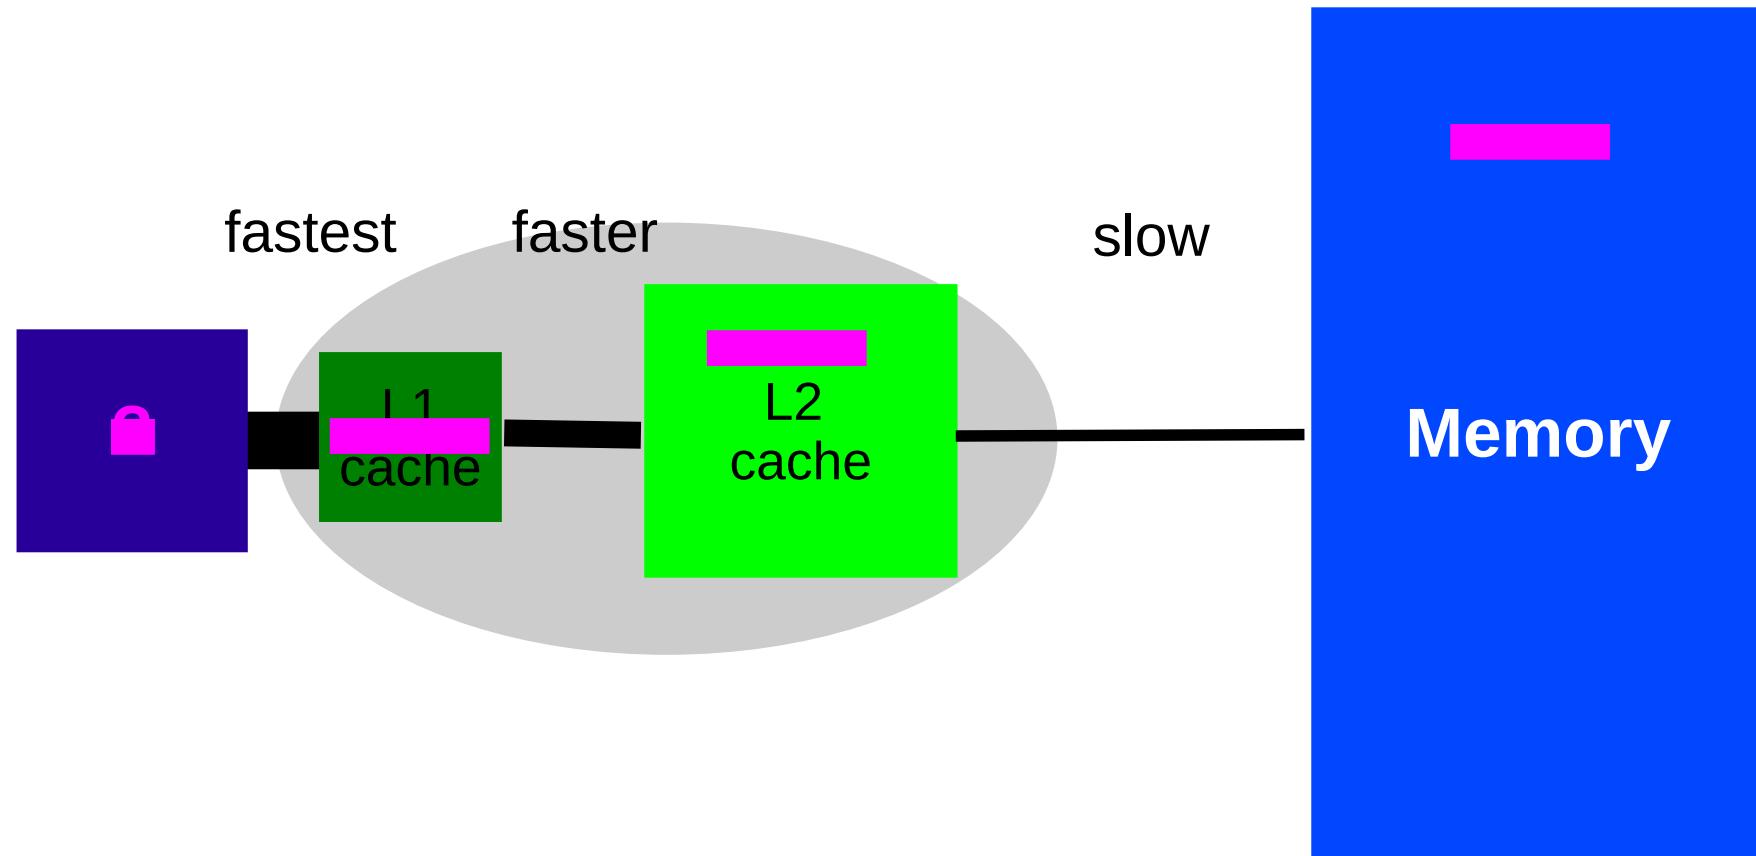
```
$ mpirun -np 4 ./hello_mpi  
Hello world from 1!  
Hello world from 3!  
Hello world from 0!  
Hello world from 2!
```

Automatic Parallelization

- ❑ Some compilers are capable of generating parallel code for loops that can be safely executed in parallel.
- ❑ This is always loop based!
- ❑ Compilers:
 - ❑ Intel compilers: -parallel
 - ❑ Oracle Studio compilers: -xautopar
 - ❑ GCC: -floop-parallelize-all -ftree-parallelize-loops=N
 - ❑ limited, sets no. of threads (N) at compile time!
- ❑ For more information see the manual pages

Caches revisited

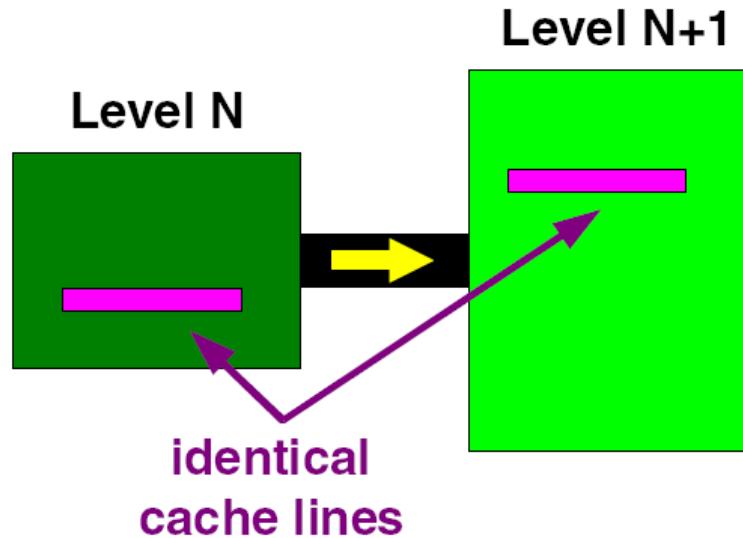
Typical cache based system



How are caches organized?

- ❑ Caches contain partial images of memory
- ❑ If data gets modified, the state of that data, i.e. the whole cache line, changes
- ❑ This has to be made known to the system
- ❑ There are two common approaches:
 - ❑ Write-through
 - ❑ Write-back

Write-through cache

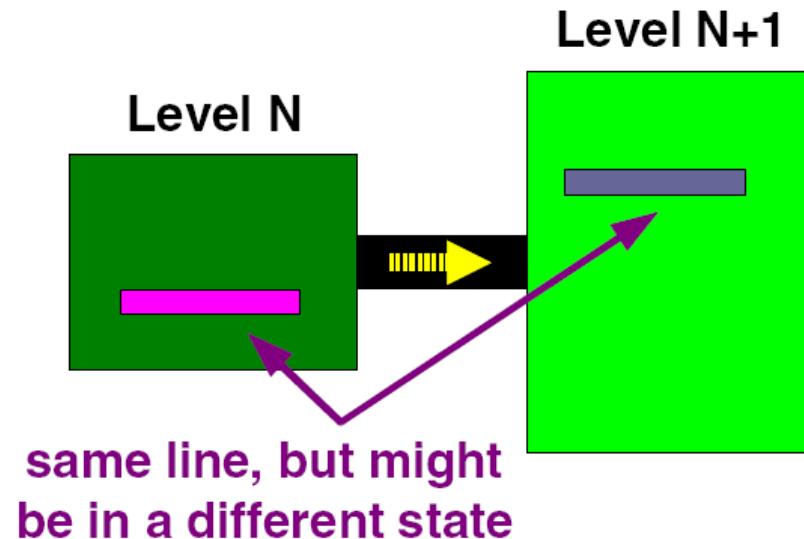


Notes:

- simple to implement
- easy to find the right copy
- can result in waste of bandwidth

- ❑ Always flushes a modified cache line back to a higher level in the memory hierarchy
 - ❑ e.g. from L1-cache to main memory
- ❑ This assures, that the system always knows where to access the correct cache line

Write-back cache



Notes:

- minimizes cache traffic
- need to keep track of status
- this mechanism is called '*cache coherency*'

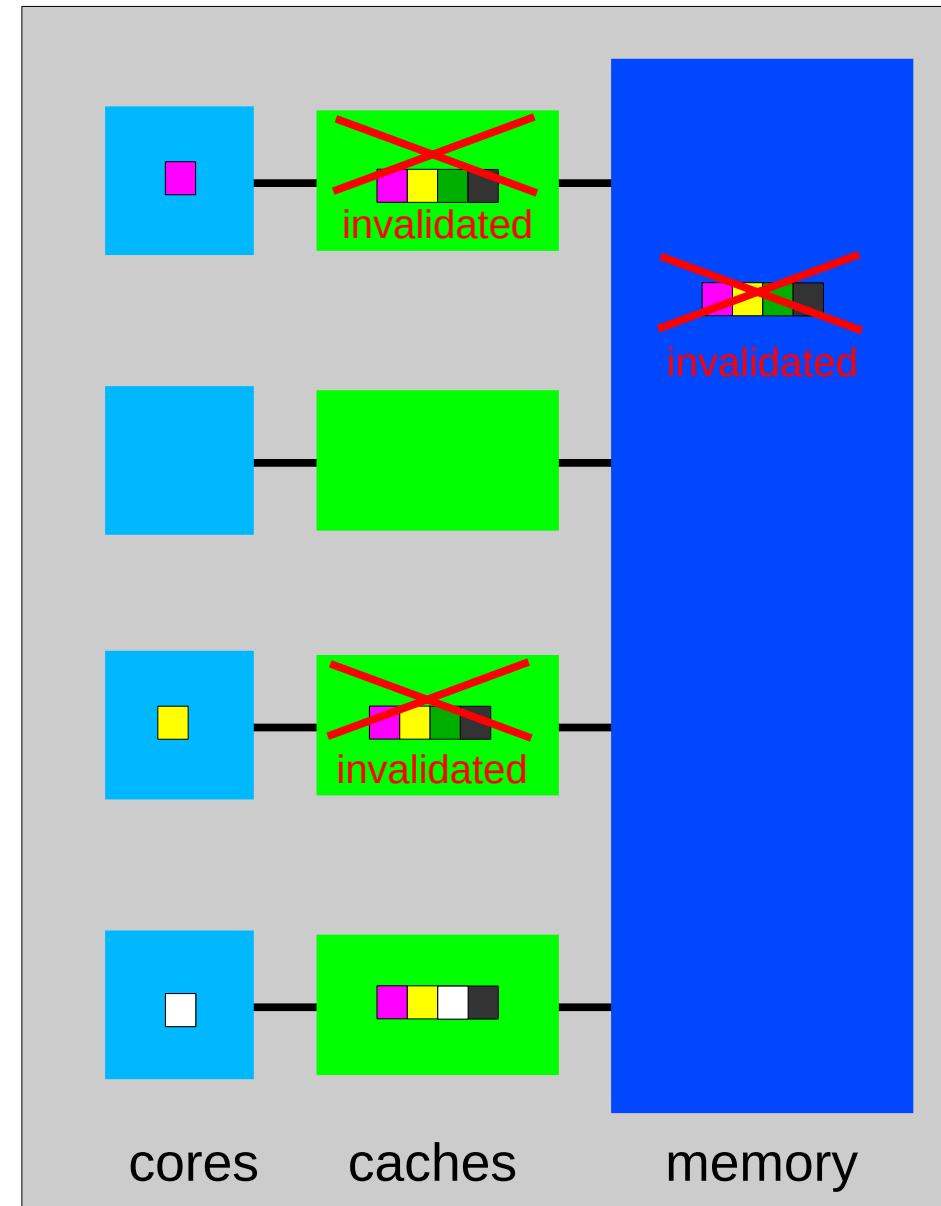
- ❑ Write a modified cache line back only if needed
 - ❑ capacity issues
 - ❑ another cache line maps onto this line
 - ❑ another CPU might need this cache line

Caches in MP/multi-core systems

- ❑ A cache line always starts in memory
- ❑ Over time multiple copies may exist

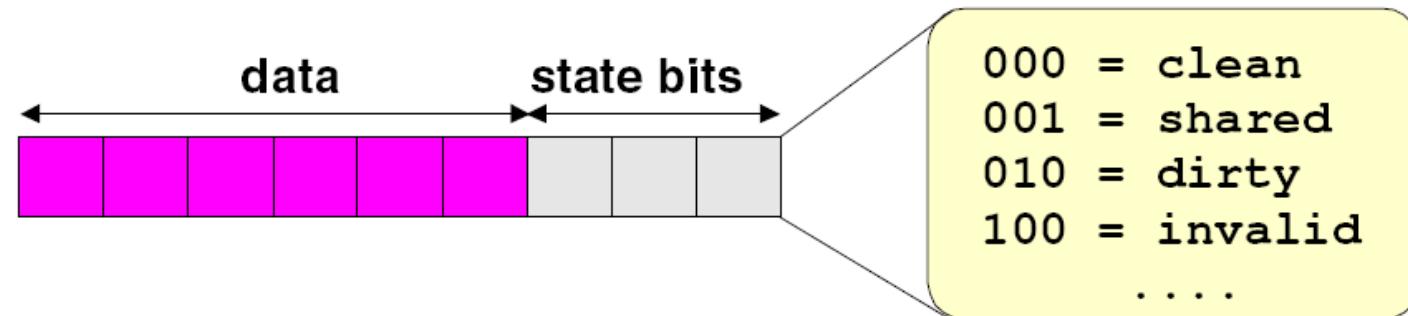
Cache coherency ('cc')

- ❑ tracks changes in copies
- ❑ assures correct cache line is used
- ❑ many implementations
- ❑ hardware support to be efficient



Cache coherency ('cc')

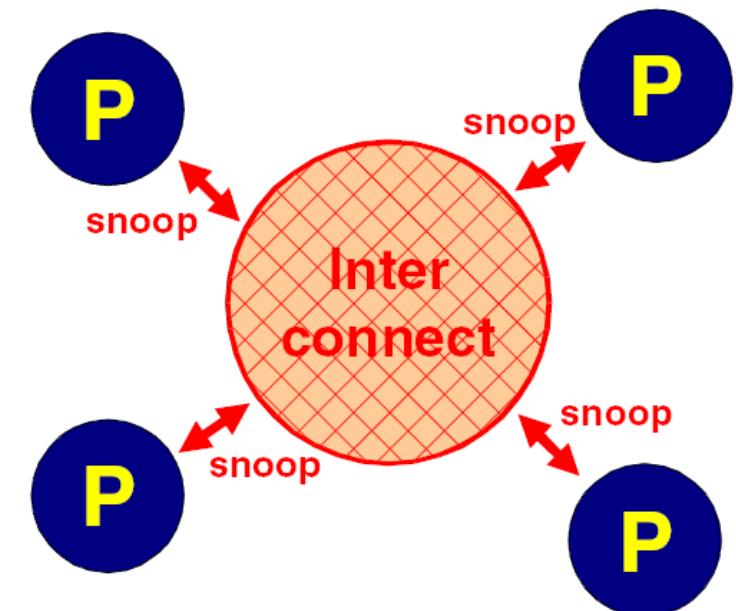
- ❑ Needed in write-back cache systems
- ❑ Keeps track of the status of all cache lines
- ❑ State information:



- ❑ Signals (“coherence traffic”) are used to update the state bits of the cache lines
- ❑ This allows to build efficient SMP systems

Snoopy based cache coherence

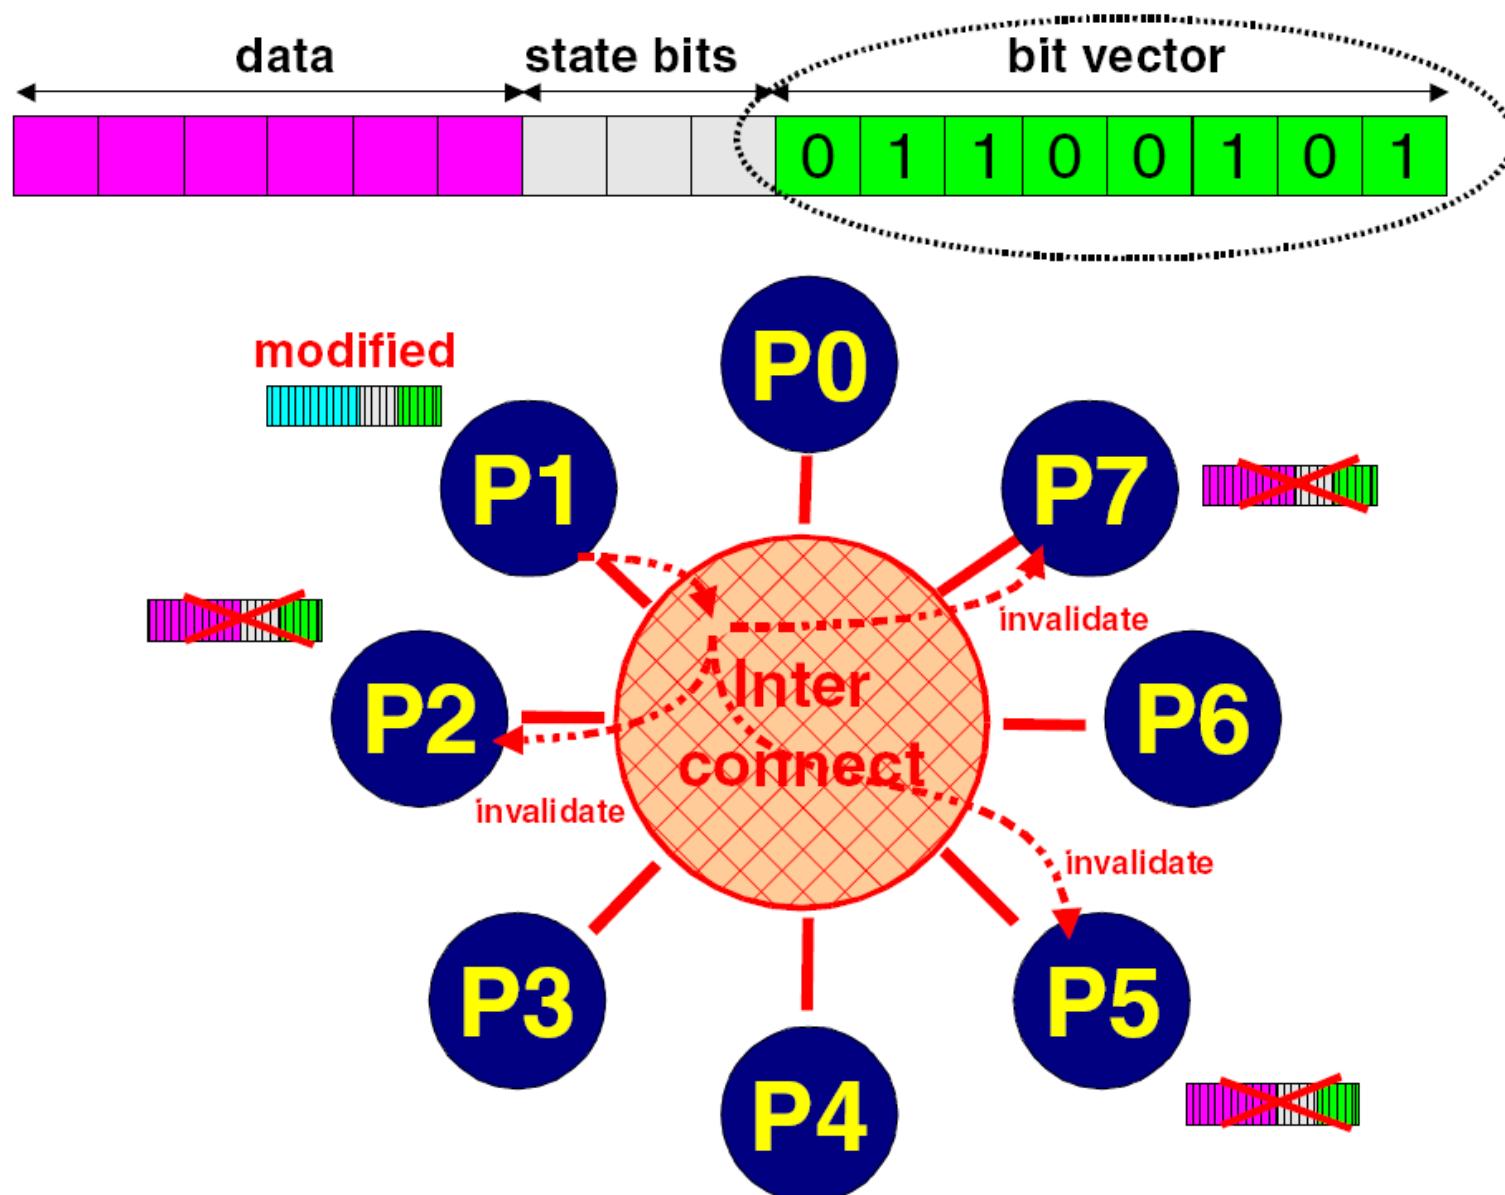
- ❑ Also known as “broadcast cache coherence”
 - ❑ all addresses are sent to all CPUs
 - ❑ result takes only a few cycles
- ❑ Advantages:
 - ❑ low latency
 - ❑ fast cache-to-cache transfer
- ❑ Disadvantages:
 - ❑ data bandwidth limited by snoop bandwidth
 - ❑ difficult to scale to many CPUs



Directory based cache coherence

- ❑ Also known as SSM (Scalable Shared Memory)
- ❑ point-to-point protocol, i.e. a directory keeps track which CPUs are involved with a particular cache line
- ❑ requests are sent to the involved CPUs only
- ❑ Advantages:
 - ❑ larger bandwidth & scalable to many CPUs
- ❑ Disadvantages:
 - ❑ longer & non-uniform latency
 - ❑ slower cache-to-cache transfer
 - ❑ need to store the directory entries

SSM example:

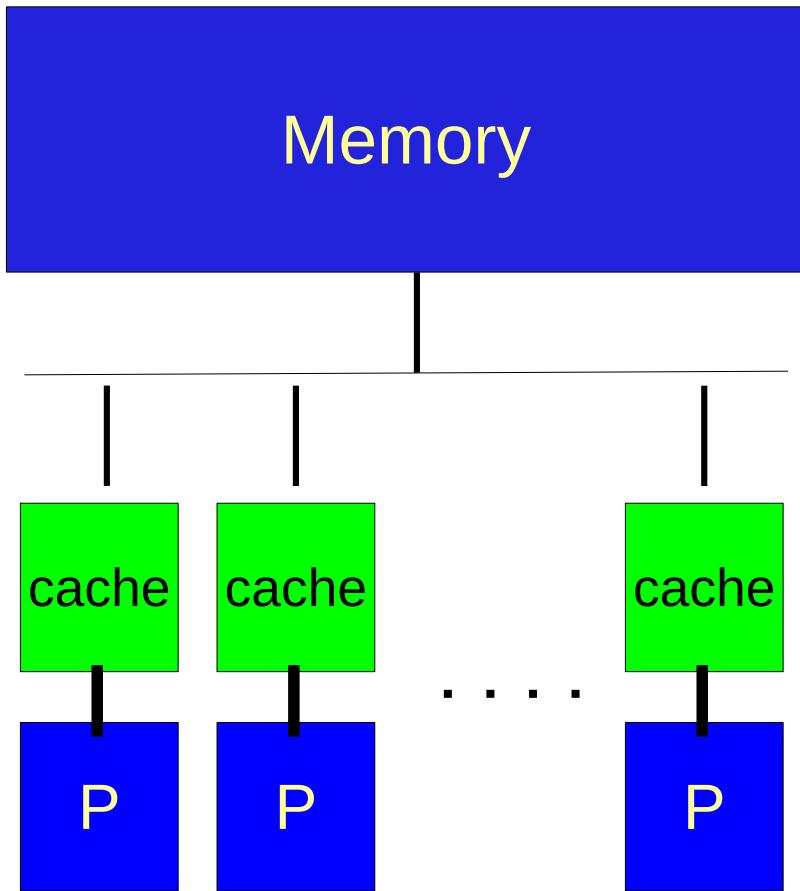


Parallel Architectures

Parallel Architectures

- ❑ It is difficult to label systems:
 - ❑ most systems share some characteristics, but not all
 - ❑ the variety of systems is increasing
- ❑ In the (“*historical*”) overview presented here, systems are labelled based on main memory:
 - ❑ Shared or Distributed:
 - ❑ can all CPUs access all memory, or only a subset?
 - ❑ Memory access times
 - ❑ uniform vs non-uniform

Uniform Memory Access (UMA)



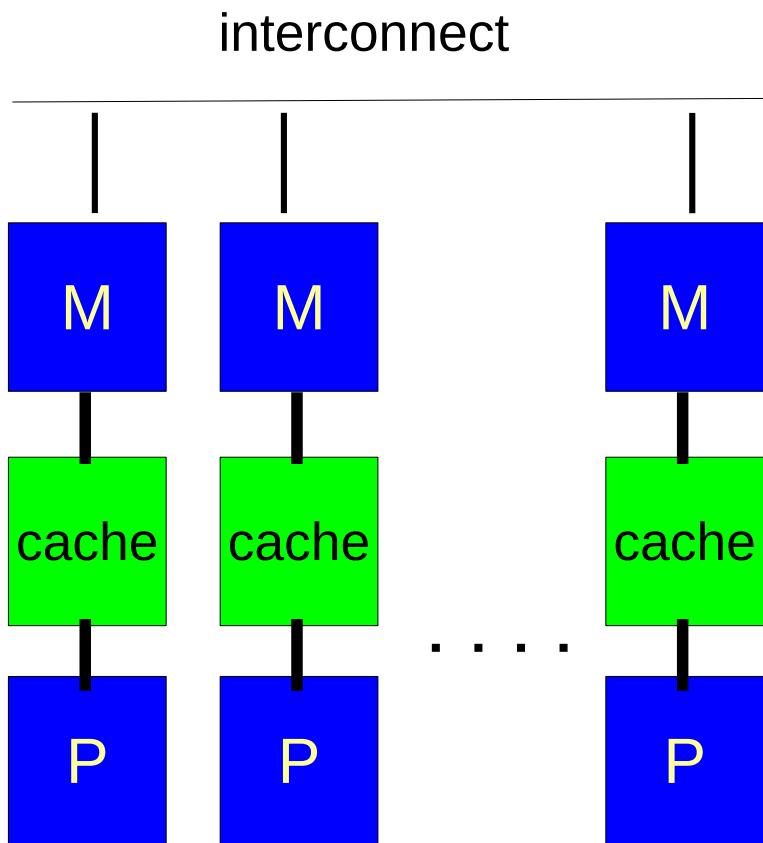
Memory access:

- ❑ uniform for all P

Example systems:

- ❑ single-socket multi-core CPU (e.g. your laptop)

Non-Uniform Memory Access (NUMA)



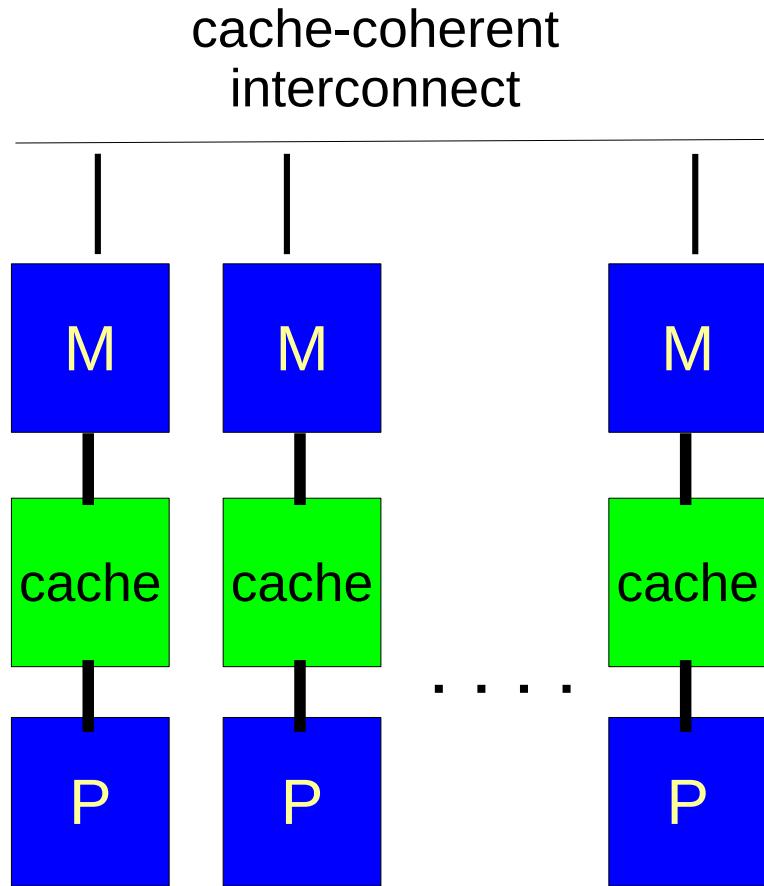
Memory access:

- ❑ non-uniform across P

Example systems:

- ❑ cluster of nodes,
connected by a (fast)
network

cc-NUMA



Memory access:

- ❑ non-uniform across P
- ❑ ... but cache-coherency on the interconnect

Example systems:

- ❑ (almost) all multi-socket x64 servers

Multi-core – everywhere!

Welcome to a “threaded” world

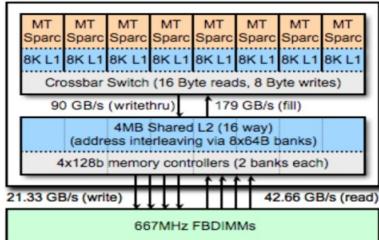
The first multi-core chips

- ❑ 2004 – multi-core arrives:
 - ❑ IBM POWER5
 - ❑ Sun UltraSPARC-IV
- ❑ 2005 is the year of the x86 dual-core CPUs:
 - ❑ AMD Opteron “Denmark” (August 2005)
 - ❑ Intel Xeon “Paxville DP” (October 2005)
- ❑ 2008/2009: quad-cores
 - ❑ AMD Opteron 'Barcelona'
 - ❑ Intel Xeon 'Nehalem'

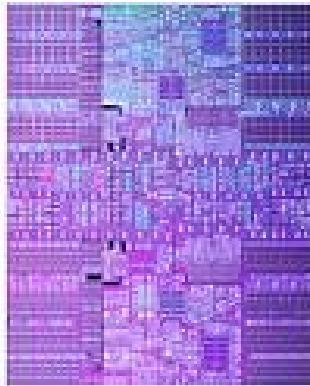
2009's Multi-cores

99% of Top500 Systems Are Based on Multi-core

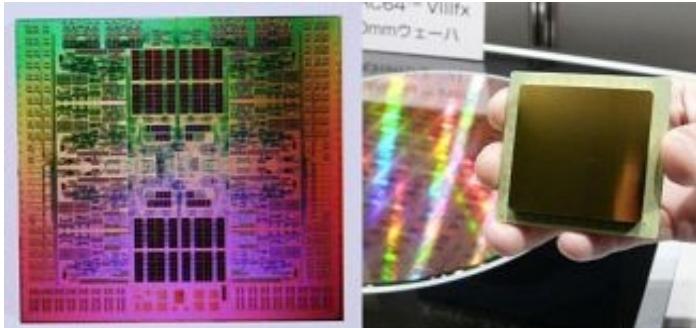
282 use Quad-Core
204 use Dual-Core
3 use Nona-core



Sun Niagara2 (8 cores)



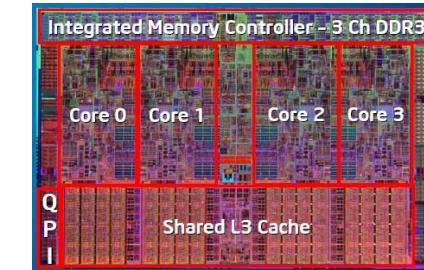
IBM Power 7 (8 cores)



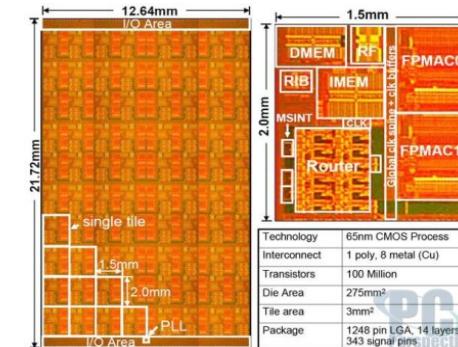
Fujitsu Venus (8 cores)



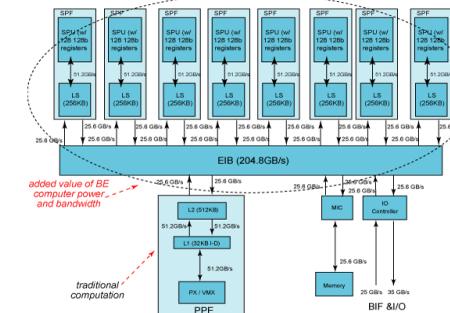
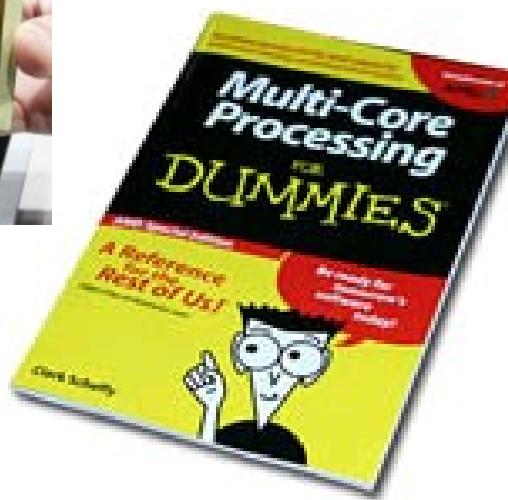
AMD Istanbul
(6 cores)



Intel Nehalem (4 cores)

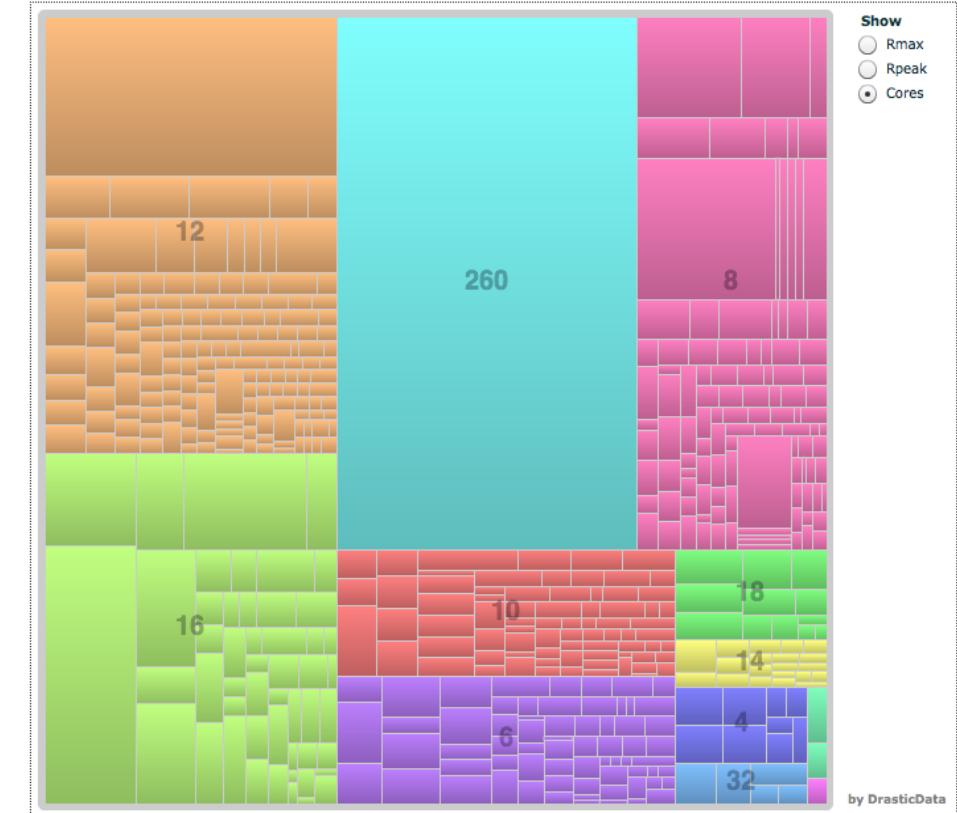
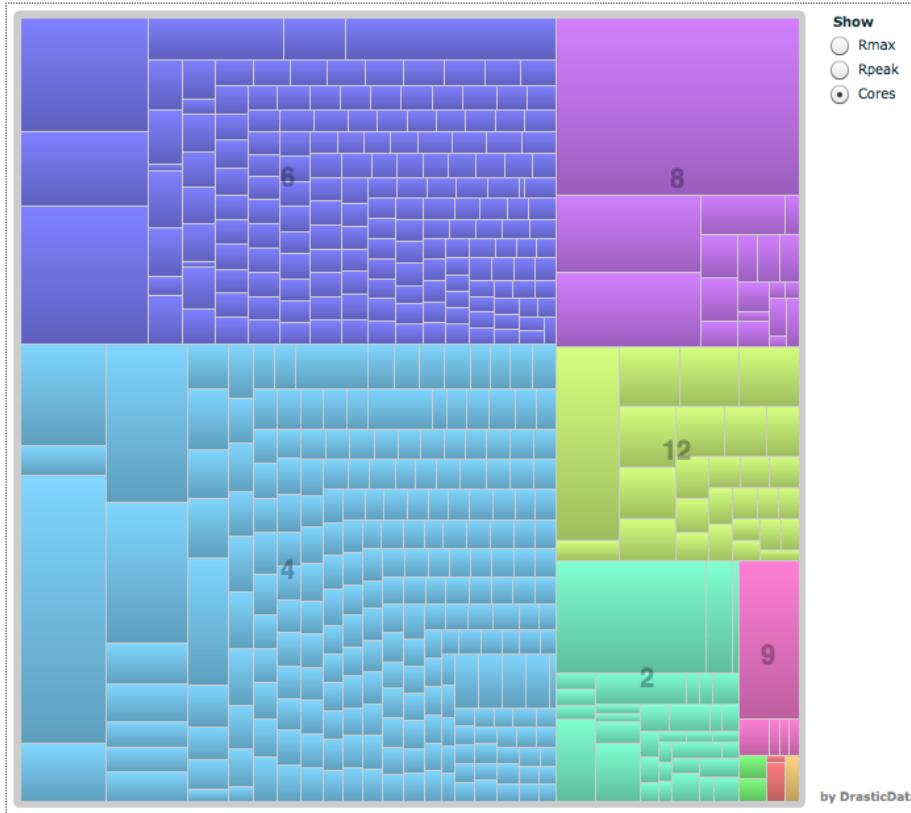


Intel Polaris [experimental]
(80 cores)



IBM Cell
(9 cores)

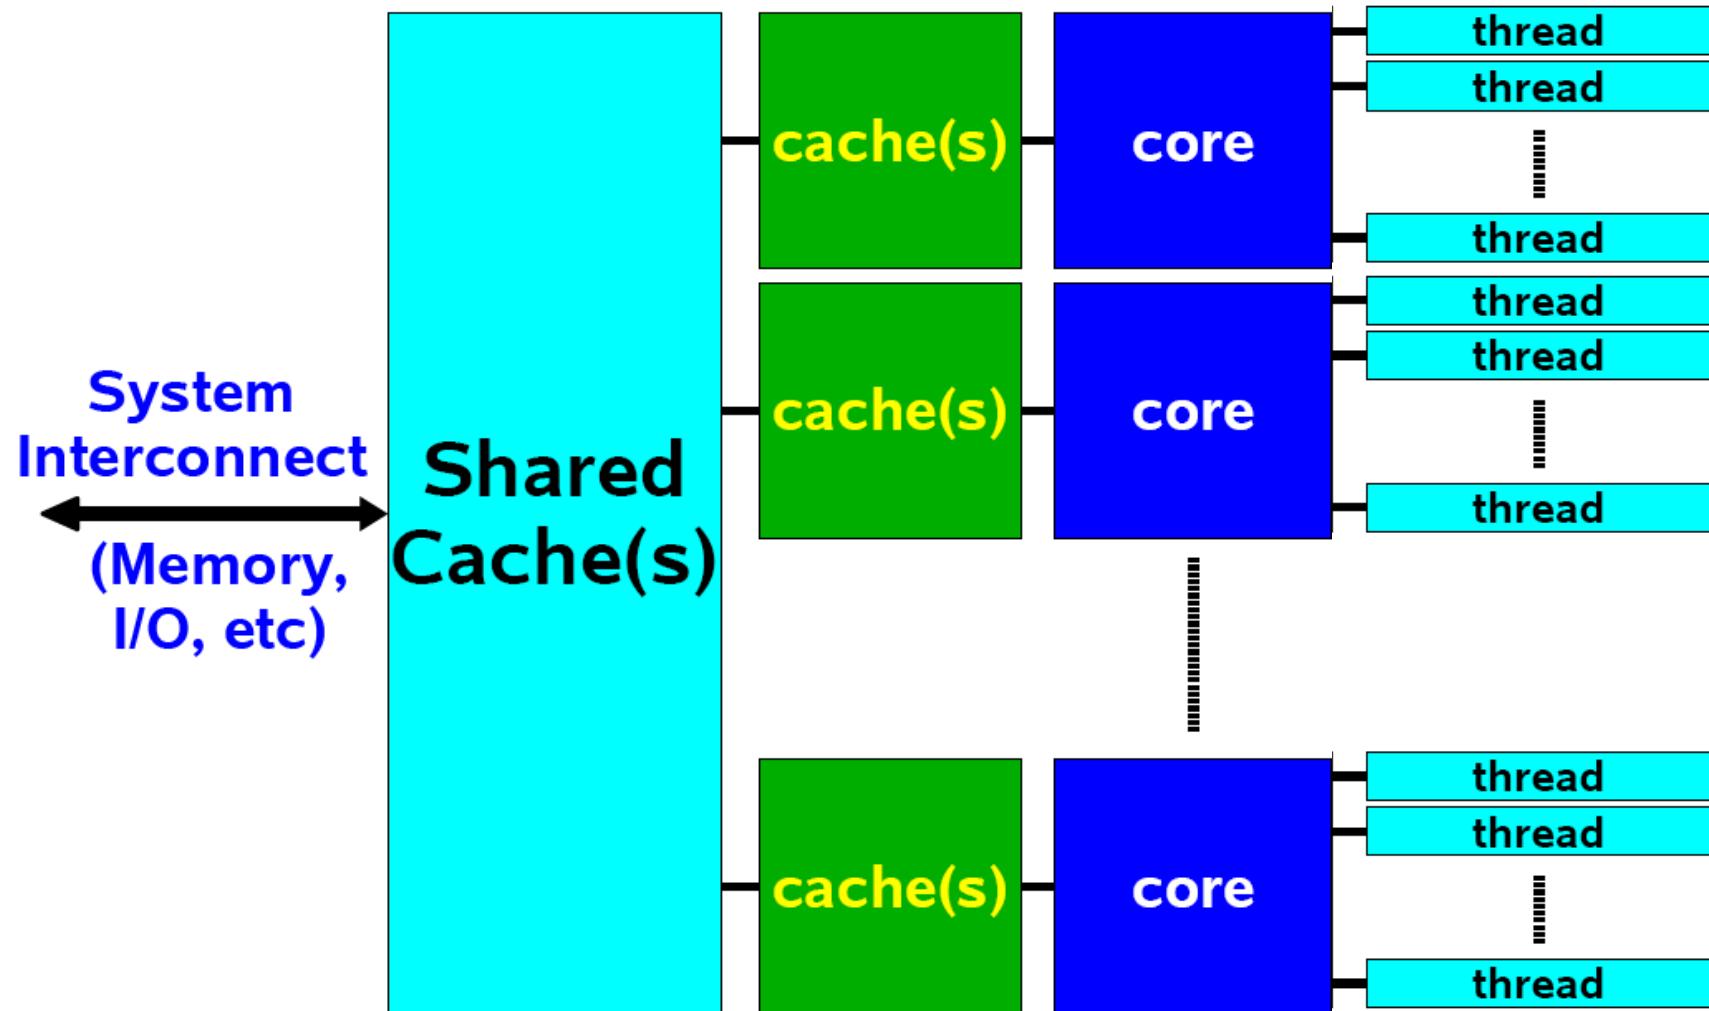
TOP500: multi-cores 2011 and 2016



What is a multi-core chip?

- ❑ A “core” is not well-defined – let us assume it covers the processing units and the L1 caches (a very simplified CPU).
- ❑ Different implementations are possible – and available (examples follow), e.g. multi-threaded cores
- ❑ Cache hierarchy of private and shared caches
- ❑ For software developers it matters that there is parallelism in the hardware, they can take advantage of

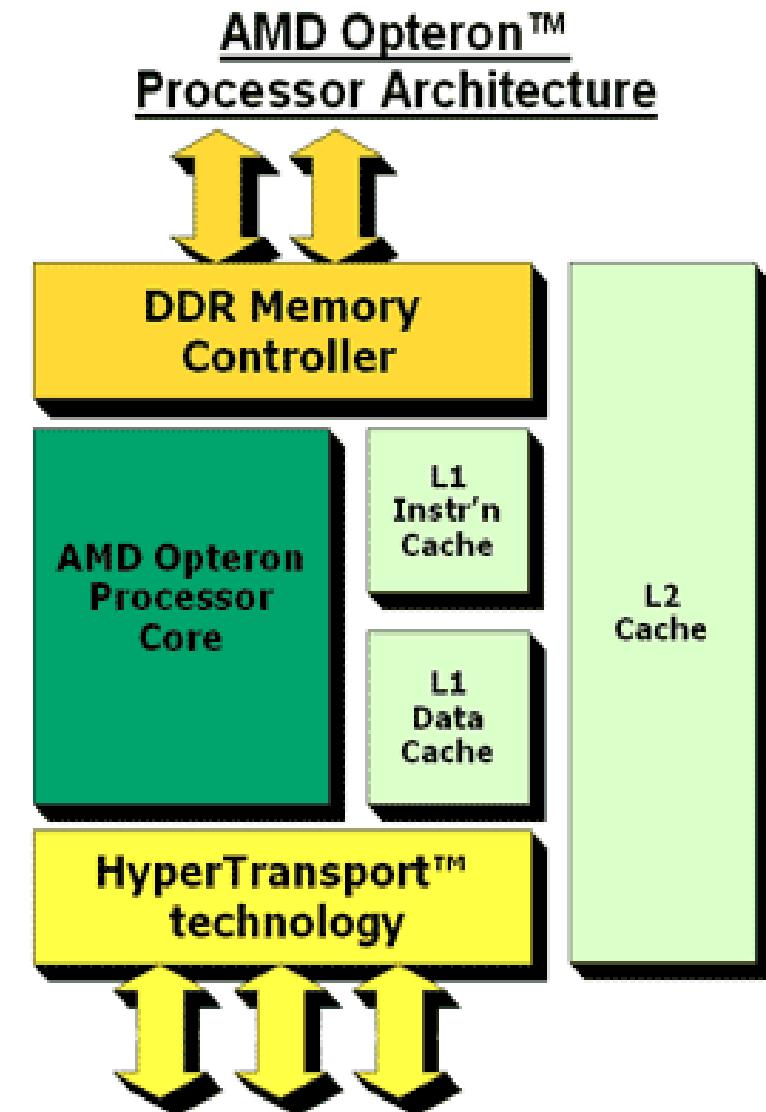
A generic multi-core design



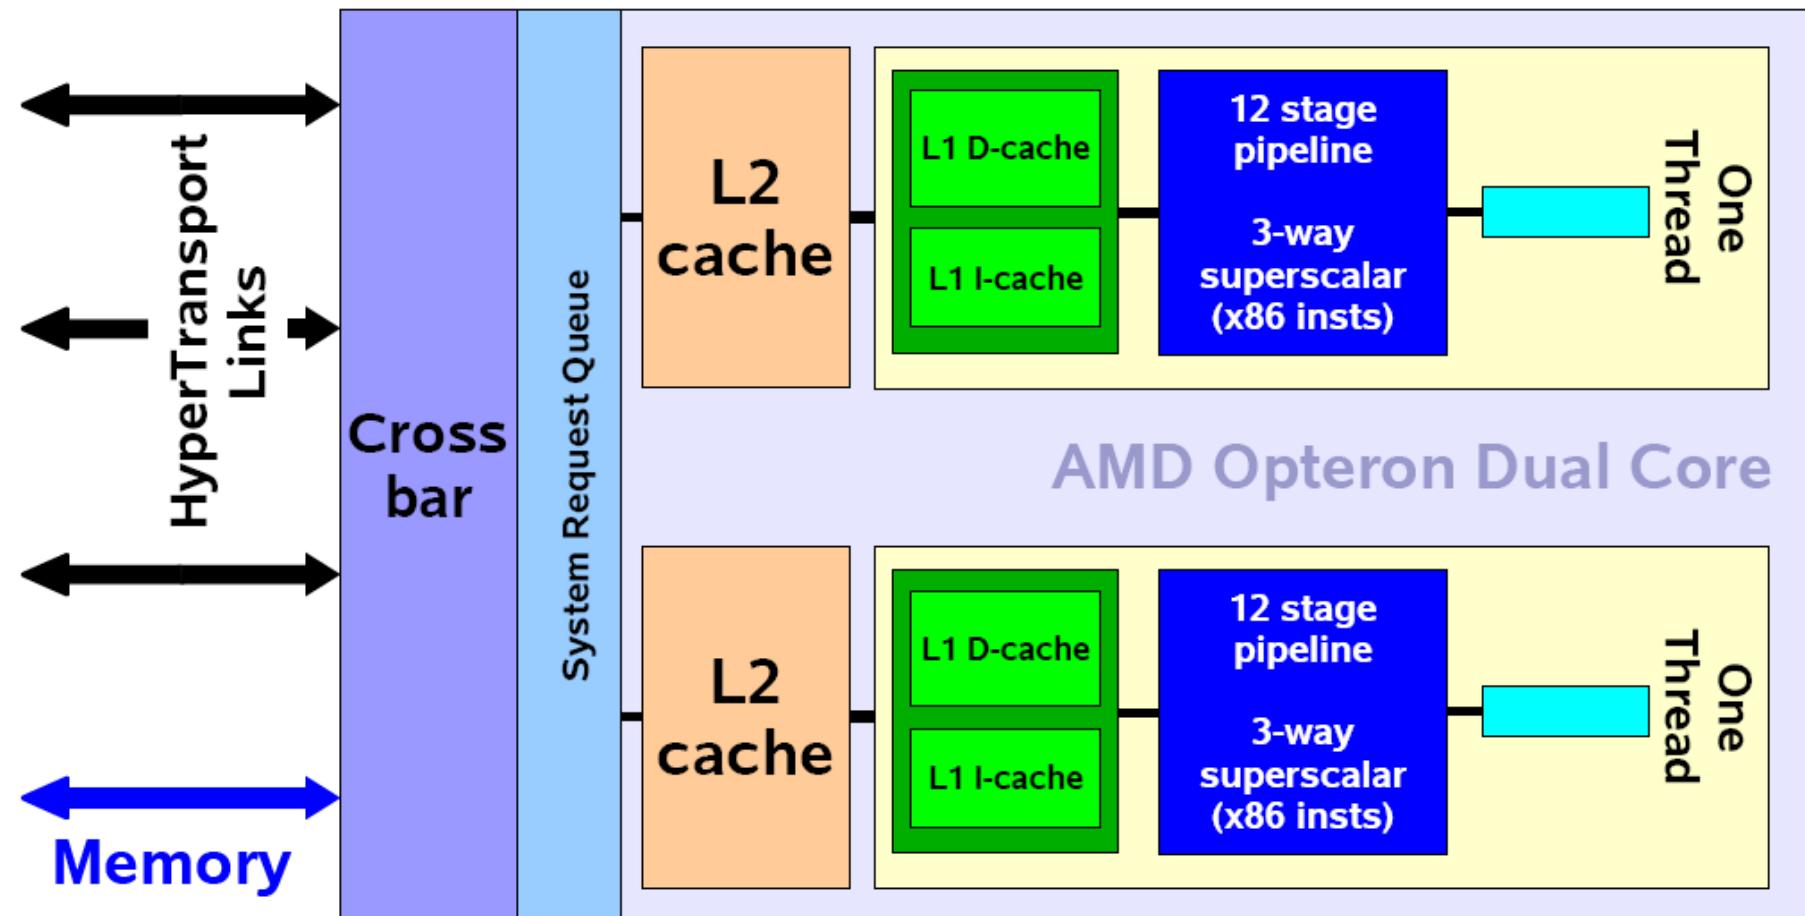
The AMD Opteron – single core

On-chip:

- ❑ Memory controller
- ❑ L2 cache
- ❑ 3 fast HyperTransport links: 6.4 GB/s per link

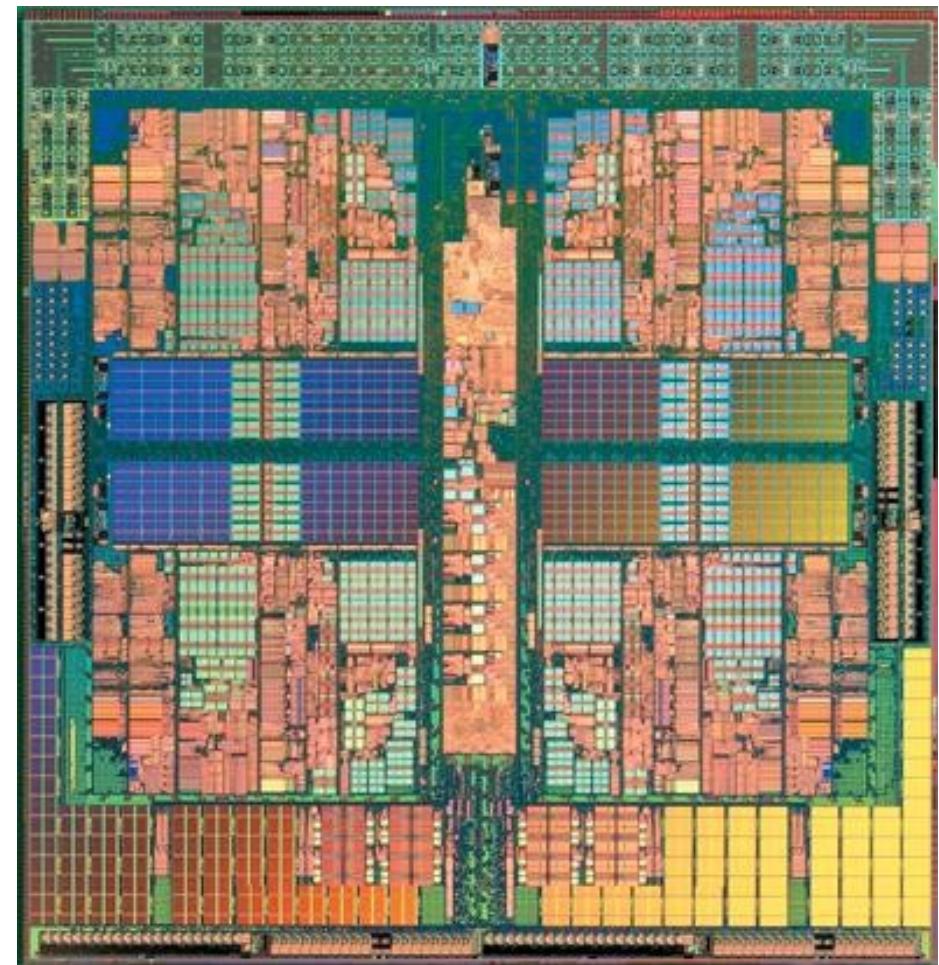
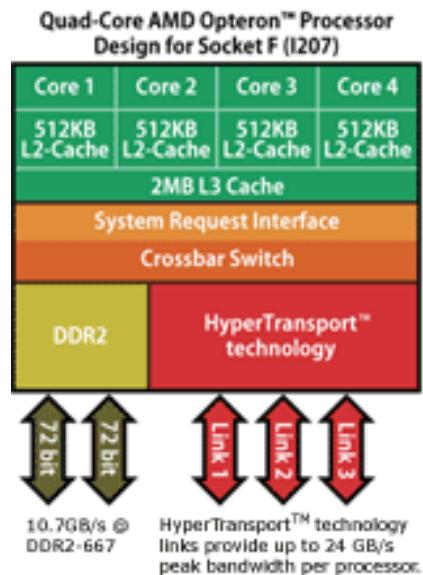


AMD Opteron - dual-core

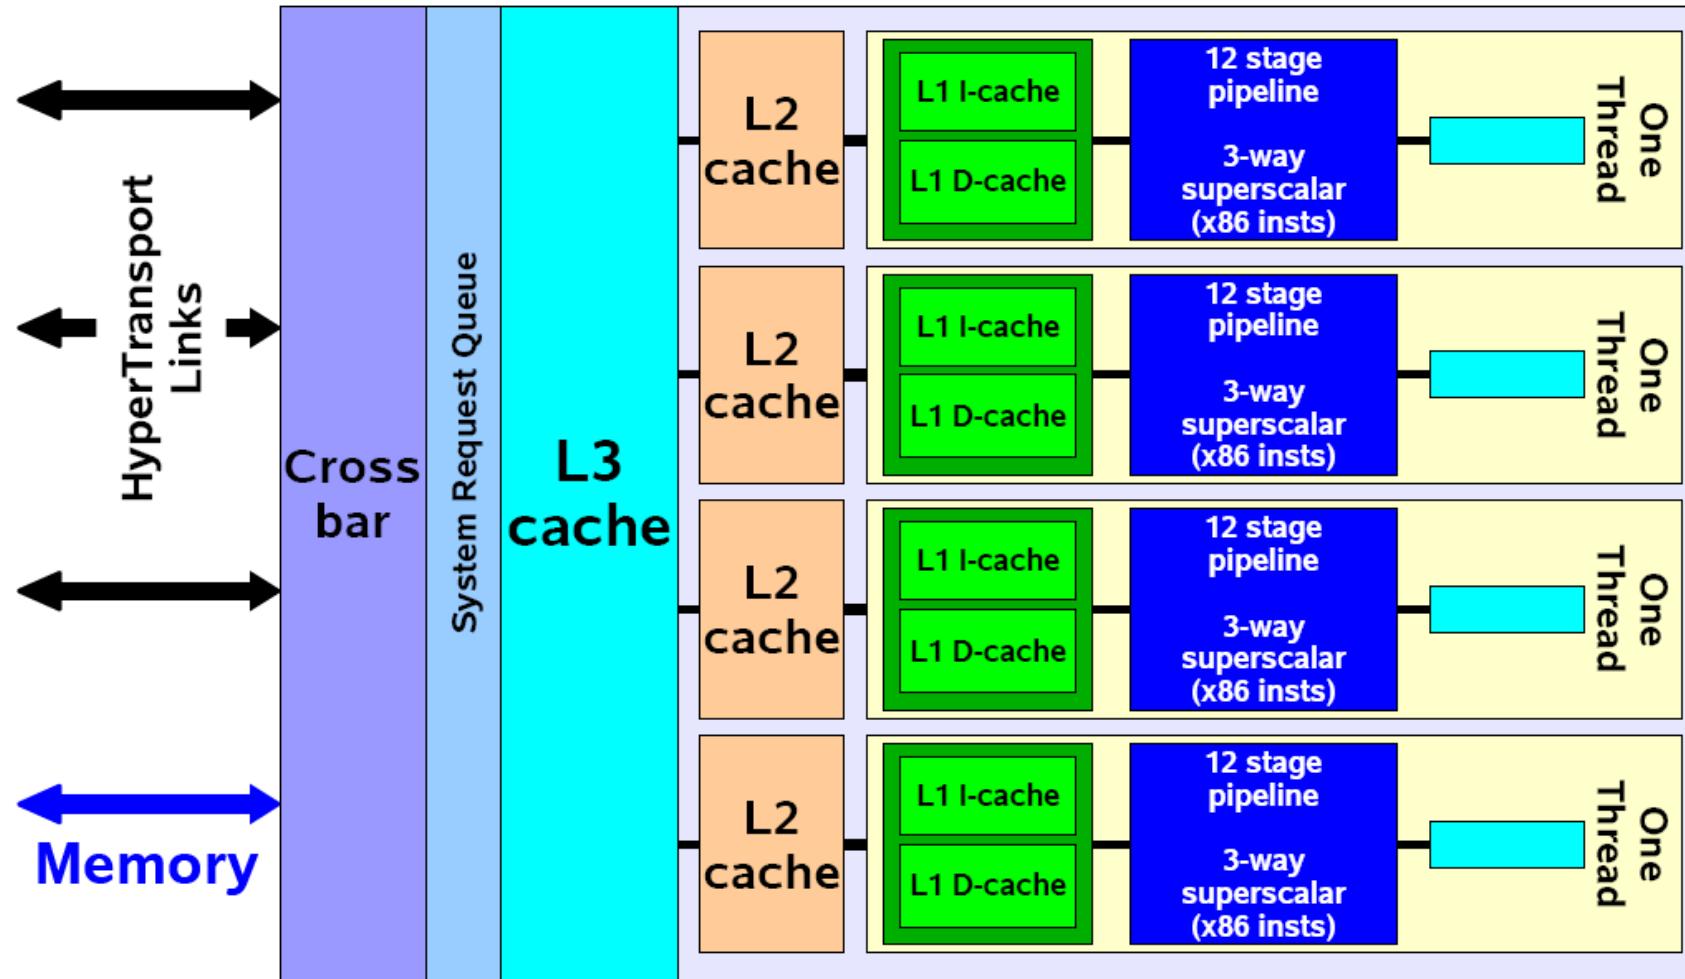


AMD Opteron – quad-core

- ❑ dedicated L2 caches
- ❑ shared L3 cache

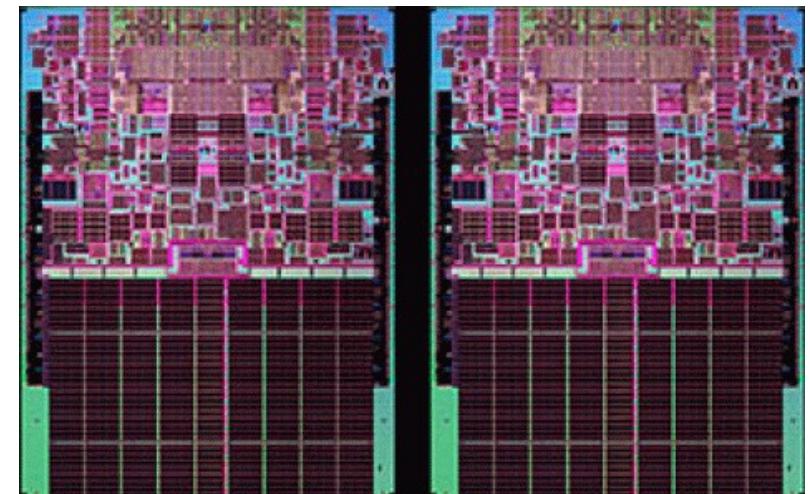
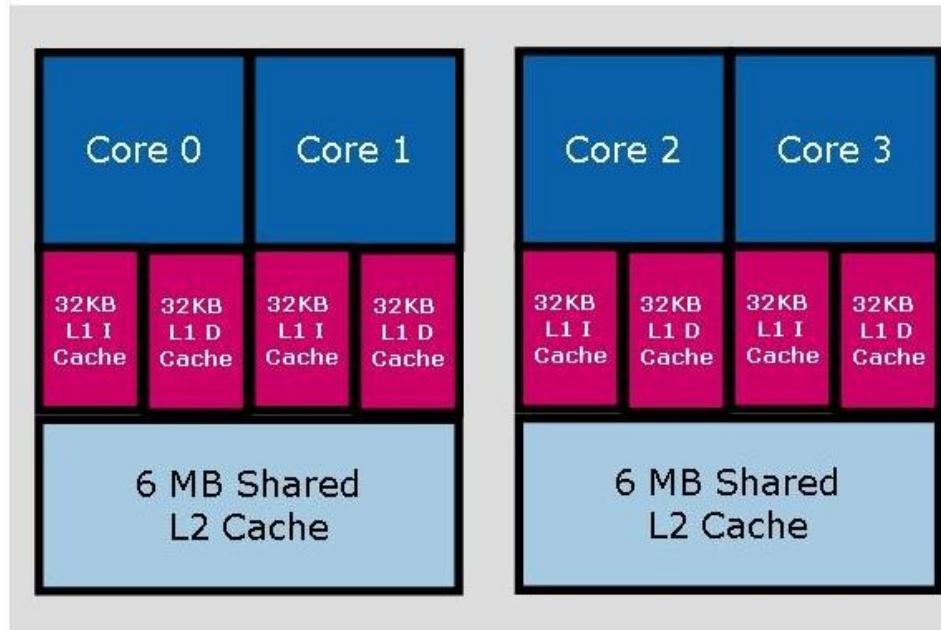


AMD Opteron – quad-core



Quad-core Intel Xeon

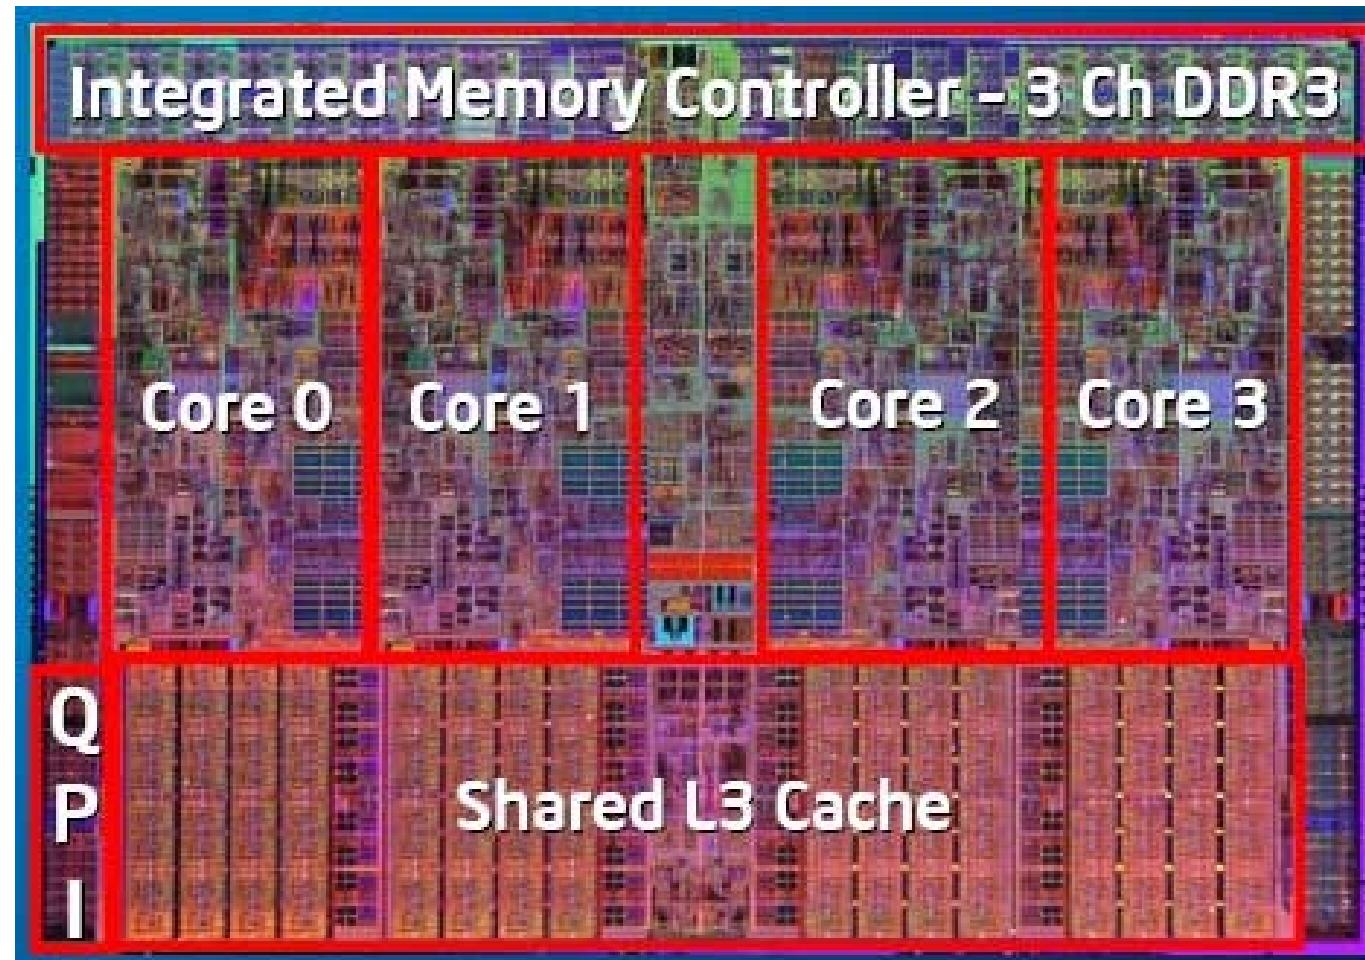
Intel® Xeon® processor 5400 series (Codename 'Harpertown')



Two dual-core chips “glued” together

The Intel “Nehalem” CPU

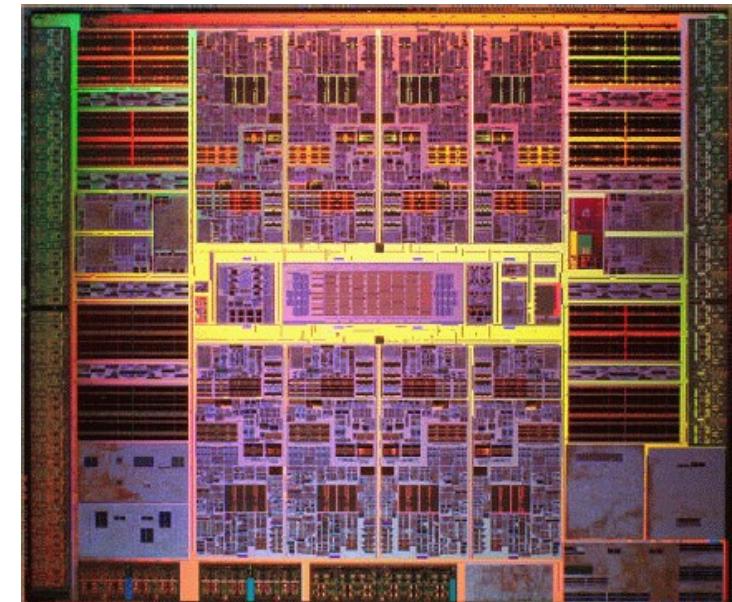
First quad-core CPU with QPI



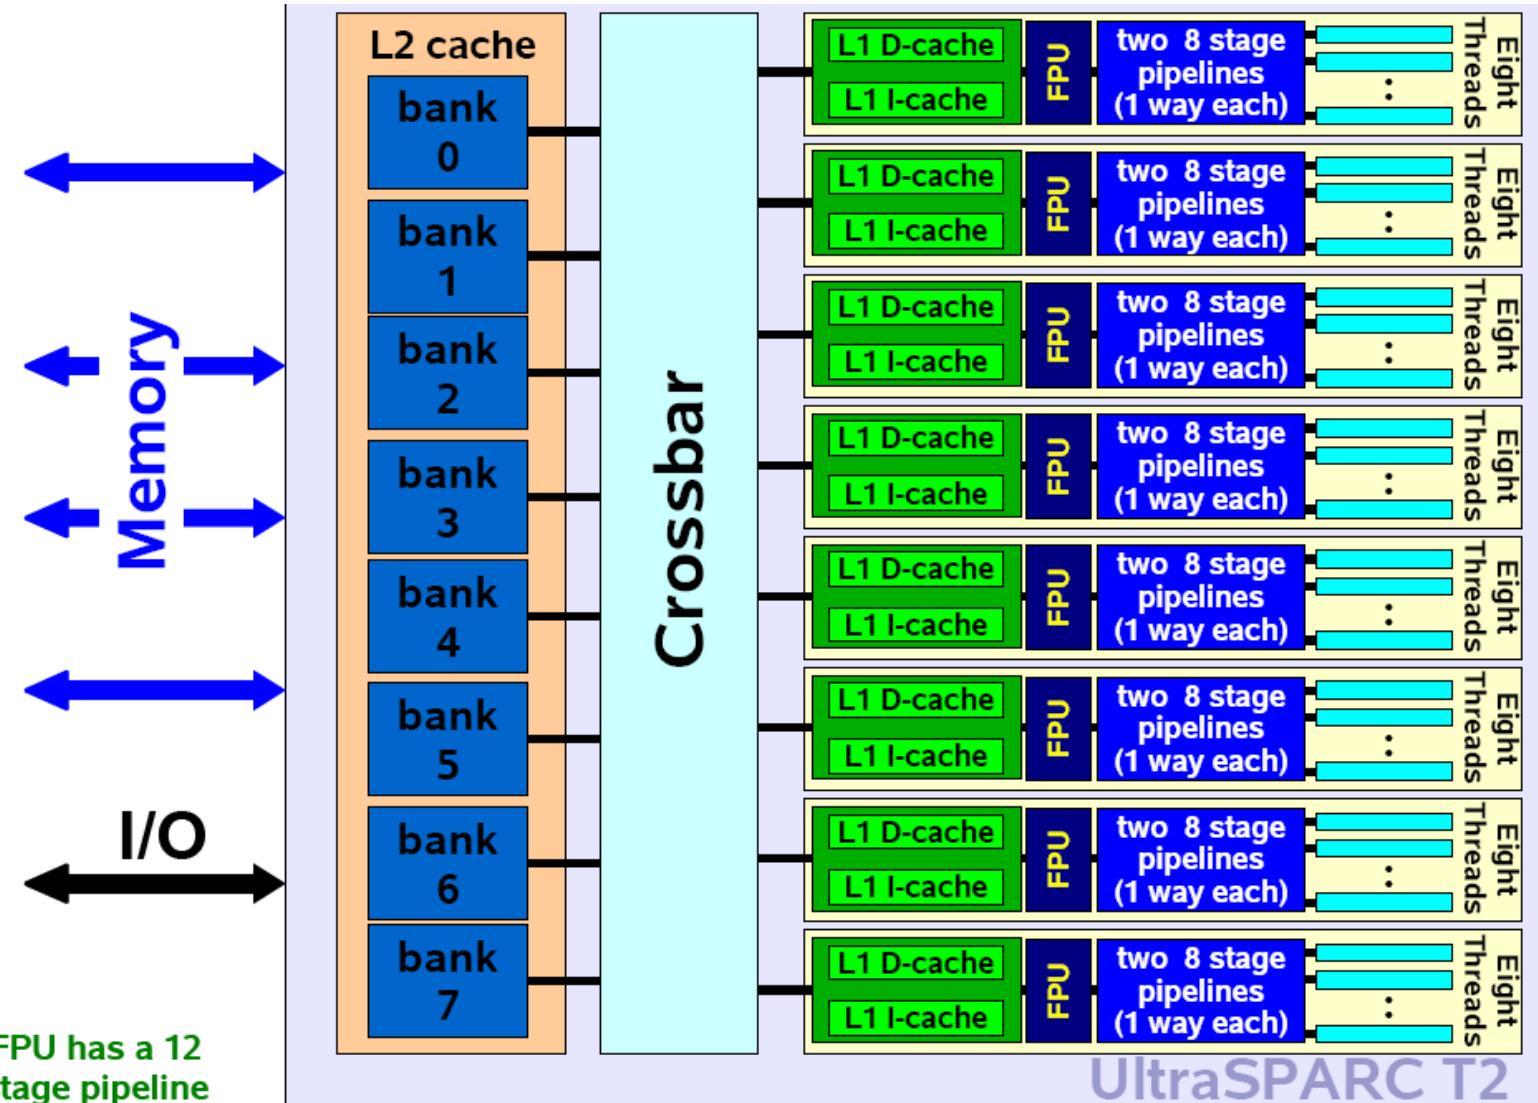
UltraSPARC-T2

System on a chip:

- ❑ 8 cores with 8 threads = 64 threads
- ❑ integrated multi-threaded 10 Gb/s Ethernet
- ❑ integrated crypto-unit per core
- ❑ low power (< 95W)
- ❑ < 1.5W/thread



UltraSPARC-T2



FPU has a 12 stage pipeline

Why adding threads to a core?

Execution of two threads:



Interleaving the work – better utilization:



Keyword: “Throughput Computing”

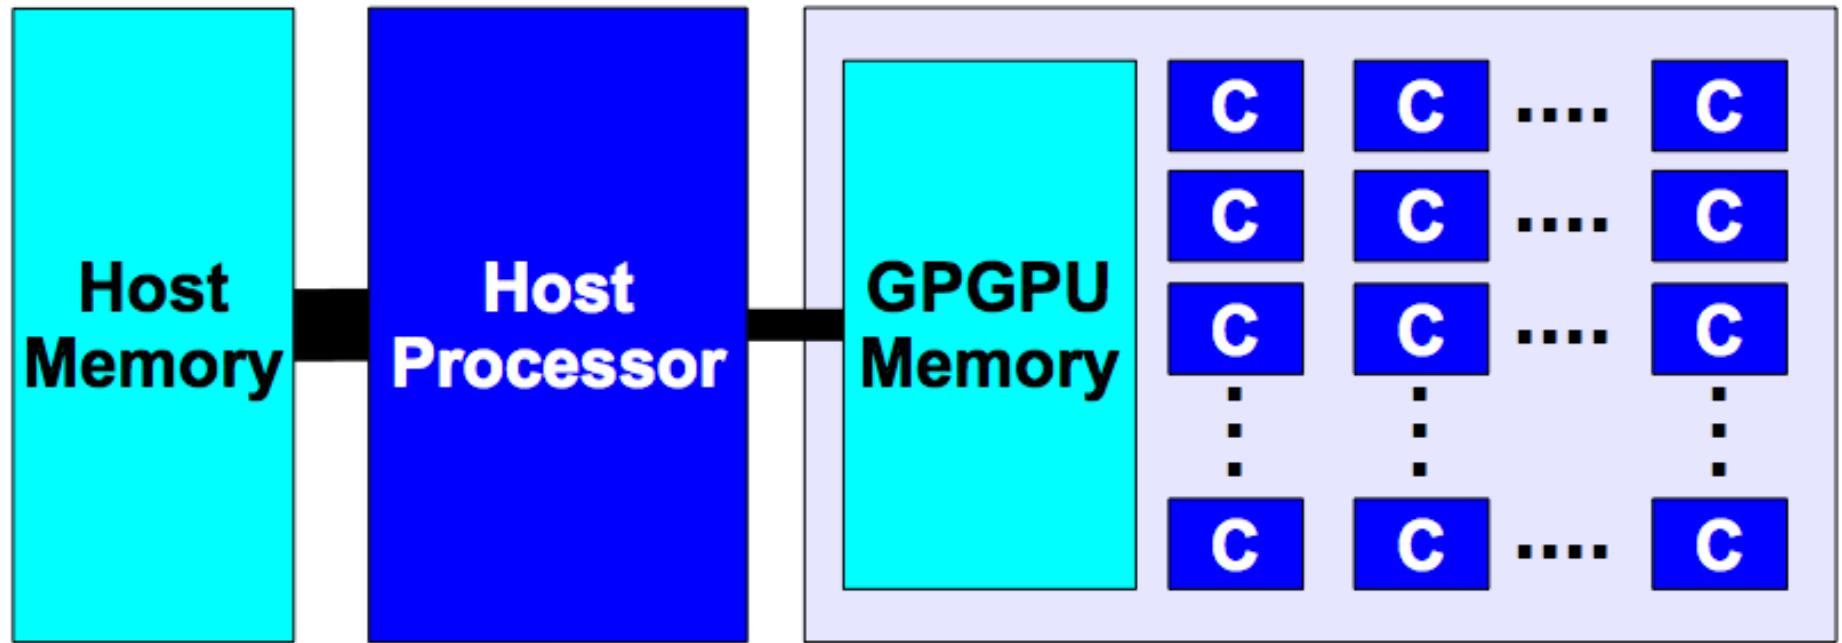
GPU computing - Accelerators



Graphics Processors

- ❑ Specialized hardware for operations typical for graphics rendering
- ❑ lots of cores (SPs – scalar processors)
- ❑ very fast memory (expensive!) - limited in size (compared to main memory)
- ❑ more instructions have been added over the last years to do more general purpose computing
- ❑ programming environments (CUDA, OpenCL) to harness the power of the GPUs

A generic GPGPU



GPU “features”

- ❑ every “core” is a very simple processor
- ❑ “cores” cannot work independently
- ❑ no independent execution of threads, but SIMD
(Single Instruction, Multiple Threads)
- ❑ no global address space, neither within the GPU,
nor with the CPU
- ❑ no cache-coherency
- ❑ latency hiding by executing many threads
- ❑ ==> more next week

Multi-core and memory bandwidth

Remember from “Serial Tuning”:

- ❑ The TPP is easy to calculate:
 - ❑ 2 GHz CPU/core
 - ❑ 4 Flops per clock cycle
 - ❑ TPP: 8 GFlop/s (per core!)
- ❑ To obtain that peak performance, we need to be able to feed the core with the right amount of data!

Multi-core and memory bandwidth

How much data is that?

- ❑ 4 floating point operations (add, mult) need 8 floating numbers => 64 bytes (double prec.)
- ❑ That is 64 bytes / clock cycle or 128 GB/s
 - ❑ 128 GB is the content of more than 27 DVDs!!!
- ❑ But what is the memory bandwidth in modern machines?

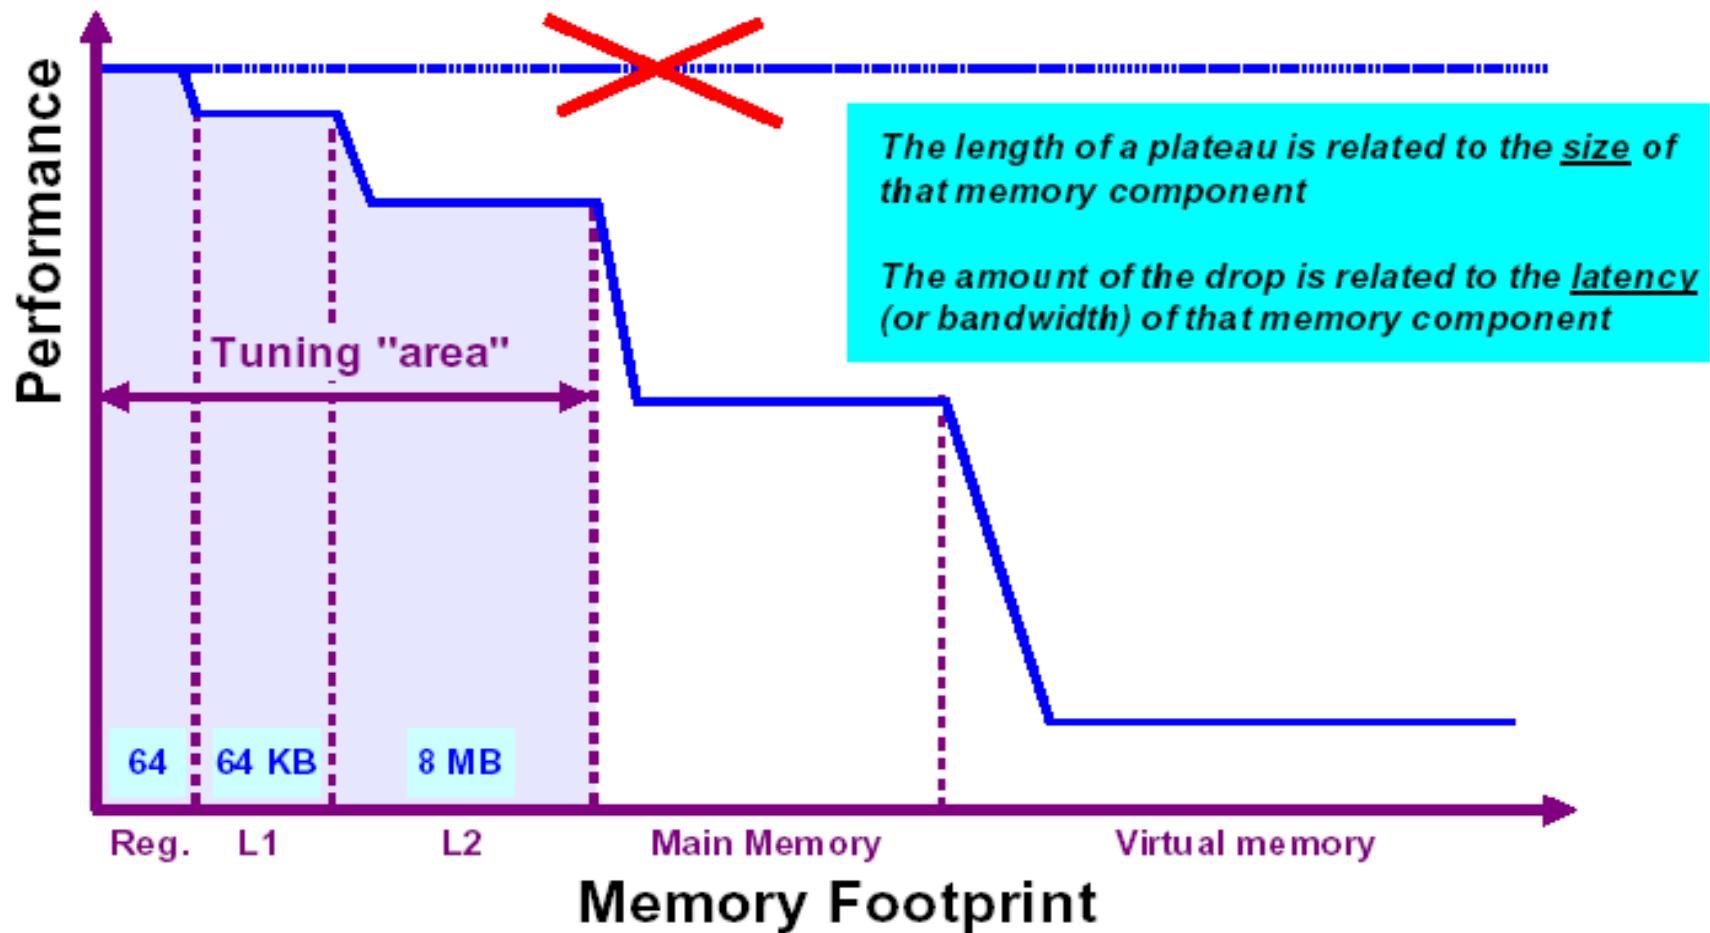
Multi-core and memory bandwidth

Memory bandwidth of the Xeon 5500 family

- ❑ equipped with DDR3-1333 DIMMS:
 - ❑ 1333 MT/s => 10.6 GB/s per memory channel
 - ❑ or 32 GB/s maximum (all channels equipped)
- ❑ with DDR3-1066 DIMMS:
 - ❑ max. 25.5 GB/s
 - ❑ that's per socket – but each CPU has 4 cores!
 - ❑ => the memory bandwidth per core is less!!!!

Multi-core and memory bandwidth

Flashback (single core):



Multi-core and memory bandwidth

Compute-bound vs memory-bound

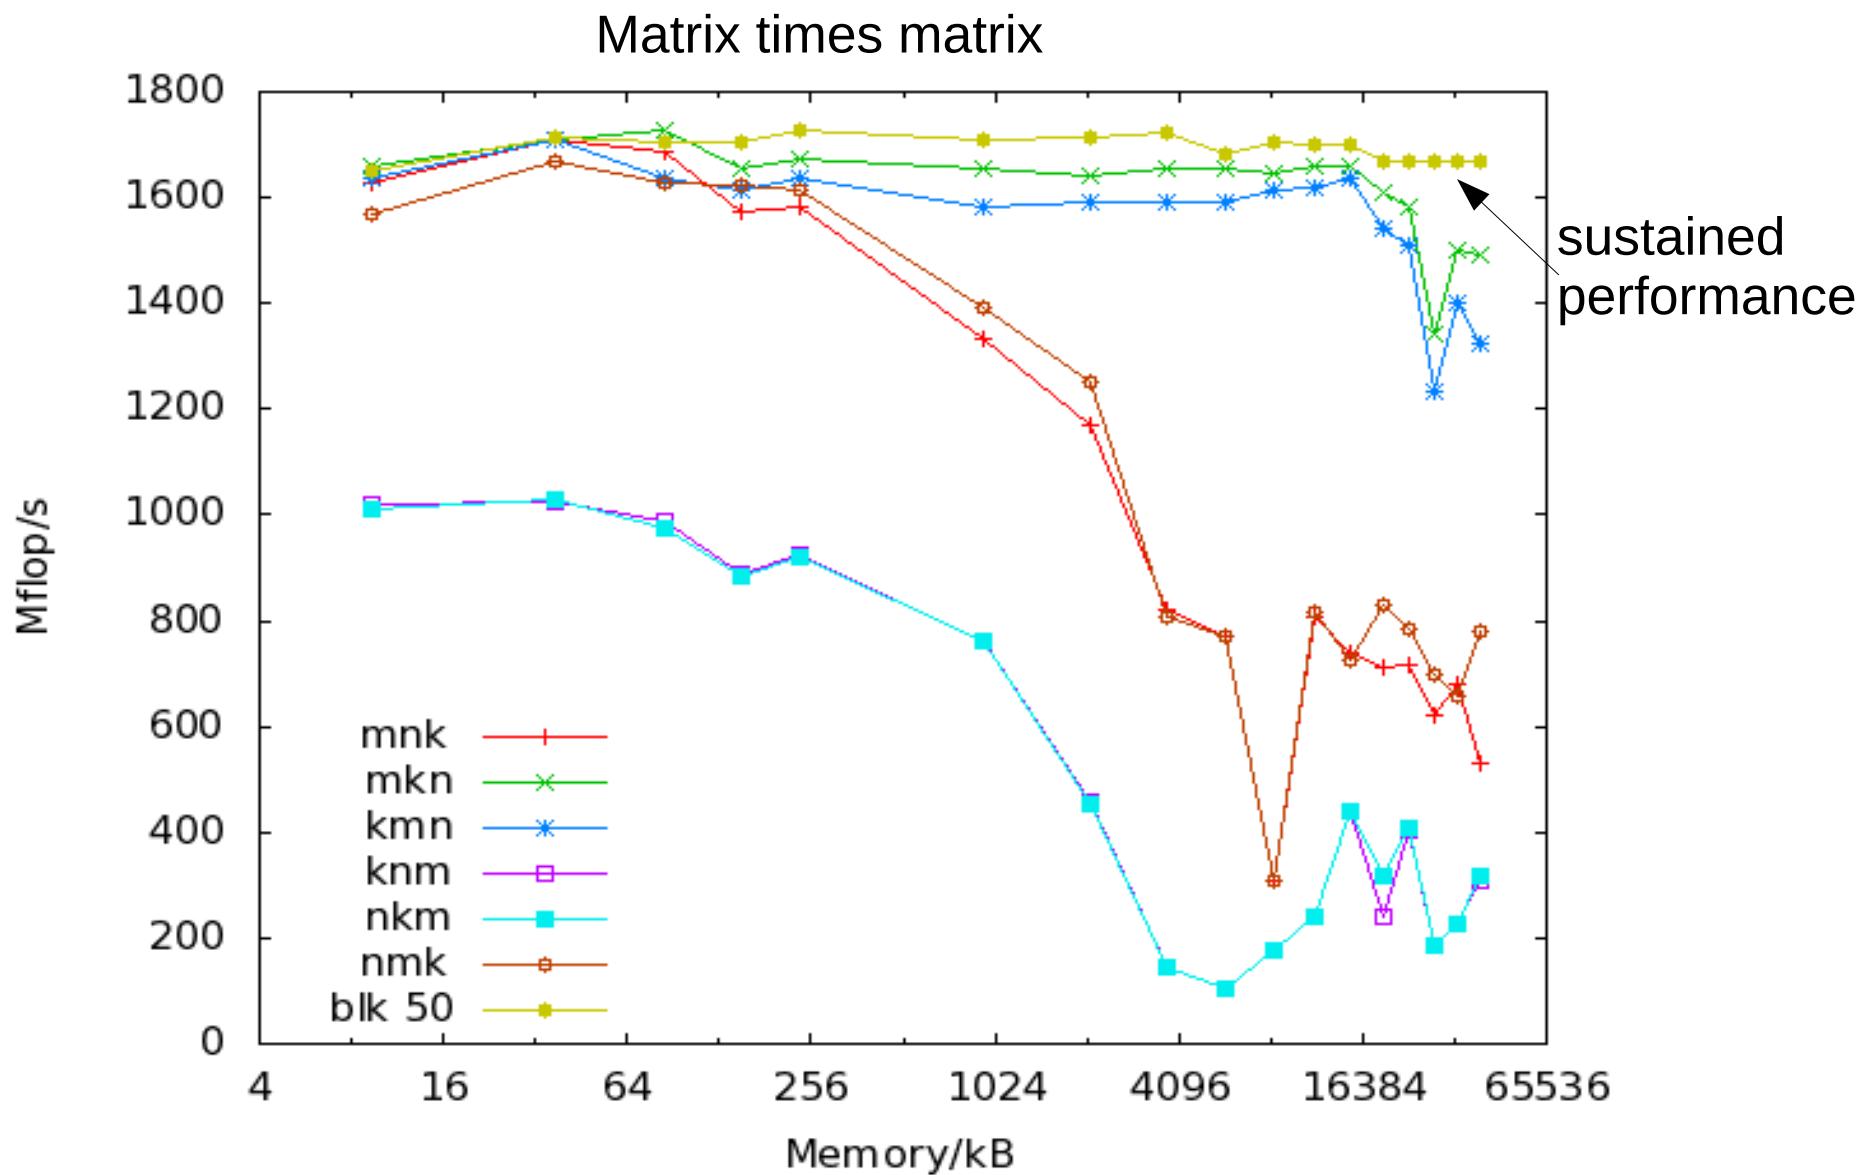
- ❑ Compute-bound:
 - ❑ the number of flops is higher than the number of mem-ops (load/store)
 - ❑ example: matrix times matrix
- ❑ Memory bound:
 - ❑ the number of mem-ops is dominating the flops
 - ❑ example: matrix times vector

Multi-core and memory bandwidth

Compute-bound vs memory-bound

- ❑ but ...
- ❑ for large problems, even matrix times matrix gets dominated by the memory operations, it turns into a memory bound problem
- ❑ we know how to solve that – e.g. by blocking algorithms, to keep the problem compute bound

Multi-core and memory bandwidth



Multi-core and memory bandwidth

Scaling of memory-bound applications

- ❑ What happens with memory bound applications on multi-core systems?
- ❑ N cores – one CPU, e.g. $N = 4$
- ❑ N threads execute a mem-bound kernel
 - ❑ => all N threads fight for the same mem-bw
 - ❑ => this is a performance bottleneck
 - ❑ => this has a negative effect on scaling (speed-up)

Multi-core and memory bandwidth

Scaling of memory-bound applications (cont'd)

- ❑ this effect is visible for all parallel applications, regardless of the programming model, i.e.
 - ❑ MPI applications
 - ❑ OpenMP applications
- ❑ adding more cores to the CPU makes the situation even worse!
- ❑ adding more sockets helps, but one still can't make use of all cores!

The Roofline Model

- ❑ Amdahl's law, i.e. the definition of speed-up, doesn't take bottlenecks into account!
- ❑ Where are those bottlenecks?
 - ❑ arithmetic operations: flop/s
 - ❑ memory access: bandwidth/latency
 - ❑ network access: communication or I/O (bw/latency)

The Roofline Model

Example: part of the STREAM benchmark (Triad)

```
for (i=0; i < N; ++i) {  
    a[i] = b[i] + s * c[i];  
}
```

- ❑ How ‘expensive’ is one iteration of this loop?
 - ❑ 2 floating point operations: add & mult
 - ❑ 3 memory operations (2 loads, 1 store)
 - ❑ s is kept in a register

The Roofline Model

Arithmetic operations:

- ❑ time for a Flop: about one nanosecond (1 ns)
- ❑ modern CPUs can do several Flops per cycle (vectorization, FMA)

Memory operations:

- ❑ latency: 90ns (DRAM)
- ❑ bandwidth: a few bytes per cycle

=> the Triad code is memory bound!

The Roofline Model

Operational (or arithmetic) intensity:

$$I = \frac{\text{arithmetic operations} [\text{flops}]}{\text{data transfers } (ld, st) [\text{words}]}$$

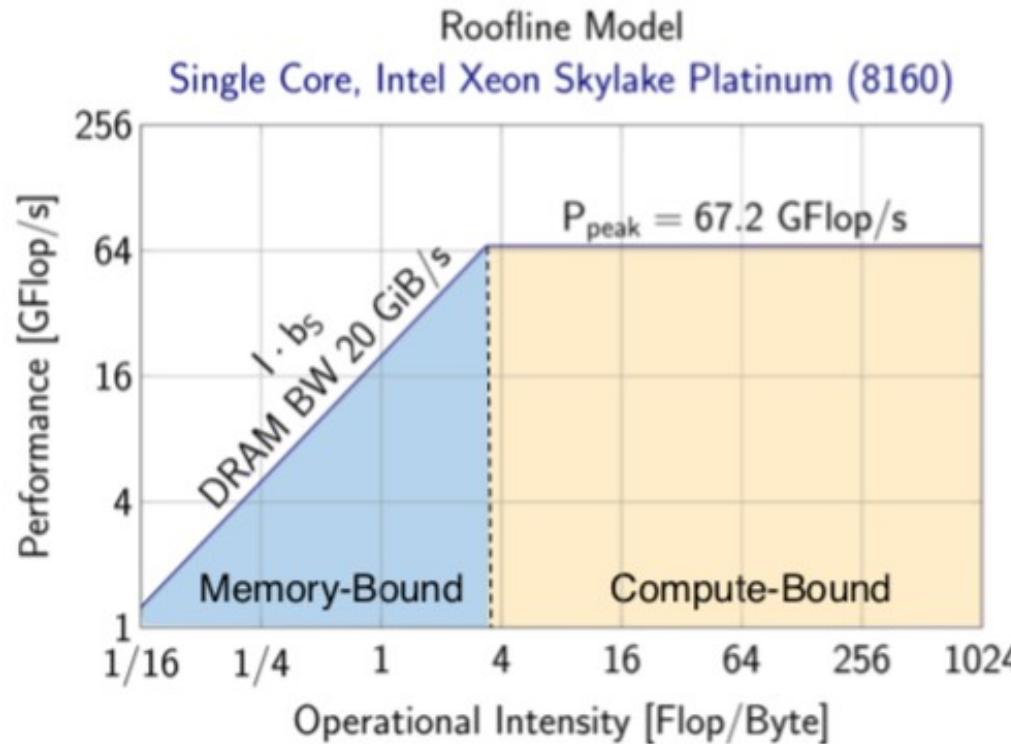
- ❑ e.g. for the triad problem:
 - ❑ $I = 2 \text{ flops}/3 \text{ words} = 0.66 \text{ flops}/\text{word}$
 - ❑ with 8 bytes/word (double): $I = 0.083 \text{ flops}/\text{byte}$
- ❑ now we also need to consider the bandwidth of the slowest data path, b_s (in bytes/second)

The Roofline Model

Achievable Performance P :

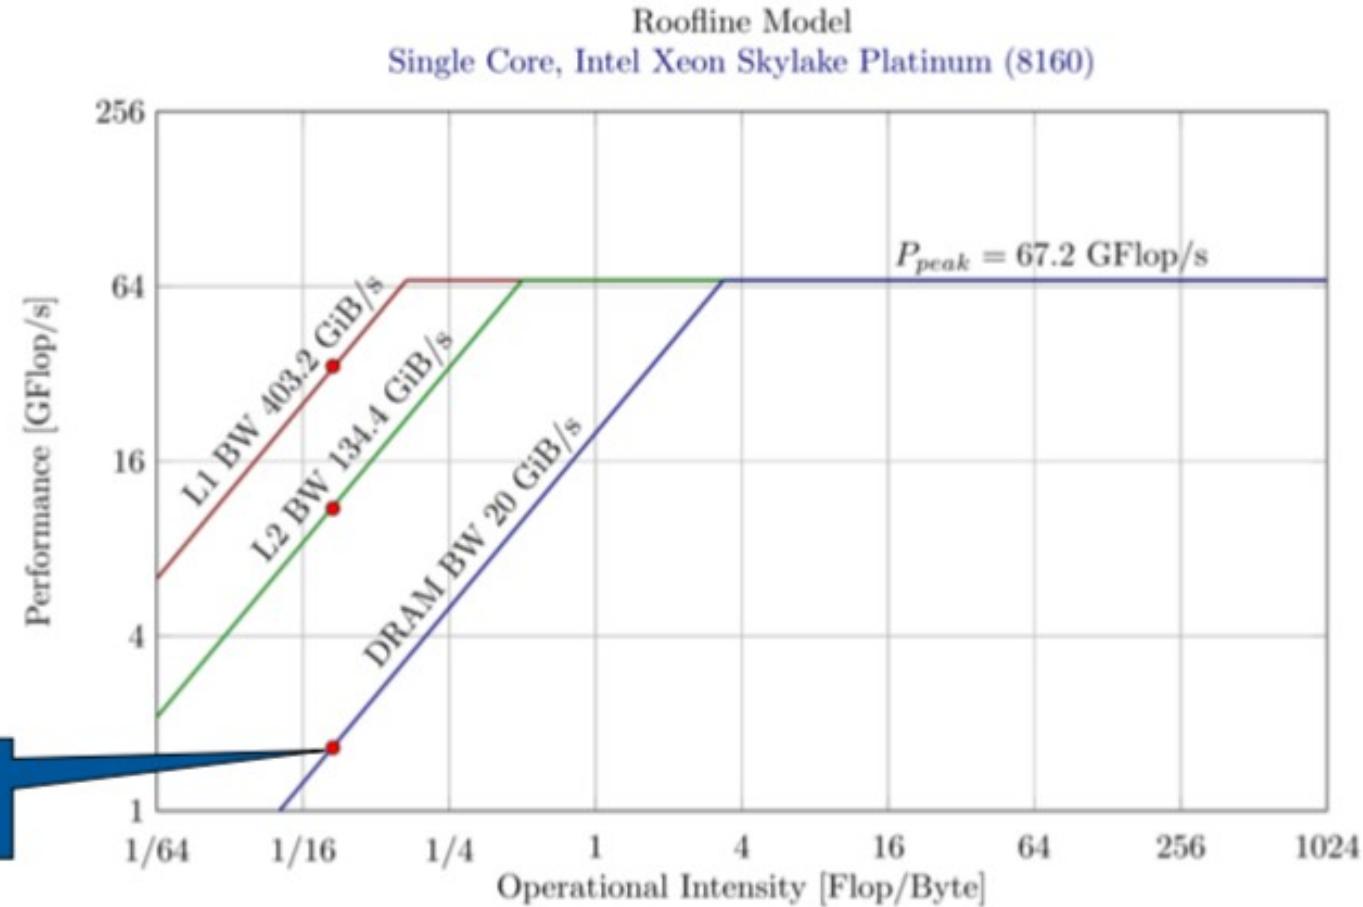
$$P := \min(I * b_s, P_{\text{peak}})$$

- ❑ This is called the ‘Roofline (Performance) Model’



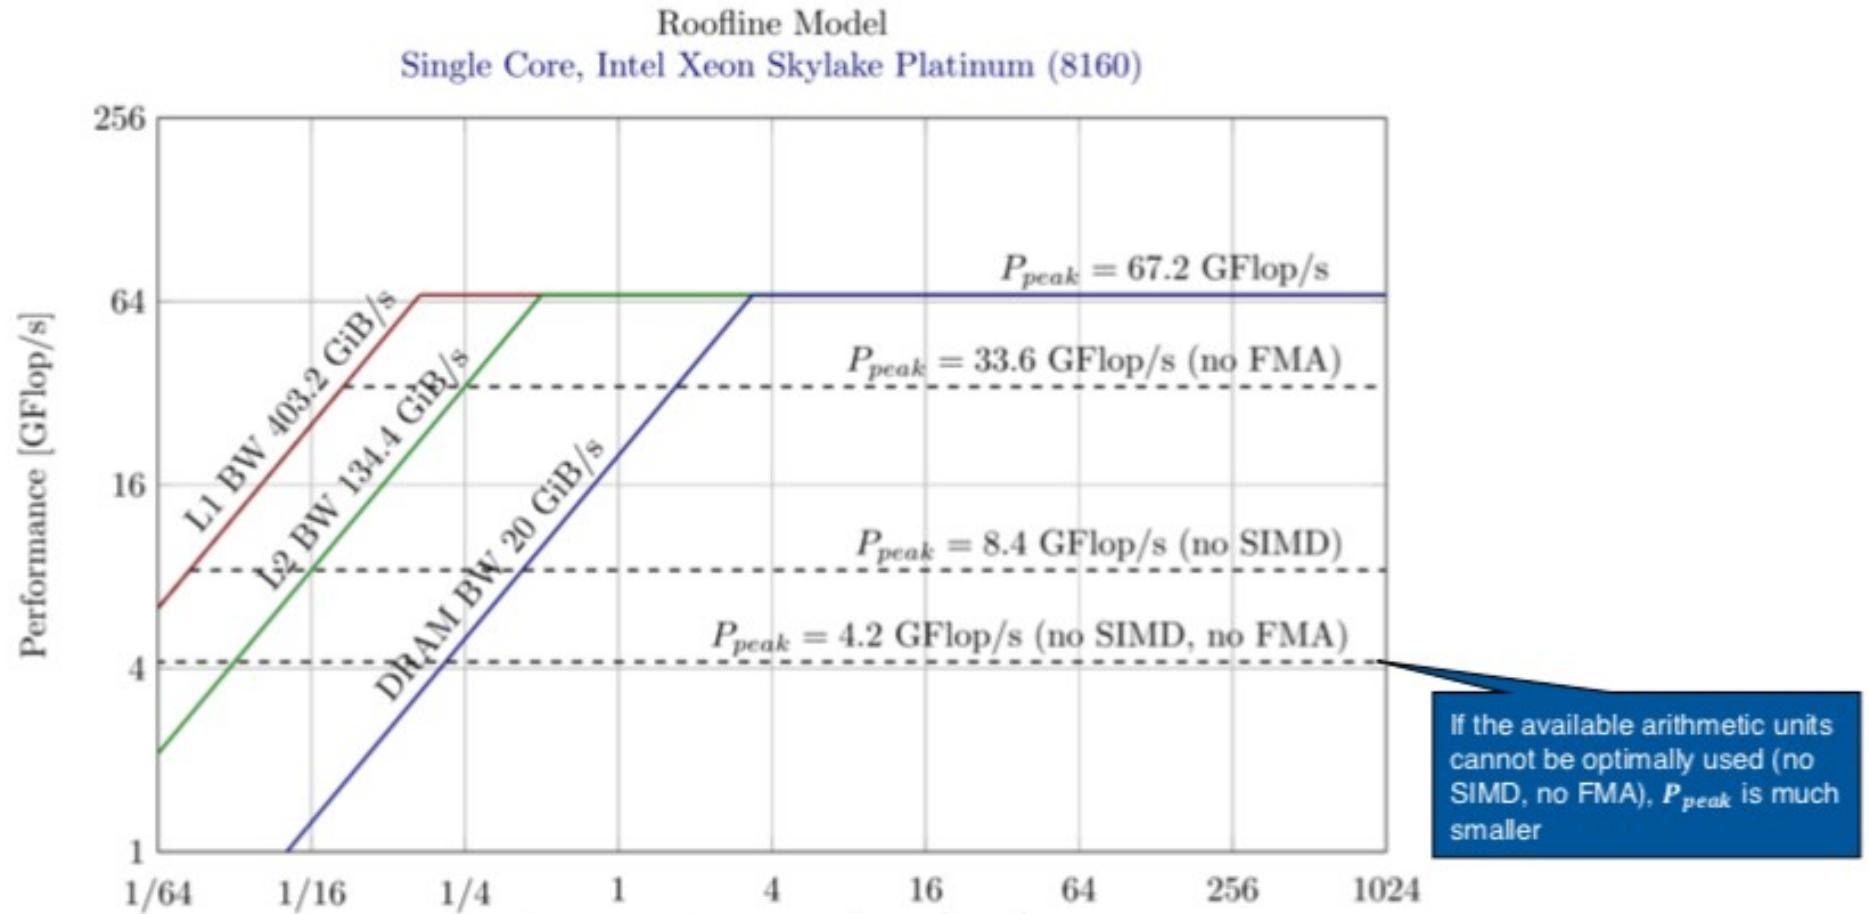
from: Simon Schwitanski
RWTH Aachen (2022)

The Roofline Model



from: Simon Schwitanski
RWTH Aachen (2022)

The Roofline Model



from: Simon Schwitanski
RWTH Aachen (2022)

The Roofline Model

Classification of problems:

Arithmetic oper.	Data transfers	Classification	Examples
$O(N)$	$O(N)$	memory bound	scalar product, vector ops, sparse matrix times vector
$O(N^2)$	$O(N^2)$	memory bound (but can benefit from cache)	dense matrix times vector, matrix addition
$O(N^x)$ ($x > 2$)	$O(N^2)$	compute bound	dense matrix multiplication

Summary

- ❑ You have heard about:
 - ❑ Parallel programming models and basic concepts
 - ❑ Parallel architectures (shared / distributed memory)
 - ❑ Cache-coherency
 - ❑ Multi-core CPUs
 - ❑ GPUs/accelerators
 - ❑ Multi-core and memory bandwidth
- ❑ Next 3 lectures:

Portable programming of shared memory systems
with OpenMP