

Parallel Programming in OpenMP

Outline

- ❑ OpenMP Overview
- ❑ OpenMP
 - ❑ Basics
 - ❑ Syntax: Directives
 - ❑ Syntax: Clauses
 - ❑ Example

OpenMP – Overview

What is OpenMP?

OpenMP – Overview

- ❑ OpenMP: stands for

Open Multi Processing

- ❑ parallel programming model for shared memory multiprocessors
- ❑ 'de-facto' standard, not an industry *standard*
- ❑ not a new language, but
 - ❑ compiler directives
 - ❑ support function library

OpenMP – Overview

- ❑ OpenMP development is community driven
- ❑ Architecture Review Board (ARB):
 - ❑ hardware and software vendors
 - ❑ government and academia
- ❑ Founded back in 1997
 - ❑ unification of different vendor “standards”
- ❑ Official OpenMP website:
 - ❑ <https://www.openmp.org/>

OpenMP – Overview

- ❑ standard versions:
 - ❑ C/C++: version 2.0 (March 2002)
 - ❑ Fortran: version 2.0 (November 2000)
 - ❑ version 2.5 (Fortran and C/C++ / May 2005)
 - ❑ version 3.0 (May 2008)
 - ❑ version 3.1 (July 2011)
 - ❑ ...

OpenMP – Overview

- ❑ 4.0 (Jul 2013)
 - ❑ exception handling
- ❑ 4.5 – update (Nov 2015)
 - ❑ first support for accelerators (“GPU offloading”)
- ❑ 5.0 (Nov 2018), 5.2 – update (Nov 2021)
 - ❑ is implemented in most compilers now
- ❑ 6.0 (Nov 2024)
- ❑ The OpenMP standard specifications:
 - ❑ <http://www.openmp.org/specifications/>

OpenMP – Overview

OpenMP Literature:

- ❑ “Using OpenMP – The Next Step” by R. van der Pas, E. Stotzer and C. Terboven, MIT Press (2017)
– available via findit.dtu.dk
- ❑ “The OpenMP Common Core” by T.G. Mattson, Y. He, and A.E. Koniges, MIT Press (2019)
- ❑ “Using OpenMP”, B. Chapman, G. Jost, R. van der Pas, MIT Press (2007)

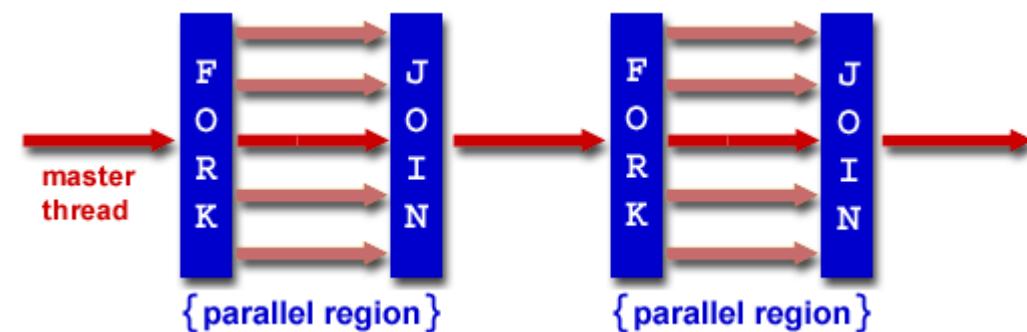
OpenMP Basics

Basic elements of OpenMP

OpenMP Basics

OpenMP uses the “Fork-Join Model”:

- ❑ All programs begin as a single process: the **master thread**.
- ❑ FORK: the master thread creates a team of parallel threads (**parallel region**).
- ❑ JOIN: synchronization and termination of the **worker threads**.



OpenMP Basics

OpenMP is mostly based on compiler directives:

- ❑ C/C++:

```
#pragma omp directive [clause]  
{ ...code block... }
```

- ❑ Fortran:

```
!$OMP directive [clause]  
...code block...  
!$OMP end directive
```

OpenMP Basics

The OpenMP API has also

- ❑ a set of support library functions:
`omp_... ()`
e.g. `omp_get_thread_num()`
- ❑ control via environment variables:
`OMP_...`
e.g. `OMP_NUM_THREADS`

OpenMP Basics

First OpenMP version of “Hello world”:

```
#include <stdio.h>

int main(int argc, char *argv[] ) {
    #pragma omp parallel
    {
        printf("Hello world!\n");
    } /* end parallel */
    return(0);
}
```

OpenMP Basics

Second version of “Hello world”:

```
#include <stdio.h>
#ifndef _OPENMP
#include <omp.h>
#endif

int main(int argc, char *argv[] ) {
    int t_id = 0;
#pragma omp parallel private(t_id)
{
    #ifdef _OPENMP
    t_id = omp_get_thread_num();
    #endif
    printf("Hello world from %d!\n", t_id);
} /* end parallel */
return(0);
}
```



OpenMP Basics

```
$ ./hello2  
Hello world from 0!
```

```
$ OMP_NUM_THREADS=4 ./hello2  
Hello world from 0!  
Hello world from 3!  
Hello world from 1!  
Hello world from 2!
```

- ❑ Note: The order of execution will be different from run to run!
- ❑ The default no. of threads depends on the OpenMP implementation

OpenMP Components

Directives

- ❑ Parallel regions
- ❑ Worksharing
- ❑ Synchronization
- ❑ Data scoping
- ❑ no. of threads
- ❑ Orphaning

Environment variables

- ❑ no. of threads
- ❑ Scheduling
- ❑ Dynamic thread adjustment
- ❑ Nested parallelism

Runtime

- ❑ no. of threads
- ❑ Scheduling
- ❑ Dynamic thread adjustment
- ❑ Nested parallelism
- ❑ API for timers & locking

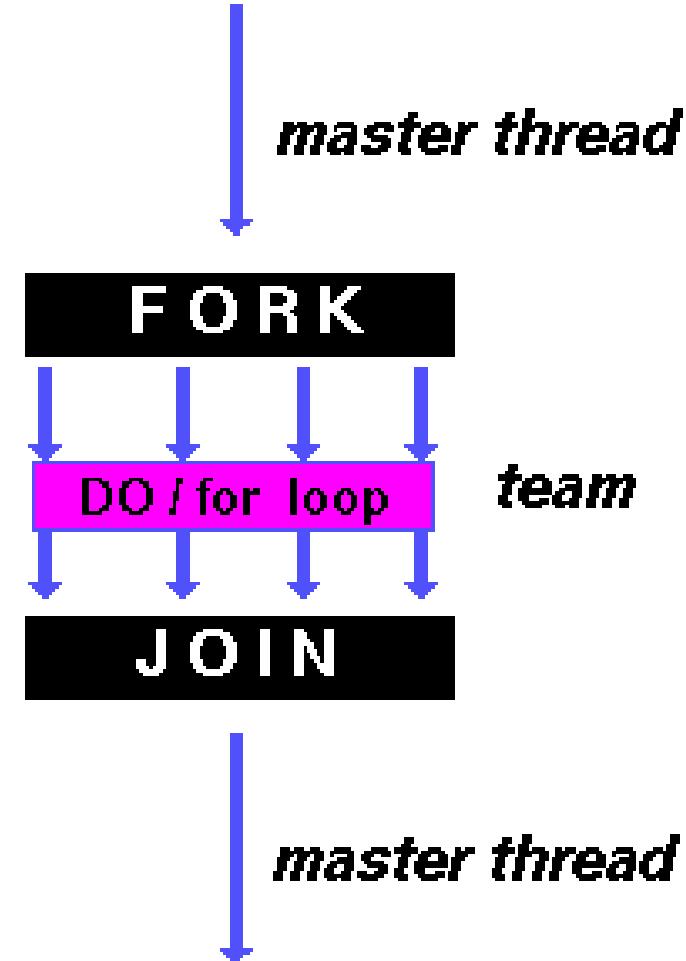
OpenMP Basics

Work-sharing

OpenMP Basics

Work-sharing constructs – 1:

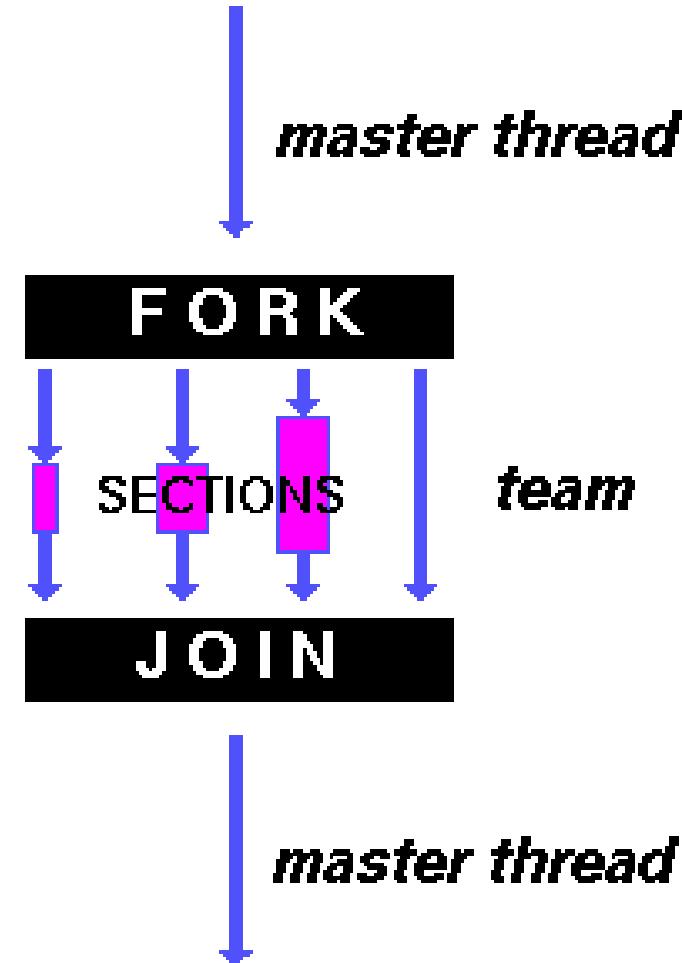
- ❑ do/for
- ❑ loop parallelism
- ❑ most common



OpenMP Basics

Work-sharing constructs – 2:

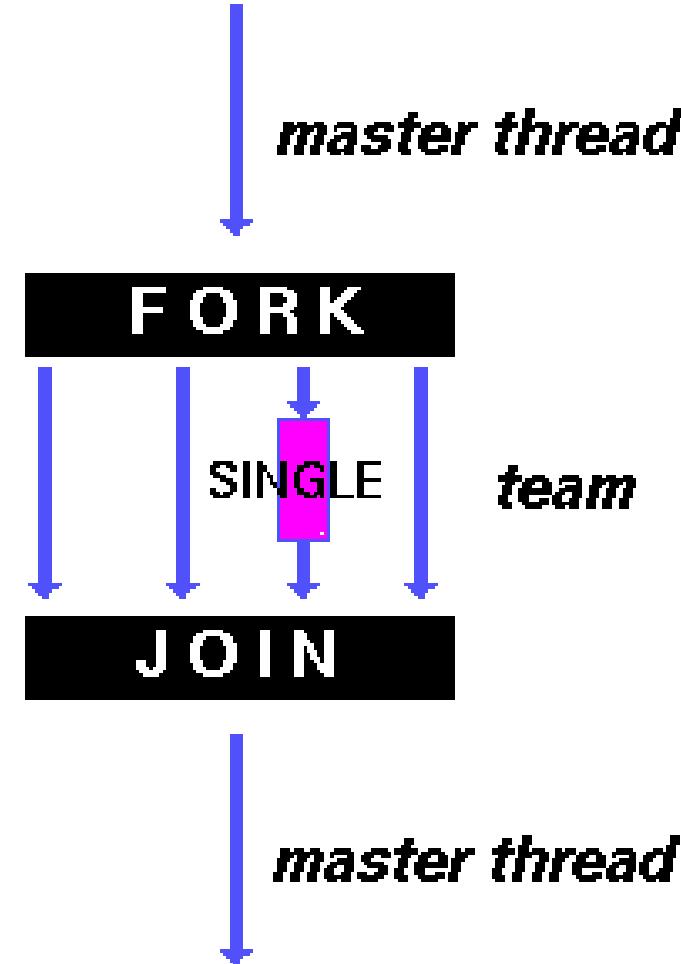
- ❑ sections
- ❑ functional parallelism
- ❑ typically independent calculations



OpenMP Basics

Work-sharing constructs – 3:

- ❑ **single**
- ❑ work assigned to one thread only
- ❑ typically I/O



OpenMP Basics

Important rules for work-sharing constructs:

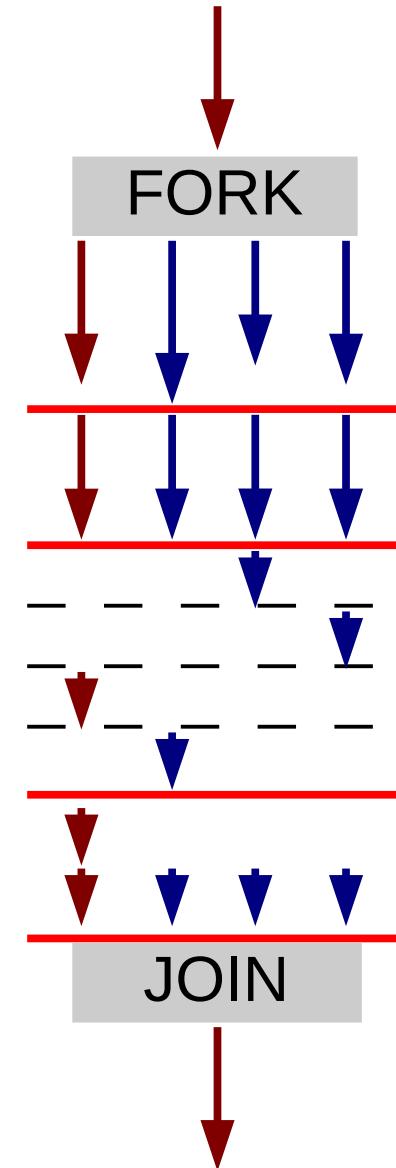
- ❑ must be enclosed in a parallel region
- ❑ must be encountered by all team members
- ❑ must be encountered in the same order

OpenMP Basics

Synchronization

OpenMP Basics

- ❑ most synchronization in OpenMP is implicit, but sometimes explicit synchronization is needed:
- ❑ barriers
- ❑ critical regions
- ❑ master only
- ❑ explicit locking



OpenMP Syntax

OpenMP programming in C/C++
Part I: Directives

OpenMP Syntax

OpenMP directives – general form:

```
#pragma omp directive [[clause] \
[clause] ...]

{
    <statements>
} /* end of omp directive */
```

- ❑ **Note:** There is no “omp end” pragma!
- ❑ Best practice: add a comment at the end of the structured block!

OpenMP Syntax

Parallel region:

- ❑ Starts a team of parallel threads
- ❑ executes code in parallel
- ❑ synchronize/terminate threads

```
main() {  
    A();  
    #pragma omp parallel  
    {  
        B(); // all threads do B()!  
    } /* end omp parallel */  
    C();  
}
```

OpenMP Syntax

Work-sharing – Loop parallelism:

- ❑ OpenMP implements parallel do/for-loops only!

```
int i;  
float a[N], b[N], c[N];  
  
for (i=0; i < N; i++)  
    a[i] = b[i] = i * 1.0;  
  
#pragma omp parallel shared(a,b,c) private(i)  
{  
    #pragma omp for  
    for (i=0; i < N; i++)  
        c[i] = a[i] + b[i];  
  
} /* end of parallel region */
```

*for has to follow
the pragma – no
{...}!*

OpenMP Syntax

Work-sharing – Loop parallelism:

- ❑ Another version: combined “parallel for”

```
int i;  
float a[N], b[N], c[N];  
  
for (i=0; i < N; i++)  
    a[i] = b[i] = i * 1.0;  
  
#pragma omp parallel for shared(a,b,c) \  
                           private(i)  
for (i=0; i < N; i++)  
    c[i] = a[i] + b[i];
```

OpenMP Syntax

Work-sharing – Fortran 95 array syntax

- ❑ Fortran 95 allows to address parts of or whole arrays – and the compiler will translate this into loops.
- ❑ A special Fortran directive:

```
double precision, dimension() :: A, B  
double precision, dimension(N,M) :: C
```

```
!$OMP WORKSHARE  
...  
A(1:M) = A(1:M) + B(1:M)  
C = 0.00  
...  
!$OMP END WORKSHARE
```

OpenMP Syntax

Work-sharing – Functional parallelism:

- Parallel sections:

```
#pragma omp parallel shared(a,b,c) private(i)
{
    #pragma omp sections
    {
        #pragma omp section
        for (i=0; i < N/2; i++)
            c[i] = a[i] + b[i];

        #pragma omp section
        for (i=N/2; i < N; i++)
            c[i] = a[i] + b[i];
    } /* end of sections */
} /* end of parallel region */
```

OpenMP Syntax

Work-sharing – Single thread execution:

- ❑ Work done by one thread only

```
#pragma omp parallel shared(a,b,c) private(i)
{
    #pragma omp single
    { read_array(a); read_array(b); }

    #pragma omp for
    for (i=0; i < N; i++)
        c[i] = a[i] + b[i];

    #pragma omp single
    { write_array(c); }
} /* end of parallel section */
```

OpenMP Syntax

Work-sharing – conditional parallelism:

- ❑ the `if(...)` clause

```
int i;  
float a[N], b[N], c[N];  
  
for (i=0; i < N; i++)  
    a[i] = b[i] = i * 1.0;  
  
#pragma omp parallel if (N > 10000) \  
    shared(a,b,c) private(i)  
{  
    #pragma omp for  
    for (i=0; i < N; i++)  
        c[i] = a[i] + b[i];  
}
```



OpenMP Syntax

The `num_threads(...)` clause:

```
#pragma omp parallel ... num_threads(int_expr)
{
    ...
}
```

- ❑ only one `num_threads` clause per parallel directive
- ❑ `int_expr` is evaluated at runtime, before the parallel region is entered

OpenMP syntax

The `collapse(n)` clause

- ❑ a way to parallelize loop nests

```
subroutine sub()  
  
!$omp do collapse(2) private(i,j,k)  
    do k = kl, ku, ks  
        do j = jl, ju, js  
            do i = il, iu, is  
                call bar(a,i,j,k)  
            enddo  
        enddo  
    enddo  
!$omp end do  
end subroutine
```

collapse the two outer loops over k and j

OpenMP Syntax

Synchronization – Critical region:

- ❑ specifies a region of code that must be executed by only **one** thread at a time!
- ❑ can be named

```
#pragma omp parallel private(loc_sum)
{
    ...
#pragma omp for
for(int i = 0; i < n; i++)
    loc_sum += x[i];

#pragma omp critical (cr_sum)
sum += loc_sum;
} /* end of parallel section */
```

OpenMP Syntax

Synchronization – Atomic construct:

- ❑ specifies a single operation(!) that must be executed by only **one** thread at a time!
- ❑ restricted syntax (see OpenMP standard)

```
int x = 0;

#pragma omp parallel shared(x)
{
    ...
    #pragma omp atomic
    x = x + 1;
    ...
} /* end of parallel section */
```

OpenMP Syntax

Synchronization – Master region:

- ❑ specifies a region of code that is executed by the master thread only!
- ❑ ignored by others – no implicit barriers!

```
#pragma omp parallel
{
    ...
#pragma omp master
{
    printf("Hello\n");
}
...
} /* end of parallel section */
```

OpenMP Syntax

Synchronization – Ordered:

- ❑ executes code block in the sequential order
- ❑ everything outside can run in parallel
- ❑ only within a parallel do/for loop

```
#pragma omp parallel for \
    private(i,myf) ordered
for(i = 0; i < n; i++) {
    myf = i+1;
    #pragma omp ordered
    {
        sum = sum + myf;
    }
} /* end of parallel for */ /*end of parallel for*/
```

```
#pragma omp parallel for \
    private(i,myf)
for(i = 0; i < n; i++) {
    myf = i+1;
    #pragma omp critical
    {
        sum = sum + myf;
    }
}
```

OpenMP Syntax

Output:

non-ordered

vs.

ordered

sum [1] = 1

sum [1] = 1

sum [6] = 7

sum [2] = 3

sum [2] = 9

sum [3] = 6

sum [7] = 16

sum [4] = 10

sum [3] = 19

sum [5] = 15

sum [8] = 27

sum [6] = 21

sum [4] = 31

sum [7] = 28

sum [9] = 40

sum [8] = 36

sum [10] = 50

sum [9] = 45

sum [5] = 55

sum [10] = 55

Result: 55

Result: 55

OpenMP Syntax

Ordered with dependences (OpenMP 4.5)

- ❑ new clause `depend(...)` enhances ordered by adding dependences
 - ❑ `#pragma omp ordered depend(sink: vec)`
 - ❑ start the ordered region with the dependences given in `vec`
 - ❑ `#pragma omp ordered depend(source)`
 - ❑ ends the ordered execution defined by sink above
- ❑ this new feature allows to parallelize e.g. loops with dependences (see next slide)

OpenMP Syntax

doacross loops:

- ❑ loops with dependencies

```
#pragma omp parallel
{
    #pragma omp for ordered(1)
    for (int i=1; i<n; i++) {
        #pragma omp ordered depend(sink:i-1)
        a[i] = a[i-1] + b[i];
        #pragma omp ordered depend(source)

        c[i] = 2*a[i];
    } // End of for-loop
} // End of parallel region
```

of dependences

dependence vector

OpenMP Syntax

ordered depend(...) and doacross loops:

- ❑ provides logic in the OpenMP runtime, that else would have to be done manually
- ❑ helps to “resolve” static dependences and allows part of the code to run in parallel
- ❑ useful, if the parallel part of the loop(s) is more computational intensive than the ordered part
- ❑ might need a different loop scheduling (depends on compiler/runtime)

OpenMP Syntax

Synchronization – Barrier:

- ❑ synchronizes all threads in a team

```
#pragma omp parallel
{
    ...
#pragma omp barrier
    ...
} /* end of parallel section */
```

OpenMP Syntax

Synchronization – Implied barriers:

- ❑ exit from parallel region
- ❑ exit from `omp for/omp do/omp workshare`
- ❑ exit from sections
- ❑ exit from single

**No *implied* barrier on the master construct,
neither on entry nor on exit!**

OpenMP Syntax

OpenMP programming in C/C++
Part II: Clauses

OpenMP Syntax

Data scoping clauses:

- ❑ Understanding and the use of data scoping is really essential.
- ❑ Most problems/errors are due to wrong data scoping.
- ❑ Most variables are shared by default (shared memory programming model).
- ❑ Private variables: loop indices, stack of subroutines.

OpenMP Syntax

OpenMP Data scope attribute clauses:

- ❑ private
- ❑ shared
- ❑ default
- ❑ reduction
- ❑ firstprivate
- ❑ lastprivate
- ❑ copyin

OpenMP Syntax

The “private” clause:

- ❑ declares variables private to each thread:

```
#pragma omp directive private (list)
```

- ❑ a new variable is declared once for each thread
- ❑ all references are replaced with references to the newly declared variable
- ❑ variables declared private are uninitialized for each thread!

OpenMP Syntax

The “shared” clause:

- ❑ declares variables to be shared among all threads:

```
#pragma omp directive shared (list)
```

- ❑ a shared variable exists in only one memory location and all threads have read/write access to that address
- ❑ proper access to the variable is left to the programmer – that's YOU!

OpenMP Syntax

The “default” clause:

- ❑ allows the programmer to specify the default scope for all variables:

```
#pragma omp dir default(shared|none)
```

- ❑ C/C++ knows only those two types
- ❑ only one default clause per parallel region
- ❑ Best practice: use default(None) and scope all your variables explicitly

OpenMP Example

Two examples

OpenMP Example

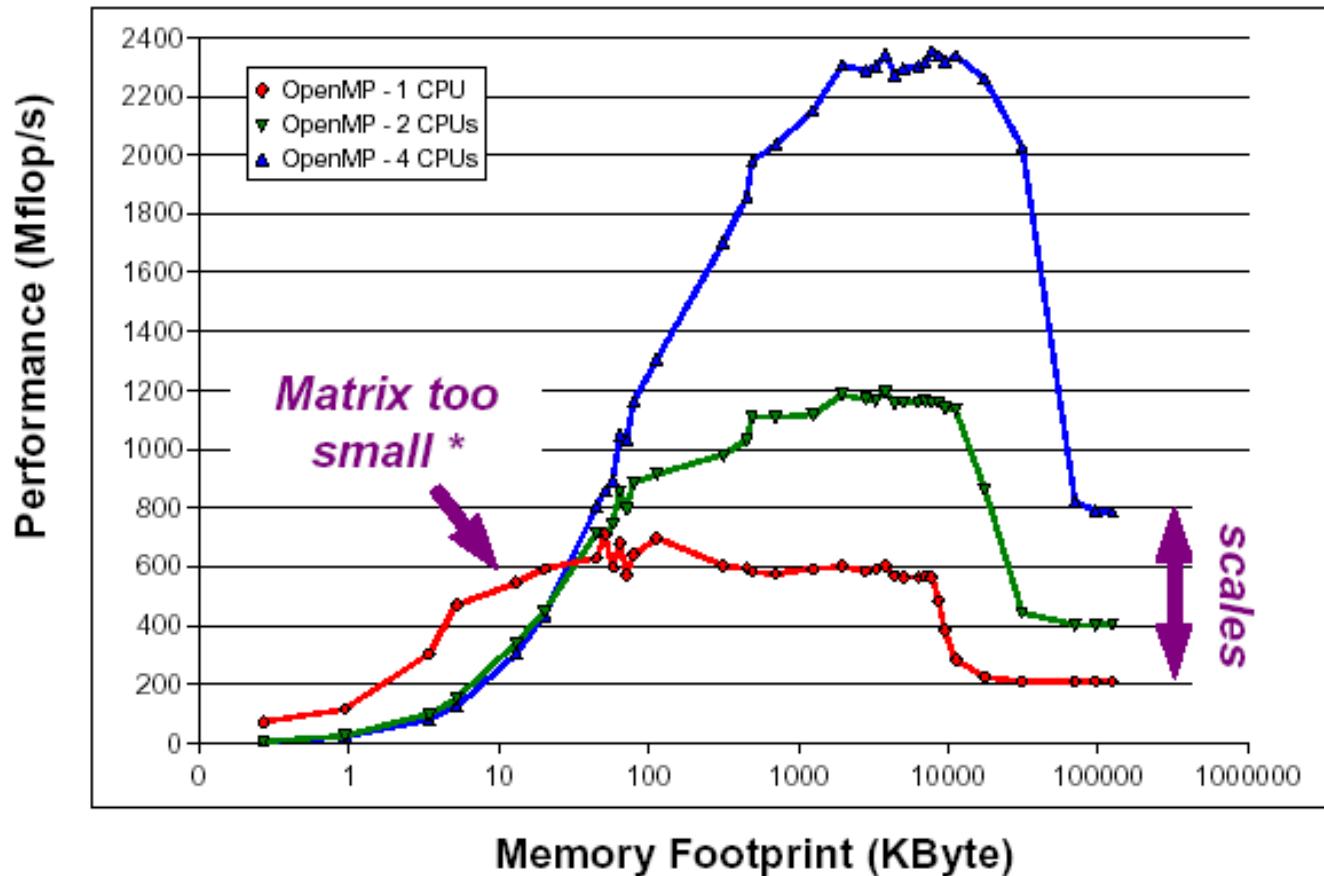
Matrix times vector

```
void
mxv(int m, int n, double *a, double *b, double *c)
{
    int i, j;
    double sum;

    #pragma omp parallel for default(none) \
                    shared(m,n,a,b,c) private(i,j,sum)
    for (i=0; i<m; i++) {
        sum = 0.0;
        for (j=0; j<n; j++)
            sum += b[i*n+j] * c[j];
        a[i] = sum;
    }
}
```



OpenMP Example



SunFire 6800
UltraSPARC III Cu @ 900 MHz
8 MB L2-cache

*) With the IF-clause in OpenMP this performance degradation can be avoided

OpenMP Example

Matrix times vector (using C99 standard)

```
void
mxv(int m, int n, double *a, double *b, double *c)
{
    double sum;

    #pragma omp parallel for default(none) \
        shared(m,n,a,b,c) private(sum)
    for (int i=0; i<m; i++) {
        sum = 0.0;
        for (int j=0; j<n; j++)
            sum += b[i*n+j] * c[j];
        a[i] = sum;
    }
}
```



variables declared inside a parallel region are already private!

OpenMP Example

Example: numerical integration of $f(x)$

```
int i, n;  
double h, x, sum;  
  
h = 1.0 / (double)n;  
sum = 0.0;  
  
#pragma omp parallel for default(none) \  
    shared(n,h,sum) private(i,x)  
  
for(i=1; i<=n; i++) {  
    x = h * ((double)i + 0.5);  
    #pragma omp critical  
    sum += f(x);  
}  
  
Race condition!  
sequential code!
```

OpenMP Example

Example: numerical integration of $f(x)$

❑ Improvement 1

```
int i, n;  
double h, x, fx, sum;  
  
h = 1.0 / (double)n;  
sum = 0.0;  
  
#pragma omp parallel for default(none) \  
    shared(n,h,sum) private(i,x,fx)  
for(i=1; i<=n; i++) {  
    x = h * ((double)i + 0.5);  
    fx = f(x); ← function evalution in parallel  
#pragma omp critical  
    sum += fx;  
}
```

OpenMP Example

Example: numerical integration of $f(x)$

❑ Improvement 2

```
int i, n; double h, x, t_sum, sum;  
  
h = 1.0 / (double)n; sum = 0.0;  
#pragma omp parallel default(none) \  
    shared(n,h,sum) private(i,x,t_sum) {  
    t_sum = 0.0;  
    #pragma omp for  
    for(i=1; i<=n; i++) {  
        x = h * ((double)i + 0.5);  
        t_sum += f(x);  
    }  
    #pragma omp critical  
    sum += t_sum;  
} // end omp parallel
```

OpenMP Syntax

The “reduction” clause:

- ❑ performs a reduction on the variables that appear on the list:

```
#pragma omp dir reduction(op: list)
```

- ❑ a private copy for each thread of all variables on the list is created
- ❑ at the end, the reduction operation is carried out and the result(s) written to the global variable(s)

OpenMP Example

Example: numerical integration of $f(x)$

❑ smart OpenMP solution

```
int i, n;
double h, x, sum;

h = 1.0 / (double)n;
sum = 0.0;

#pragma omp parallel for default(none) \
    shared(n,h) private(i,x) \
    reduction(+: sum)
for(i=1; i<=n; i++) {
    x = h * ((double)i + 0.5);
    sum += f(x);
}
```

OpenMP Exercises – I

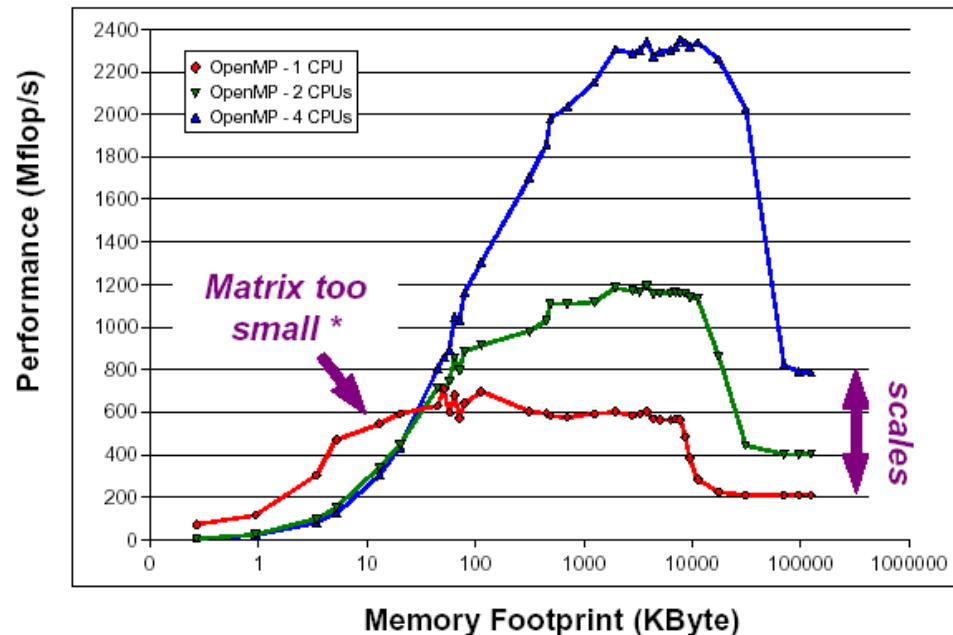
- ❑ Write an OpenMP code to calculate π , using

$$\pi = \int_0^1 \frac{4}{(1+x^2)} dx \approx \frac{1}{N} \sum_{i=1}^N \frac{4}{1 + \left(\frac{i-0.5}{N}\right)^2}$$

- ❑ implement the integrand as a function
- ❑ write your own reduction code
- ❑ use the OpenMP reduction clause
- ❑ compare the run-times

OpenMP Exercises – II

- ❑ Improve the matrix times vector example by adding an if-clause to the omp pragma – experiment with the threshold value!



SunFire 6800
UltraSPARC III Cu @ 900 MHz
8 MB L2-cache

**) With the IF-clause in OpenMP this performance degradation can be avoided*