

Parallel Programming in OpenMP – part II

Outline

- ❑ Data scoping – cont'd
- ❑ Orphaning
- ❑ Tasking
- ❑ OpenMP correctness & Data Races
- ❑ Runtime library
- ❑ Scheduling
- ❑ A real world example

OpenMP Syntax

More on data scoping

OpenMP Syntax

Reminder: the “private” clause –

- ❑ declares variables private to each thread:

```
#pragma omp directive private (list)
```

- ❑ i.e. a **new** variable is declared once for each thread
- ❑ all references are replaced with references to the newly declared variable
- ❑ variables declared private are uninitialized for each thread!

OpenMP Syntax

Consequences of private(...):

```
main() {  
    ...  
    A = 10;  
  
    #pragma omp parallel  
    {  
        #pragma omp for private(i, A, B) ...  
        for(i = 0; i < n; i++) {  
            ...  
            B = A + i; // A undefined!  
                        // unless declared firstprivate  
            ...  
        } /* end of omp for */  
        ...  
    } /* end of omp parallel */  
  
    C = B;           // B undefined!  
                    // unless declared lastprivate  
}
```

OpenMP Syntax

Solutions:

```
#pragma omp ... firstprivate(list)
```

- ❑ All variables in list are initialized with the value the original object had before entering the parallel construct.

```
#pragma omp ... lastprivate(list)
```

- ❑ The thread that executes the sequentially last iteration updates all variables in list.

OpenMP Syntax

The “threadprivate” and “copyin” clauses:

- ❑ **threadprivate(list)**: creates a private copy of global data (e.g. common blocks or global variables in modules in Fortran) for each thread
- ❑ **copyin(list)**: copies the values from the master thread into the private copies – like a ‘firstprivate’ for global variables
- ❑ subsequent modifications of list affect only the private copies – within one parallel region

OpenMP Syntax

Example 1:

```
int counter = 0;  
#pragma omp threadprivate(counter)  
  
int increment_counter()  
{  
    counter++;  
    return(counter);  
}
```

```
INTEGER FUNCTION INCREMENT_COUNTER()  
    COMMON/A22_COMMON/COUNTER  
!$OMP THREADPRIVATE (/A22_COMMON/)  
    COUNTER = COUNTER +1  
    INCREMENT_COUNTER = COUNTER  
    RETURN  
END FUNCTION INCREMENT_COUNTER
```

OpenMP Syntax

Example 2:

```
int
increment_counter()
{
    static int counter = 0;
    #pragma omp threadprivate(counter)

    counter++;
    return(counter);
}
```

OpenMP Syntax

The `copyprivate(...)` clause

- copying a value out of a single region into the private data of other threads (“broadcast”)

```
#pragma omp single copyprivate(list)
{
    ...
}
```

```
!$OMP SINGLE ....
...
!$OMP END SINGLE COPYPRIVATE (LIST)
```

OpenMP Syntax

Example:

```
int x, y; /* global data */
#pragma omp threadprivate(x, y)

void use_values(int id, int a, int b) {
    printf(" TID %d: a = %d, b = %d, c = %d, d = %d\n",
           id, a, b, x, y);
}

void init(int id, int *a, int *b) {
    int r_a, r_b;

#pragma omp single copyprivate(r_a, r_b, x, y)
{   scanf("%d %d %d %d", &r_a, &r_b, &x, &y); }

    *a = r_a; *b = r_b;
    use_values(id, *a, *b);
}
...
...
```

OpenMP Syntax

Example (cont'd):

```
int main(int argc, char *argv[] ) {  
  
    int tid = 0;  
    int a, b;  
  
#pragma omp parallel private(tid,a,b)  
{  
#ifdef _OPENMP  
    tid = omp_get_thread_num();  
#endif  
  
    init(tid,&a,&b);  
    printf("In main - TID %d: a = %d, b = %d, ",  
          " x = %d, y = %d\n",  
          tid, a, b, x, y);  
} /* end of omp parallel */  
  
    return(0);  
}
```

OpenMP Syntax

Example output:

```
$ OMP_NUM_THREADS=3 ./copypriv
1 2 3 4
          TID 2: a = 1, b = 2, c = 3, d = 4
          TID 1: a = 1, b = 2, c = 3, d = 4
          TID 0: a = 1, b = 2, c = 3, d = 4

In main - TID 0: a = 1, b = 2, x = 3, y = 4
In main - TID 1: a = 1, b = 2, x = 3, y = 4
In main - TID 2: a = 1, b = 2, x = 3, y = 4
```

```
env OMP_NUM_THREADS=3 ./copypriv
1 2 3 4
          TID 1: a = 0, b = 1, c = 3, d = 4
          TID 0: a = 0, b = 0, c = 3, d = 4
          TID 2: a = 1, b = 2, c = 3, d = 4

In main - TID 2: a = 1, b = 2, x = 3, y = 4
In main - TID 0: a = 0, b = 0, x = 3, y = 4
In main - TID 1: a = 0, b = 1, x = 3, y = 4
```

without copyprivate
on r_a and r_b

OpenMP Orphaning

Orphaning in OpenMP

OpenMP Orphaning

The OpenMP standard does not restrict worksharing and synchronization directives to be within the lexical extent of a parallel region. Those directives can be orphaned, i.e. they can appear outside a parallel region:

```
#pragma omp parallel
{
    :
    dowork();
    :
}
```

orphaned
worksharing
directive

```
void dowork(void) {
    :
    #pragma omp for
    for(i=0; i<N; i++) {
    }
    :
}
```

OpenMP Orphaning

- ❑ When an orphaned directive is detected within the dynamic extent of a parallel region, its behaviour is similar to the non-orphaned case.
- ❑ When an orphaned directive is detected in the sequential part of the program, it will be ignored.

```
dowork(); // serial for  
  
#pragma omp parallel  
{  
    :  
    dowork(); // parallel for  
    :  
}
```

```
void dowork(void) {  
    :  
    #pragma omp for  
    for(i=0; i<N; i++) {  
        :  
    }  
    :  
}
```

OpenMP syntax

Tasking (since OpenMP 3.0)

- ❑ allows parallelization of work that is generated dynamically
- ❑ provides a flexible model for irregular parallelism
- ❑ uses a “task pool” concept
- ❑ new opportunities:
 - ❑ while loops
 - ❑ recursive structures

OpenMP syntax

- ❑ Syntax C/C++:

```
#pragma omp task [clause]
{
    ...
}
```

- ❑ clause can be
 - ❑ if (int_expr)
 - ❑ default(shared|none)
 - ❑ private(list), shared(list)
 - ❑ firstprivate(list)
 - ❑ untied

OpenMP syntax

- ❑ Syntax Fortran:

```
!$OMP task [clause]  
...  
!$OMP end task
```

- ❑ where clause can be

- ❑ if (int_expr)
- ❑ default(shared|private|firstprivate|none)
- ❑ private(list), shared(list)
- ❑ firstprivate(list)
- ❑ untied

OpenMP syntax

Tasking example I:

while loop:

```
p = lhead;
while (p != NULL)
{
    do_work(p);
    p = next(p);
}
```

parallel while loop with OpenMP tasks:

```
#pragma omp parallel
{
    #pragma omp single
    {
        p = lhead;
        while (p != NULL) {
            #pragma omp task
            {
                do_work(p);
            }
            p = next(p);
        }
    } // end of single
} // end of parallel
```

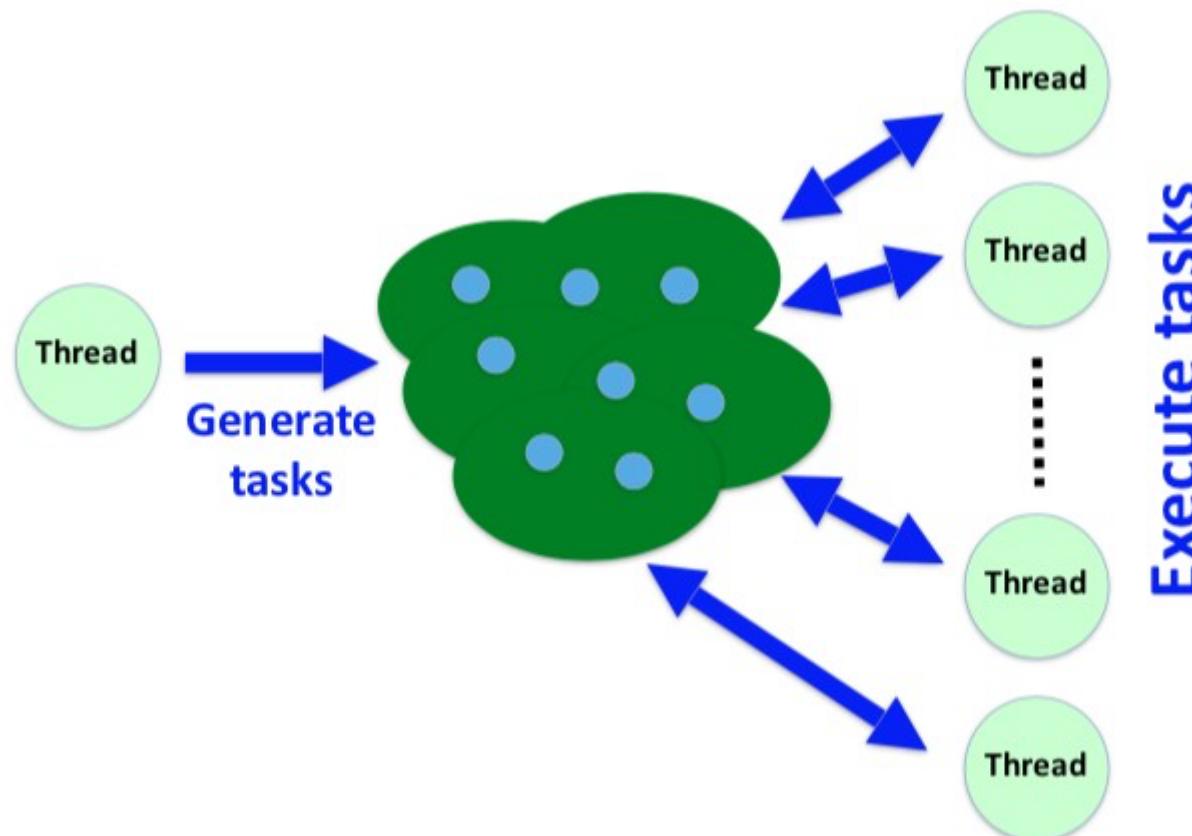
OpenMP syntax

What's going on?

```
#pragma omp parallel           ← start of parallel region
{
    #pragma omp single          ← one thread only, please
    {
        p = lhead;
        while (p != NULL) {
            #pragma omp task      ← task generation – tasks
            {                   are added to the task
                do_work(p);    list
            }
            p = next(p);
        }
    } // end of single          ← all work is done here!
} // end of parallel           ← implicit barrier – all
                                unfinished tasks have
                                to be finished
```

OpenMP syntax

The task pool concept:



courtesy: Ruud van der Pas, Oracle

OpenMP syntax

Tasks and recursion: Fibonacci numbers

- ❑ Recursive scheme to calculate the n^{th} Fibonacci number:
 - ❑ $\text{fib}(n) = \text{fib}(n-1) + \text{fib}(n-2)$
 - ❑ stopping critererion: return 1 if $n < 2$
- ❑ Caveat: this method is not very effective, but used here to demonstrate the concept of tasking!

OpenMP syntax

The sequential code:

```
int  
main(int argc, char* argv[]) {  
    [...]  
    fib(input);  
    [...]  
}
```

```
int  
fib(int n) {  
  
    int x, y;  
  
    if (n < 2) return n;  
  
    x = fib(n - 1);  
    y = fib(n - 2);  
  
    return(x + y);  
}
```

OpenMP syntax

OpenMP version of fib() with tasks:

```
int
fib(int n) {

    int x, y;

    if (n < 2) return n;

    #pragma omp task shared(x)
    x = fib(n - 1);
    #pragma omp task shared(y)
    y = fib(n - 2);

    #pragma omp taskwait
    return(x + y);
}
```

note the special
scoping rules!

generate two tasks,
calling fib() recursively

task synchronization -
to get the right results

OpenMP syntax

Scoping rules with tasks:

- ❑ Static and global variables are shared
- ❑ Local (aka automatic) variables are private
- ❑ Orphaned task variables are firstprivate
- ❑ Non-orphaned task variables inherit the shared attribute
- ❑ (Local) Task variables are firstprivate, unless declared shared
- ❑ Thus, we have to declare x and y as shared

OpenMP syntax

Task synchronization:

- ❑ `#pragma omp taskwait`
- ❑ suspends the encountering task, until all child tasks are completed
- ❑ direct children only, not descendants
- ❑ needed here, to make sure that x and y are still exist when we take the sum.

OpenMP syntax

OpenMP version of main() with tasks:

```
int
main(int argc, char* argv[] ) {
    [ . . . ]

    #pragma omp parallel ←
    {
        #pragma omp single ←
        {
            fib(input);
        }
    } // end of omp parallel

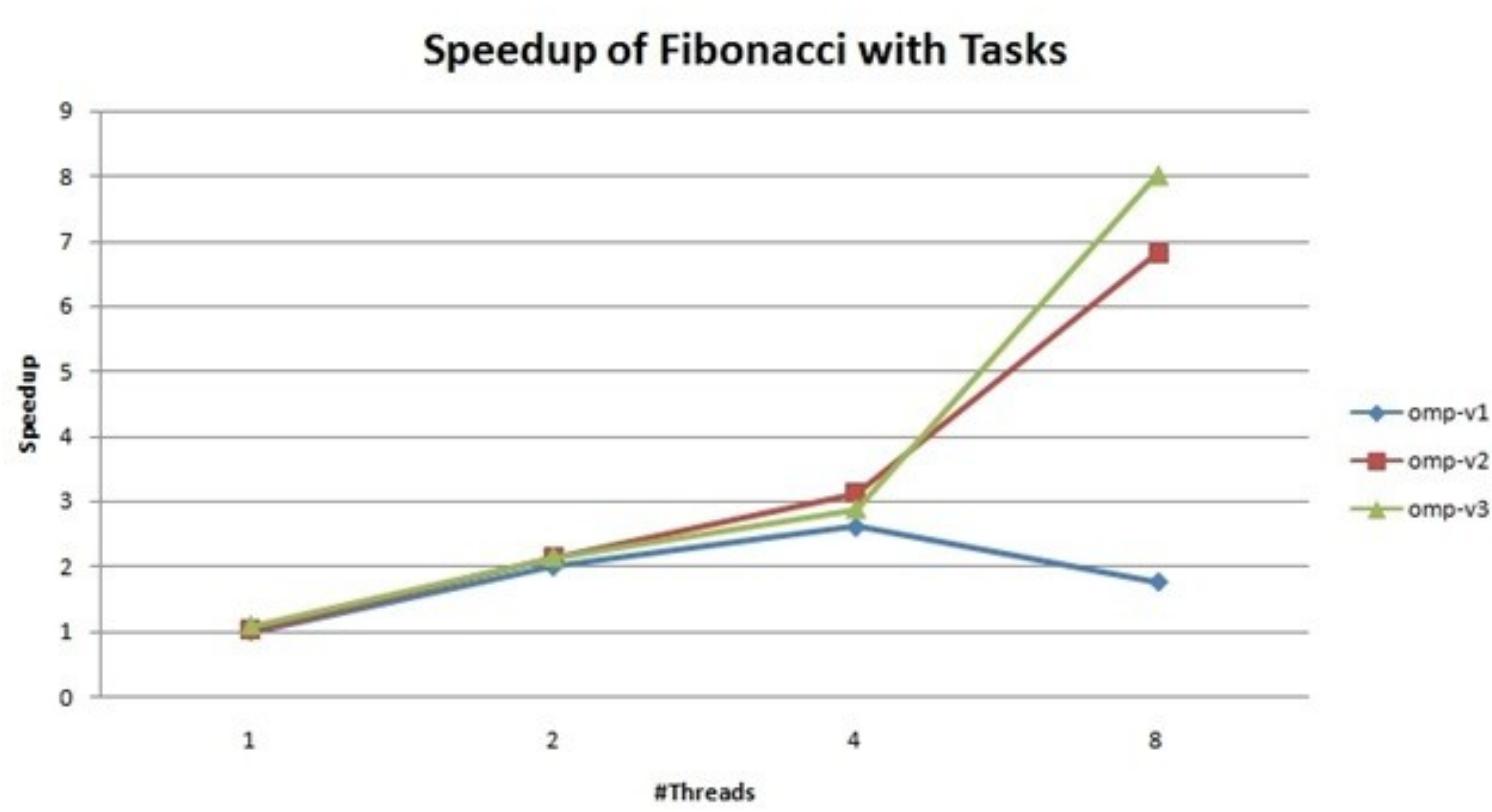
    [ . . . ]
}
```

start of parallel region -
team of worker threads

task generation by one
thread, only!

OpenMP syntax

Results of the Fibonacci program



courtesy: Chr. Terboven, RWTH Aachen

OpenMP syntax

Notes on the Fibonacci speedup results:

- ❑ The simple OpenMP version (omp-v1) doesn't scale – as expected – due to the large amount of tasks generated
- ❑ Improvement 1 (omp-v2):
 - ❑ add an if-clause to the tasks:
`#pragma omp task if(n>=30) shared(...)`
 - ❑ improves the speed-up, but still not perfect
- ❑ Improvement 2 (omp-v3): (see next slide)

OpenMP syntax

version omp-v3 of fib() with tasks:

```
int
fib(int n) {

    int x, y;

    if (n < 2) return n;
    if (n < 30) {
        return(fib(n-1) + fib(n-2));
    }

    #pragma omp task shared(x)
    x = fib(n - 1);
    #pragma omp task shared(y)
    y = fib(n - 2);

    #pragma omp taskwait
    return(x + y);
}
```

OpenMP syntax

Some notes on tasking:

- ❑ tasks allow to exploit parallelism with OpenMP, that hasn't been possible before
- ❑ this makes OpenMP a more powerful parallel programming API
 - ❑ goes beyond the scope in this course
- ❑ a nice intro about tasking in OpenMP: <https://www.openmp.org/events/webinar-how-to-get-openmp-tasks-to-do-your-work/>

OpenMP: Error detection

Tools to check your OpenMP code

OpenMP compile-time checks

- ❑ area of constant improvement, but quality is very much compiler dependent
- ❑ example: code between “#pragma omp for” and the “for(...)" statement
 - ❑ gcc: error, “for statement expected” - OK!
 - ❑ clang: error, “expected expression” - Which???
- ❑ example: “#pragma omp for schedule(run)”
 - ❑ gcc: invalid schedule kind before 'run'
 - ❑ clang: error: expected 'static', 'dynamic', 'guided', 'auto', 'runtime', 'monotonic', 'nonmonotonic' or 'simd' in OpenMP clause 'schedule'

OpenMP: Data Race Detection

There are a number of tools to detect data races in OpenMP programs:

- ❑ Intel Inspector (formerly: ThreadChecker)
- ❑ Oracle Studio Thread Analyzer (outdated)
- ❑ the Thread Sanitizer (tsan), that ships with LLVM
 - ❑ uses a special library, called “Archer” to make the tsan library and OpenMP compatible
 - ❑ developed at RWTH Aachen, Univ. of Utah, and LLNL
 - ❑ can be used with GCC, too (see next slides)

OpenMP: Data Race Detection

Example:

```
int main(int argc, char *argv[]){  
  
    int i, total = 0, N = 2000000;  
    int primes[N];  
  
    #pragma omp parallel for  
  
    for( i = 2; i < N; i++ ) {  
        if ( is_prime(i) ) {  
            primes[total] = i;  
            total++;  
        }  
    }  
  
    printf("# of prime numbers between 2 and %d: %d\n",  
          N, total);  
    return(0);  
}
```



OpenMP: Data Race Detection

Example (cont'd): compile and run

```
$ gcc -g -O3 -o prime prime.c -lm
$ time ./prime
# of prime numbers between 2 and 2000000: 148933
real      0m0.575s
user      0m0.568s
sys       0m0.004s
```



```
$ gcc -g -O3 -fopenmp -o prime prime.c -lm
$ OMP_NUM_THREADS=4 time ./prime
# of prime numbers between 2 and 2000000: 141506
real      0m0.194s <--- speed-up: 2.9x
user      0m0.579s
sys       0m0.005s
```

OpenMP: Data Race Detection

Example (cont'd): run, run, ... and use archer

- ❑ note: the archer module and its commands are DCC specific

```
$ OMP_NUM_THREADS=4 ./prime
# of prime numbers between 2 and 2000000: 139817

$ OMP_NUM_THREADS=4 ./prime
# of prime numbers between 2 and 2000000: 144512

$ module load archer

$ gcc -g -O3 -fopenmp -fsanitize=thread \
-o prime prime.c -lomp -lm
```

OpenMP: Data Race Detection

Example (cont'd): run, run, ... and use archer

- ❑ archer: shows help how to use archer
- ❑ detect_dr: wrapper to run your code with data race detection

```
$ detect_dr ./prime
Archer detected OpenMP application with TSan, supplying
OpenMP synchronization semantics
# of prime numbers between 2 and 2000000: 128202
SUMMARY: ThreadSanitizer: data race prime.c:34 in
         main._omp_fn.0
SUMMARY: ThreadSanitizer: data race prime.c:34 in
         main._omp_fn.0
ThreadSanitizer: reported 2 warnings
Full datarace detection report written
to prime_250105_151926.ddr
```

OpenMP: Data Race Detection

Example (cont'd): run, run, ... and use archer

- ❑ `show_dr`: wrapper that marks the source lines where races happened – run on the `.ddr` file created by `detect_dr`

```
$ show_dr prime_250105_151926.ddr
...
    int primes[N];

#pragma omp parallel for
for( i = 2; i < N; i++ ) {
    if ( is_prime(i) ) {
        primes[total] = i;           <--- here [rww] !
        total++;                  <--- here [w] !
    }
}
...
...
```



OpenMP: Data Race Detection

Example (cont'd): fix the bug

```
int main(int argc, char *argv[]) {  
  
    int i, total = 0, N = 2000000;  
    int primes[N];  
    #pragma omp parallel for  
  
    for( i = 2; i < N; i++ ) {  
        if ( is_prime(i) ) {  
            #pragma omp critical  
            { primes[total] = i;  
              total++;  
            }  
        }  
    }  
    printf("# of prime numbers between 2 and %d: %d\n",  
          N, total);  
    return(0);  
}
```

OpenMP: Data Race Detection

Example (cont'd): check – and recompile

```
$ gcc -g -O3 -fopenmp -fsanitize=thread \
      -o prime prime.c -lomp -lm

$ detect_dr ./prime
Archer detected OpenMP application with TSan,
supplying OpenMP synchronization semantics
# of prime numbers between 2 and 2000000: 148933
No data race report in prime_250105_153439.ddr!
No races detected!

$ gcc -g -O3 -fopenmp -o prime prime.c -lm
$ time OMP_NUM_THREADS=4 ./prime
# of prime numbers between 2 and 2000000: 148933

real 0m0.213s
user 0m0.779s
sys  0m0.019s
```

OpenMP: Data Race Detection

□ Some notes:

- Archer (librarcher.so) is part of LLVM (clang)
- it can be used with GCC, too, linking to the LLVM OpenMP runtime (libomp.so, using -lomp)
- the tools available with ‘module load archer’ are developed by DCC, and still under development
- detect_dr and show_dr are ‘convenience’ wrappers, to ease the use
- for a full documentation of Archer, see e.g. the [Archer Wiki page](#)

OpenMP Scheduling

Controlling the scheduling
of OpenMP threads

OpenMP Scheduling

Load balancing:

- ❑ Important aspect of performance
- ❑ Especially for less regular workloads, e.g.
 - ❑ transposing a matrix
 - ❑ multiplications of triangular matrices
 - ❑ parallel searches in a linked list
- ❑ The **schedule** clause provides different iteration scheduling algorithms for loops

OpenMP Scheduling

The “schedule” clause:

```
#pragma omp for schedule(static[,chunk])  
#pragma omp for schedule(dynamic[,chunk])  
#pragma omp for schedule(guided[,chunk])  
#pragma omp for schedule(auto) - new in 3.0  
#pragma omp for schedule(runtime)
```

- ❑ If there is no schedule clause, the default is static.

OpenMP Scheduling

```
#pragma omp for schedule(static[,chunk])
```

Static schedule:

- ❑ Iterations are divided into pieces of size chunk and then **statically** assigned to the threads.
- ❑ If chunk is not defined, the work (N) is equally divided among the number of threads (P), i.e. $\text{chunk} = N/P$.

OpenMP Scheduling

```
#pragma omp for schedule(dynamic[,chunk])
```

Dynamic schedule:

- ❑ Iterations are divided into pieces of size chunk and then **dynamically** assigned to the threads – i.e. when a thread has finished one chunk, it is assigned a new one.
- ❑ The default chunk size is 1.

OpenMP Scheduling

```
#pragma omp for schedule(guided[, chunk])
```

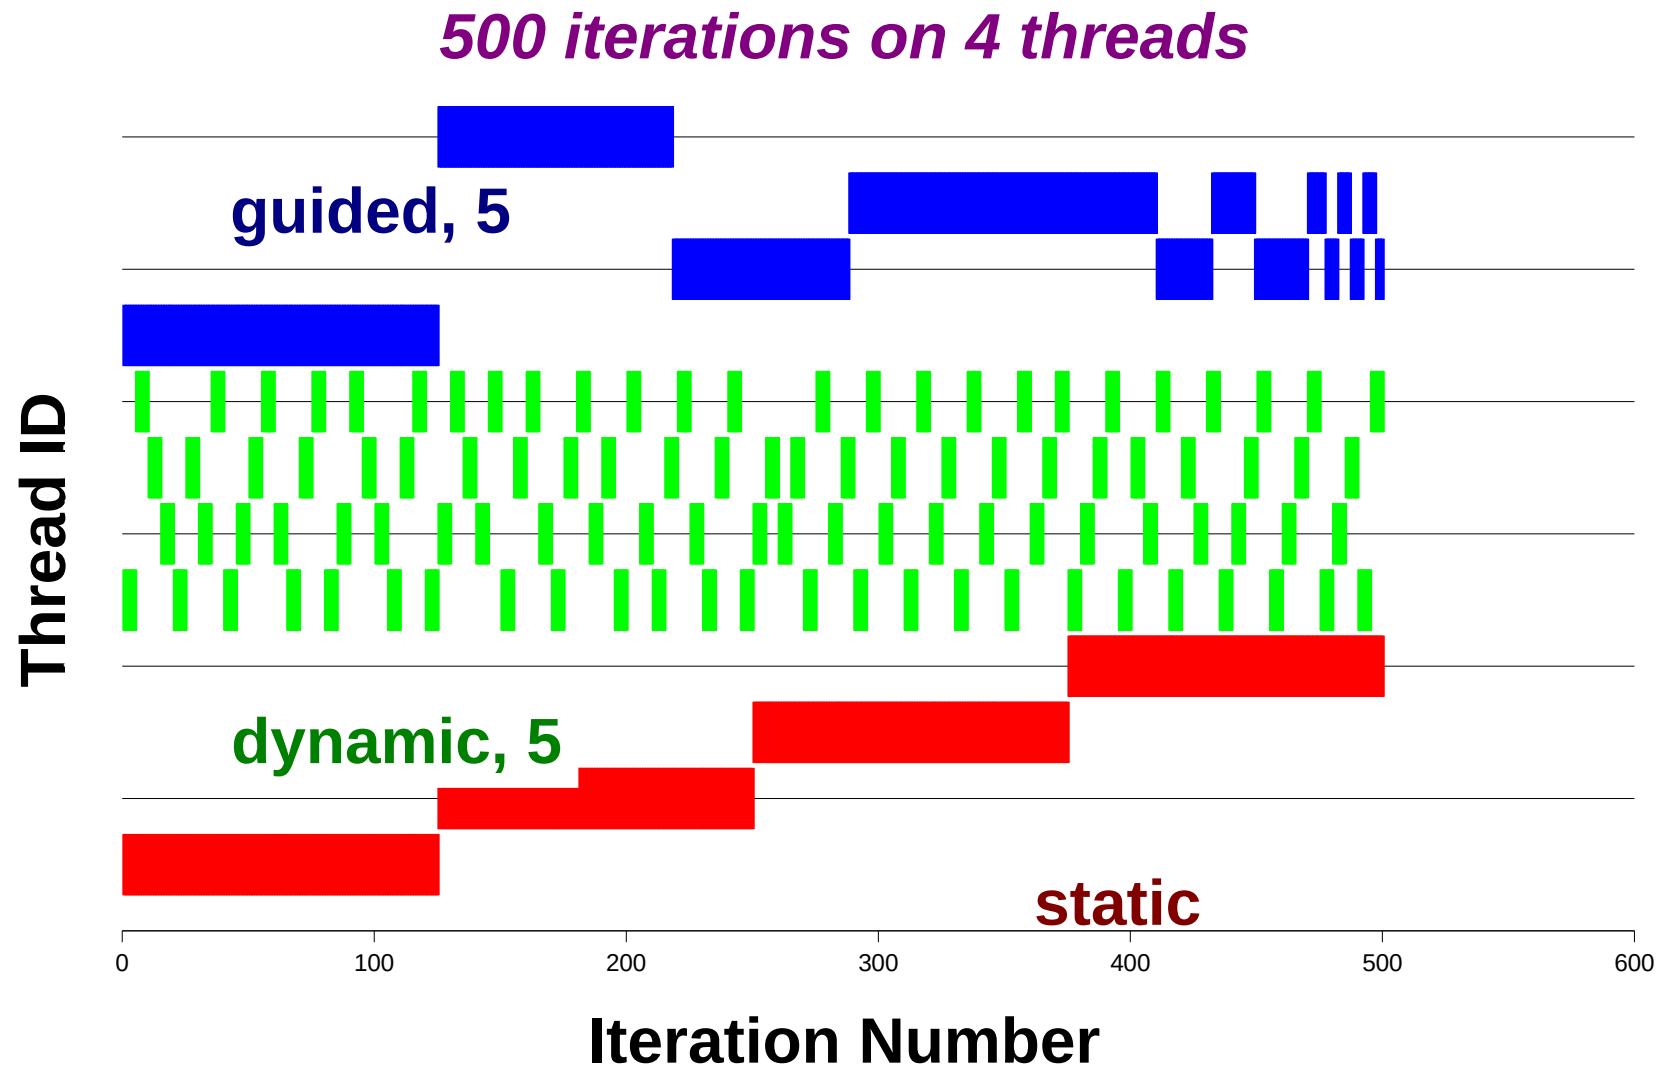
Guided schedule:

- ❑ The chunk size is exponentially reduced with each chunk that gets **dynamically** assigned to the threads; chunk defines the minimum number of iterations to assign each time.

$$\text{chunk} = \text{unass_iter} / (\text{weight} * \text{n_thr})$$

- ❑ The default minimum chunk size is 1.

OpenMP Scheduling



OpenMP Scheduling

```
#pragma omp for schedule(runtime)
```

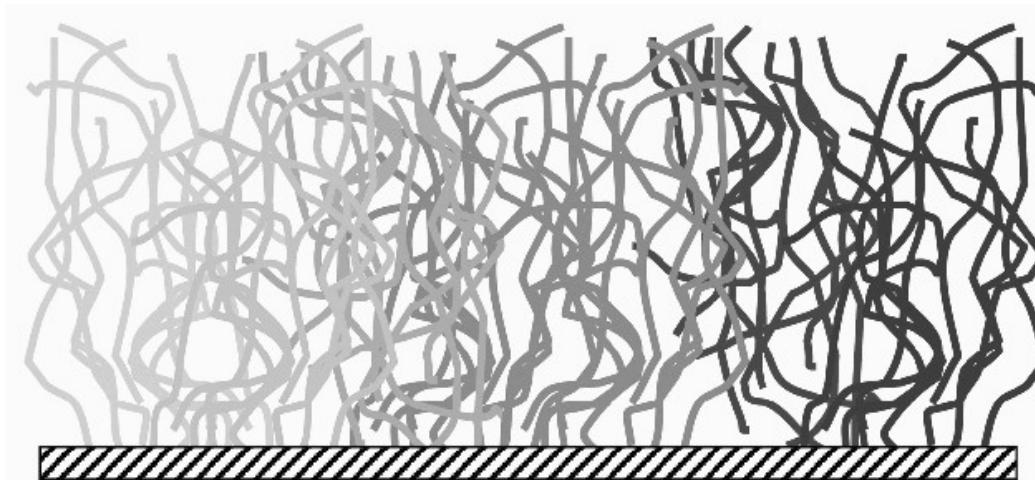
Runtime schedule:

- ❑ The schedule is detected at runtime from the setting of the OMP_SCHEDULE environment variable.
- ❑ Syntax: OMP_SCHEDULE=type,chunk

A real world example: Molecular Dynamics simulation

Example: MD simulation

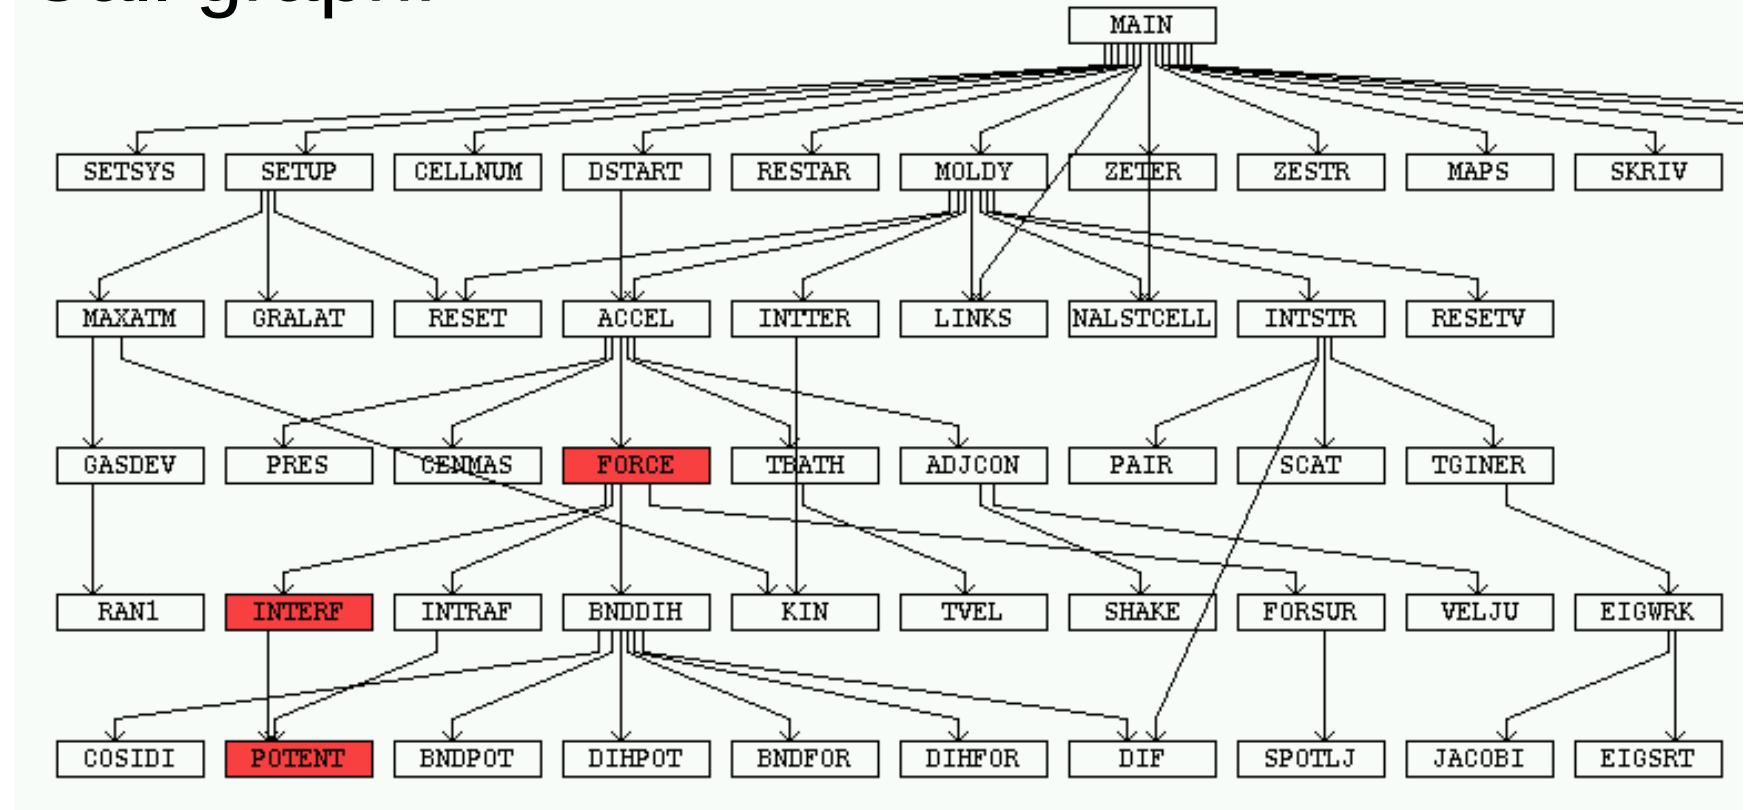
- ❑ Molecular Dynamics simulation of long carbon molecules on a surface:



- ❑ 7200+ lines of Fortran 77 code
- ❑ GOTOs, COMMON blocks, ...
- ❑ one source file

Example: MD simulation

Call graph:



more than 80% of the runtime are spent in
the red part of the call graph

Example: MD simulation

- ❑ The loop to be parallelized contains a call to another subroutine.
- ❑ Data is passed the old Fortran style via COMMON blocks
- ❑ First try: Inserted one PARALLEL DO pragma in the code, using autoscopying, i.e. a feature of the Oracle Studio compiler
- ❑ The compiler generated a parallel version!

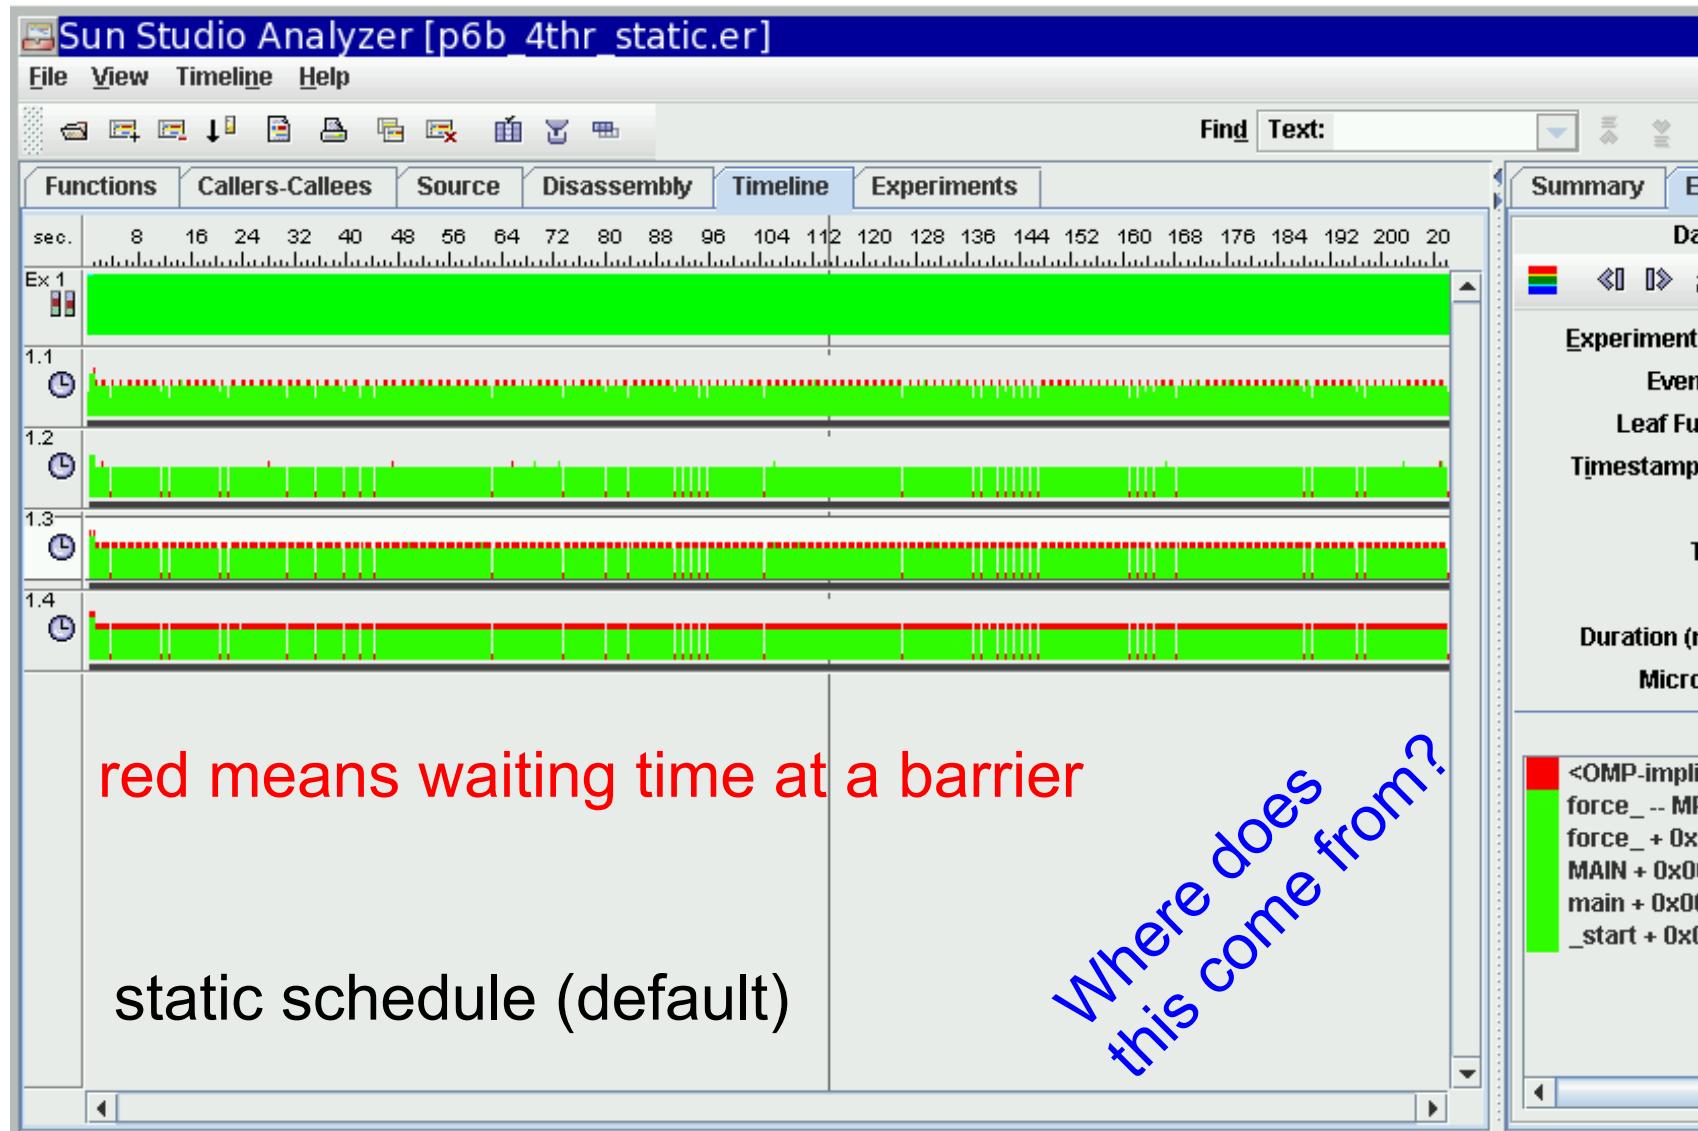
This took us by surprise!

Example: MD simulation

- ❑ First test runs:
 - ❑ It didn't scale ...
 - ❑ The results were dependent on the number of threads ...
- ❑ Thread analyzer revealed data races in two variables inside the called subroutine.
- ❑ Fix: Added additional scoping for those variables in the OpenMP pragma!
- ❑ This solved the data race problem.

Example: MD simulation

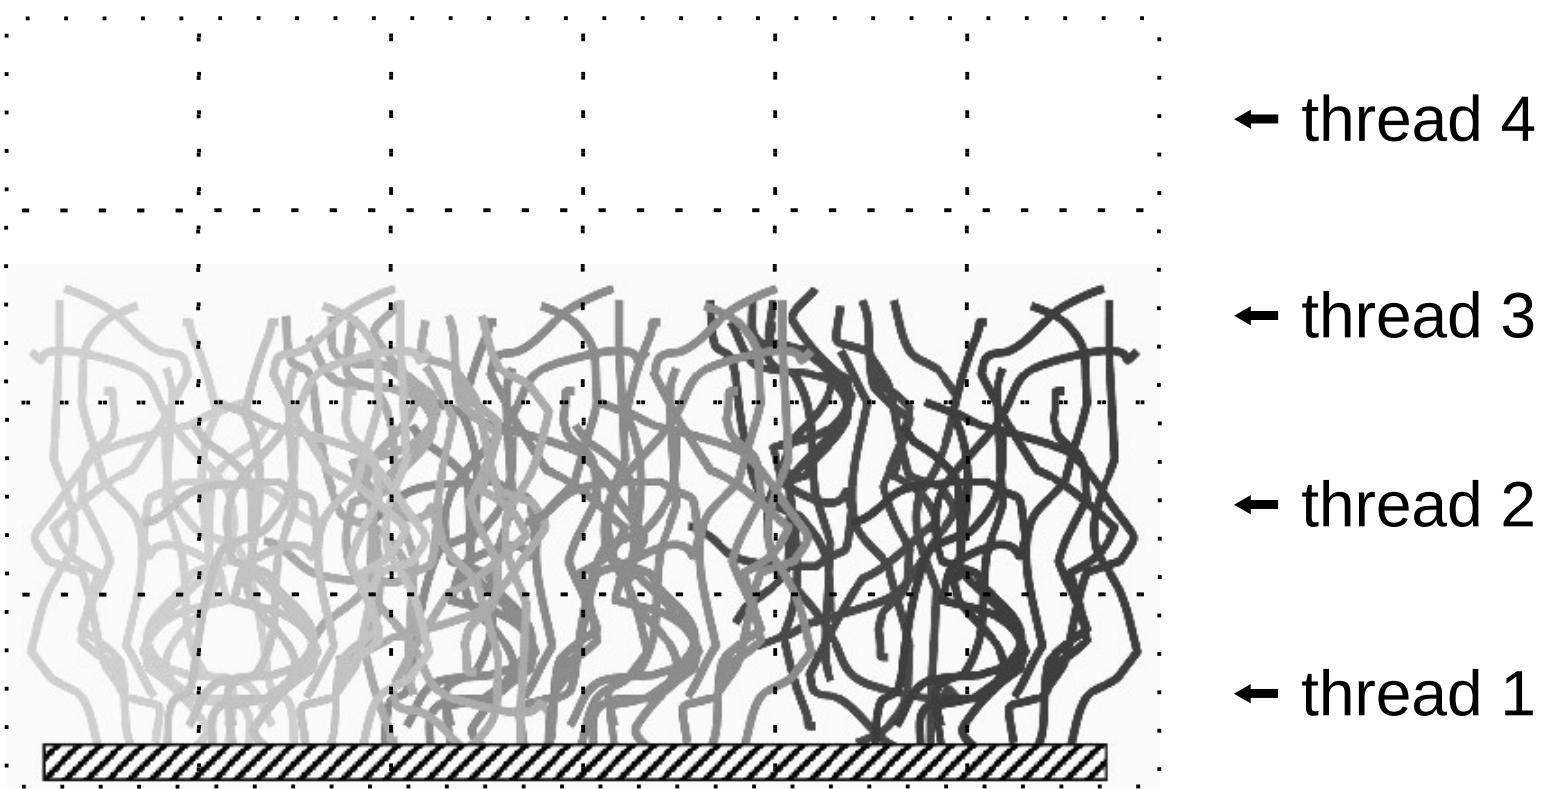
Analysis of the scaling problem:



Example: MD simulation

The simulation box:

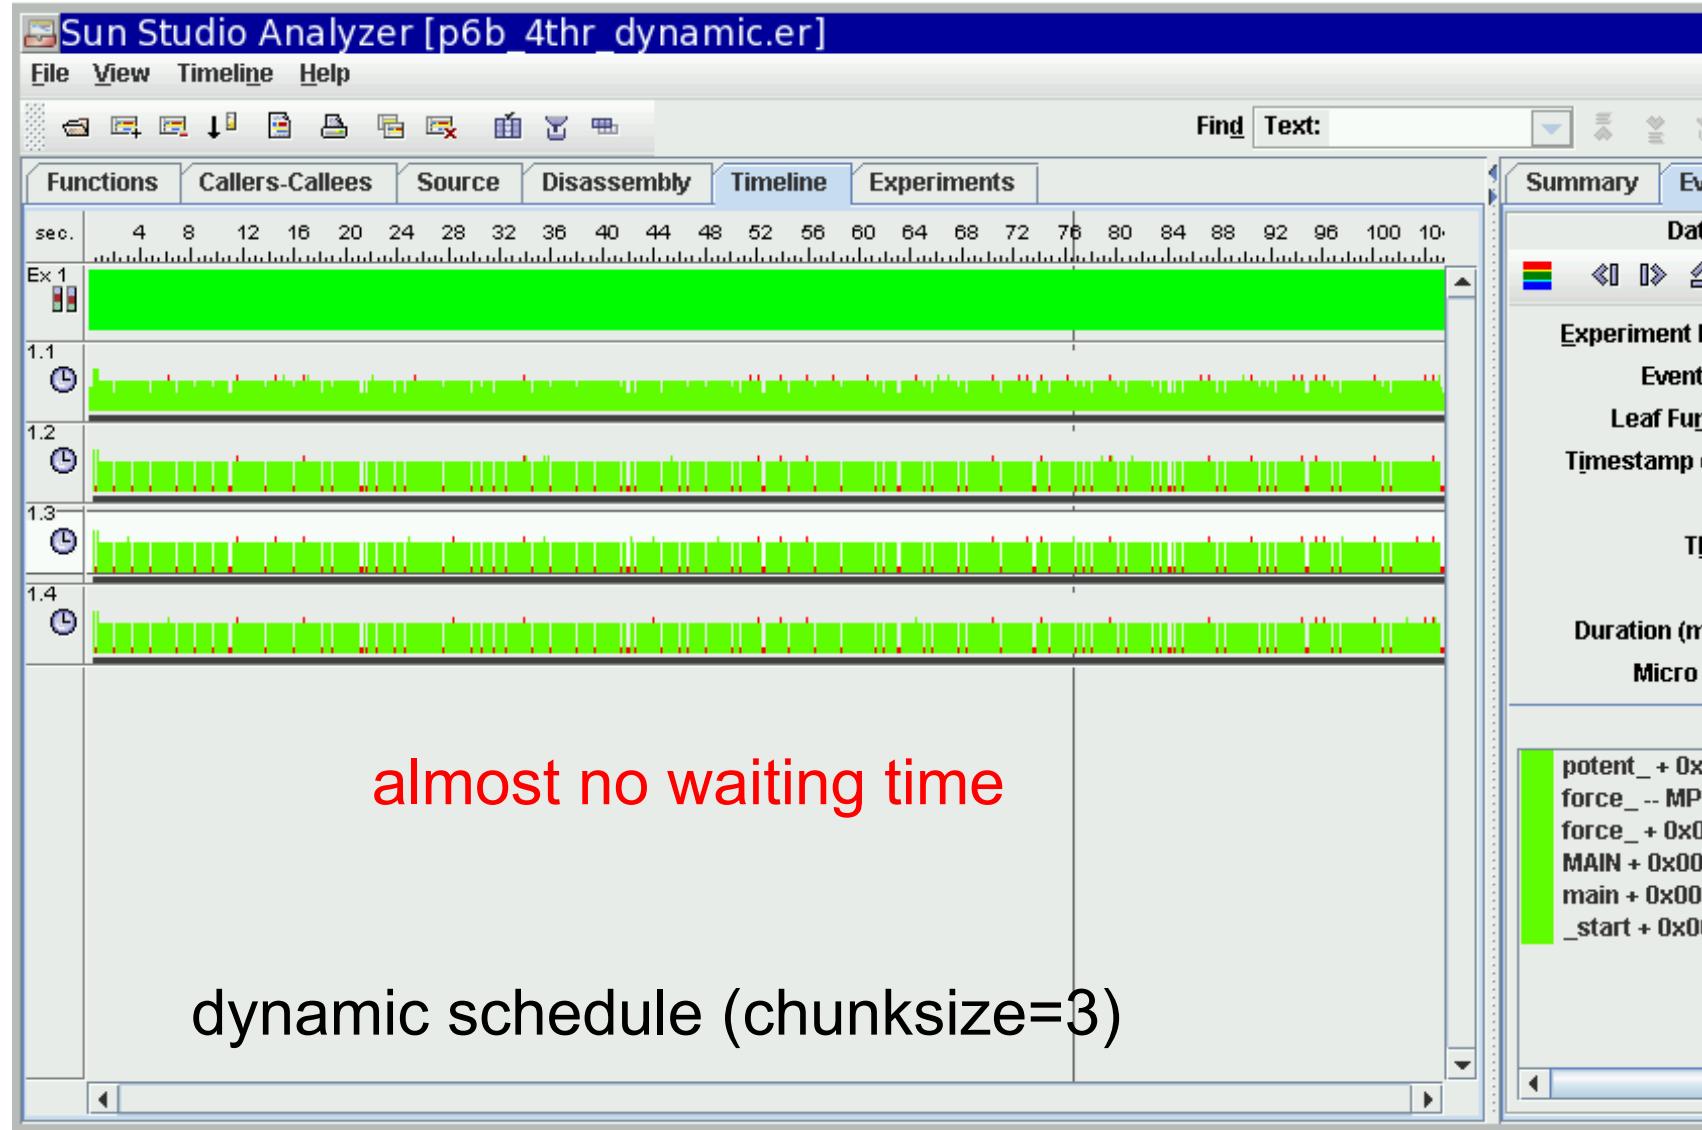
seen from the side



subdivision into smaller cells

Example: MD simulation

Adapted the schedule:



Example: MD simulation

Speed-up results:

