

Parallel Programming in OpenMP – part III

Outline

- ❑ Runtime library
- ❑ Environment variables
- ❑ OpenMP Future
- ❑ Behind the scenes
- ❑ Summary
- ❑ References

OpenMP Runtime Library

The OpenMP runtime library:
support functions

OpenMP Runtime Library

The OpenMP standard defines an API for library calls, that have a variety of functions:

- ❑ query
 - ❑ the number of threads/processors
 - ❑ thread ID, “in parallel”
- ❑ set
 - ❑ the number of threads to use
 - ❑ scheduling mode
- ❑ locking (semaphores)

OpenMP Runtime Library

Name

omp_set_num_threads
omp_get_num_threads
omp_get_max_threads
omp_get_thread_num
omp_get_num_procs
omp_in_parallel

omp_set_dynamic
omp_get_dynamic

omp_set_nested
omp_get_nested

omp_get_wtime
omp_get_wtick

Functionality

set number of threads
get number of threads in team
get max. number of threads
get thread ID
get max. number of processors
check whether in parallel region

activate dynamic thread adjustment
check for dynamic thread adjustment
 (implementation can ignore this)
activate nested parallelism
check for nested parallelism
 (implementation can ignore this)
returns wall clock time
number of second between clock ticks

OpenMP Runtime Library

Function prototypes:

```
void omp_set_num_threads(int num_threads)
int  omp_get_num_threads(void)
int  omp_get_max_threads(void)
int  omp_get_thread_num(void)
int  omp_get_num_procs(void)
int  omp_in_parallel(void)

void omp_set_dynamic(int dynamic_threads)
int  omp_get_dynamic(void)
void omp_set_nested(int nested)
int  omp_get_nested(void)

double omp_get_wtime(void)
double omp_get_wtick(void)
```

OpenMP 3.0 Runtime Library

Name

`omp_set_schedule`
`omp_get_schedule`

Functionality

set the schedule
get the schedule

`omp_get_thread_limit`

max. number of available threads

in the implementation

`omp_set_max_active_levels`

set the number of nested levels

`omp_get_max_active_levels`

get the number of nested levels

`omp_get_level`

returns the current nesting level

`omp_get_ancestor_thread_num`

returns thread id of the ancestor
thread in specified level

`omp_get_team_size`

get team size at specified level

`omp_get_active_level`

returns the number of enclosing,
active nested parallel regions

for more details see the OpenMP 3.0 specifications

OpenMP Runtime Library

- ❑ with the increasing number of features of OpenMP, the number of runtime library functions is growing, too
- ❑ OpenMP 5.x has now more than 60 runtime library functions!
- ❑ check <https://www.openmp.org/specifications/>

OpenMP Runtime Library

Usage of `omp_get_num_threads()` vs `omp_get_max_threads()`:

```
// get the number of threads  
threads = omp_get_max_threads();
```

returns value of OMP_NUM_THREADS

```
// get the number of threads  
threads = omp_get_num_threads();
```

returns 1- outside a parallel region

```
#pragma omp parallel  
{  
#pragma omp master  
{ threads = omp_get_num_threads(); }  
} // end parallel
```

returns value of threads in a parallel region

OpenMP Runtime Library

Measuring time:

- ❑ It is most useful to compare wall clock times

```
double ts, te;  
ts = omp_get_wtime();  
  
do_work();  
  
te = omp_get_wtime() - ts;  
  
printf("Elapsed time: %lf\n", te);
```

- ❑ `clock()` returns the accumulated CPU time of all threads!

OpenMP Environment Variables

Controlling OpenMP
via Environment Variables

OpenMP Environment Variables

- ❑ **OMP_NUM_THREADS = n**
 - ❑ sets the max. no of threads to n
- ❑ **OMP_SCHEDULE = schedule[,chunk]**
 - ❑ schedule: [static | guided | dynamic]
 - ❑ chunk: size of chunks
 - ❑ Note: applies to parallel do/for loops only!
 - ❑ Caveat: GCC sets default ‘dynamic,1’
- ❑ **OMP_DYNAMIC = [TRUE | FALSE]**
 - ❑ allows runtime system to adjust settings

OpenMP Environment Variables

- ❑ **OMP_STACKSIZE** = size[B|K|M|G]
 - ❑ sets the size of the stack of OpenMP threads
 - ❑ default unit: Kilobytes
- ❑ **OMP_WAIT_POLICY** = active|passive
 - ❑ controls the behaviour of idle threads
 - ❑ active: “spinning threads”, i.e. use cycles
 - ❑ passive: threads go to sleep
 - ❑ the default is implementation dependent

OpenMP Environment Variables

- ❑ **OMP_PROC_BIND** = [true|false|close|spread]
 - ❑ controls the binding of threads to cores
 - ❑ gives a hint if this should be packed or spread out over the system
- ❑ **OMP_PLACES** = [cores|sockets|<list>]
 - ❑ controls the placement of threads
 - ❑ cores: place across cores
 - ❑ sockets: place on whole sockets
 - ❑ or provide a list with core numbers
 - ❑ works in combination with binding!

OpenMP Environment Variables

- ❑ `OMP_NESTED` = [TRUE | FALSE]
- ❑ `OMP_MAX_ACTIVE_LEVELS` = n
 - ❑ controls the max. level for nested parallelism
- ❑ `OMP_THREAD_LIMIT` = n
 - ❑ sets the maximum number of threads for an OpenMP program

OpenMP Environment Variables

Notes:

- ❑ The defaults are depended on the compiler and runtime environment used.
- ❑ You can use `OMP_DISPLAY_ENV=true` to show the settings at the startup of your program. This is useful to check for differences in runtime implementations!
- ❑ On the DTU HPC systems, we set `OMP_NUM_THREADS=1` as a default, i.e. you need to adjust the value to your needs!

OpenMP Precedence

- ❑ Level of priority:
 - 1 clauses, e.g. num_threads(...)
 - 2 library calls, e.g. omp_set_num_threads(...)
 - 3 environment variables, e.g.
OMP_NUM_THREADS
- ❑ For a detailed discussion see the OpenMP specifications or check the documentation of your OpenMP implementation.

OpenMP Features

OpenMP development
and standard extensions

OpenMP Features

- New features are discussed in the OpenMP ARB and the community, and made or make it into the standard, e.g. extensions for
 - better performance
 - memory placement (4.0)
 - debugging
 - checks, both at compile- and run-time
 - exception handling (4.0)
 - access to accelerators (e.g. GPUs) (4.0)
 - ...

OpenMP: Behind the scenes

What the compiler does
with your code

OpenMP: Behind the scenes

```
#define MAX_SIZE 8000000
int main() {
    double GlobSum; /* A global variable */
    double array[MAX_SIZE];
    int nthreads;
    int i;
    /* Initialize things */
    for (i=0; i<MAX_SIZE; i++) array[i] = i;
    GlobSum = 0;
    nthreads = omp_get_max_threads();
    printf("Threads: %d\n", nthreads );
    #pragma omp parallel for private(i) \
        reduction(+ : GlobSum)
    for(i=0; i<MAX_SIZE;i++)
        GlobSum = GlobSum + array[i];

    return(EXIT_SUCCESS);
}
```

OpenMP: Behind the scenes

- ❑ Used the OMPI compiler to generate the intermediate code shown on the next slides.
- ❑ The actual implementation differs from compiler to compiler, and probably also from version to version (improvements).

OpenMP: Behind the scenes

```
int main() {  
    ...  
    int      i;  
    _omp_initialize();  
  
    for (i = 0; i < 8000000; i++) array[i] = i;  
    GlobSum = 0;  
    nthreads = omp_get_max_threads();  
    printf("Threads: %d\n", nthreads);  
  
/* #pragma omp parallel for private(i) reduction(+: GlobSum) */  
{  
    _OMP_PARALLEL_DECL_VARSTRUCT(main_parallel_0);  
    _OMP_PARALLEL_INIT_VAR(main_parallel_0, GlobSum);  
    _OMP_PARALLEL_INIT_VAR(main_parallel_0, array);  
    _omp_create_team((-1), _OMP_THREAD, main_parallel_0,  
                     (void *) &main_parallel_0_var); /* create team of  
                                         * threads */  
    _omp_destroy_team(_OMP_THREAD->parent);  
}  
  
    return 0;  
}
```

OpenMP: Behind the scenes

```
void *main_parallel_0(void *_omp_thread_data) {
    int      _omp_dummy = _omp_assign_key(_omp_thread_data);
    double   (*array)[8000000] = &_OMP_VARREF(main_parallel_0,array);
{
    int      i;
    double  GlobSum = 0;
    int      _omp_start, _omp_end, _omp_incr, _omp_last_iter = 0;
    int      _omp_for_id = _omp_module.for_ofs + 0;
    int      (*_omp_sched_bounds_func)(int, int, int, int,
                                         int, int *, int *, int, int *) ;
    /* static with chunksize or runtime */
    int      _omp_init_start, _omp_nchunks, _omp_c = 0,
            _omp_chunksize;
    _omp_incr = (1);
    _omp_init_directive(_OMP_FOR, _omp_for_id, 0,
                        _omp_incr, 0, 115);
    _omp_sched_bounds_func = _omp_static_bounds;
    _omp_static_bounds_default(8000000, 0, _omp_incr,
                               &_omp_start, &_omp_end);
    ...
}
```

OpenMP: Behind the scenes

...

```
while (((*_omp_sched_bounds_func) (8000000, 0, _omp_for_id,
    _omp_incr, -1, &_omp_start, &_omp_end, 1, 0, &_omp_c)) {
    if (_omp_start < (8000000) && _omp_end == (8000000))
        _omp_last_iter = 1;

    for (i = _omp_start; i < _omp_end; i++) {
        GlobSum = GlobSum + (*array)[i];
    } /* for */
}

if (_omp_last_iter) { /* lastprivate assignments */ }

/* reduction operation (+:GlobSum) */
othread_set_lock(&_omp_module.reduction_lock[0]);
_OMP_VARREF(main_parallel_0, GlobSum) += GlobSum;
othread_unset_lock(&_omp_module.reduction_lock[0]);
}
return 0;
}
```

OpenMP vs POSIX threads

A possible POSIX threads solution:

```
main() {  
    int i,retval;  
    pthread_t tid;  
  
    /* Initialize things */  
    pthread_attr_init(&attr);  
    pthread_mutex_init (&my_mutex, NULL);  
    pthread_attr_setscope(&attr, PTHREAD_SCOPE_SYSTEM);  
  
    for (i=0; i<MAX_SIZE; i++) array[i] = i;  
    GlobSum = 0;  
  
    for(i=0;i<ThreadCount;i++) {  
        index[i] = i;  
        retval = pthread_create(&tid,&attr,SumFunc,  
                               (void *)index[i]);  
        thread_id[i] = tid;  
    }  
    for(i=0;i<ThreadCount;i++)  
        retval = pthread_join(thread_id[i],NULL);  
}
```



OpenMP vs POSIX threads

```
void *SumFunc(void *parm) {
    int i,me,chunk,start,end;
    double LocSum;

    /* Decide which iterations belong to me */
    me = (int) parm;
    chunk = MAX_SIZE / ThreadCount;
    start = me * chunk;
    end = start + chunk; /* C-Style - actual element + 1 */
    if ( me == (ThreadCount-1) ) end = MAX_SIZE;

    /* Compute sum of our subset*/
    LocSum = 0;
    for(i=start;i<end;i++) LocSum = LocSum + array[i];

    /* Update the global sum and return */
    pthread_mutex_lock (&my_mutex);
    GlobSum = GlobSum + LocSum;
    pthread_mutex_unlock (&my_mutex);
}
```



Note: Variable definitions are omitted in this example!

OpenMP Summary

Short summary
of the three lectures

OpenMP Summary

- ❑ OpenMP: a parallel programming model for multi-core computers
- ❑ compiler directives, support functions, environment variables
- ❑ easy to implement, also “little by little”

- ❑ next lecture: “OpenMP & Performance”
 - ❑ special guest: Ruud van der Pas, Oracle

OpenMP References

- ❑ Useful Websites:
 - ❑ <http://www.openmp.org/>
 - ❑ check for webinars and tutorials
- ❑ Tutorial from LLNL:
 - ❑ <https://hpc.llnl.gov/documentation/tutorials>
 - ❑ OpenMP: <https://hpc-tutorials.llnl.gov/openmp/>
- ❑ OpenMP specifications:
 - ❑ <https://www.openmp.org/specifications/>
 - ❑ C/C++ reference card for OpenMP 4.5
 - ❑ FORTRAN reference card for OpenMP 4.5