

GPU Programming with OpenMP

Part1: OpenMP offload basics



Hans Henrik Brandenborg Sørensen

DTU Computing Center

DTU Compute

<hhbs@dtu.dk>

$$f(x+\Delta x) = \sum_{i=0}^{\infty} \frac{(\Delta x)^i}{i!} f^{(i)}(x)$$

$\Theta^{\sqrt{17}} + \Omega \int \delta e^{i\pi} =$

$\epsilon^{\infty} = \{2.7182818284590452353602874713526624977572470636231870738$

$\Sigma!$




Overview

- GPU programming languages
- GPU programming model
- OpenMP offload basics
 - Target constructs
 - Work-sharing constructs
 - Synchronization
- Runtime library API / Environment variables

Many ways to program GPUs



- CUDA (2007)
- OpenCL (2009)
- OpenACC (2012)
- OpenMP 4.0 (2013)
- Others (ROCM, Numba, oneAPI, Matlab,...)



What is CUDA?



- [Compute Unified Device Architecture]
- A parallel computing standard and API proposed by NVIDIA for general-purpose computations on CUDA-enabled GPUs
 - Priority #1: Make things easy (Sell GPUs)
 - Priority #2: Get performance
- Result: Low level – but mainly C++ syntax
 - Requires expert knowledge to get best performance
- Scalable
- Well documented and free to use (!)

What is OpenCL?



- **Open Computing Language** (current v3.0)
- Khronos group (non-profit organization):
 - “*OpenCL is an open, royalty-free standard for cross-platform, parallel programming of modern processors found in parallel computers, servers and handheld/embedded devices.*”
- Open standard for heterogeneous computing
- Priority #1: Become the *industry-wide future standard* for heterogeneous computing
- Priority #2: Use all computational resources in the system efficiently
- Up to vendors to provide support! **K H R O N O S**
GROUP

CUDA vs OpenCL?

- Most CUDA features map one-to-one to OpenCL features (only the syntax is different)
- CUDA comes with a mature software framework
- CUDA comes with tuned high-performance libs
 - cuBLAS – CUDA Basic Linear Algebra Subroutines library
 - cuFFT – CUDA Fast Fourier Transform library
 - cuSPARSE – CUDA Sparse Matrix library and many more!
- OpenCL have had much less effort in this direction
- CUDA is well documented (by NVIDIA)
- NVIDIA products are widely used in HPC (>90%)
 - OpenCL still lags in performance for NVIDIA products

What is OpenACC?



- OpenACC is an open specification for compiler directives for parallel programming
`#pragma acc directive [clause]`
- Developed by PGI, Cray, CAPS, and NVIDIA
- High-level directives: Minimal modifications to the code, fewer lines than with CUDA, OpenCL,..
- Supports CPUs, GPU accelerators and co-processors from multiple vendors
- Compiler support: gcc, nvc++ (pgi), cc, clacc

OpenMP for GPUs



- OpenMP is an open specification for compiler directives for parallel programming

```
#pragma omp target ... [clause]
```

- OpenMP 4.0 – 6.0: Standardizes established practice for heterogeneous device programming
- Support in all common compilers
- Behind the scenes: Intermediate level of CUDA or OpenCL is typically used for GPUs

OpenMP vs OpenACC?

- OpenMP is the established ‘de-facto’ standard
 - You know the directives and terminology from week 2
 - Let’s not switch and complicate things...
- PGI OpenACC was known for good performance
 - Few years ago: OpenACC + `pgi` compiler had generally better performance than OpenMP + `gcc` / `clang`!
 - Now: OpenMP + `nvc++` gives similar high performance
- Why not CUDA?
 - Low level (!) – week 3 tends to be debug, debug..
 - If you want to learn CUDA please sign up for the **CUDA special course** that will run in the spring semester 2026

NVIDIA HPC compiler

- The NVIDIA HPC C++ compiler is called
 - nvcc++ [options] [path] filename [...]
- Many options are common with gcc: -g -fast ..
- Compiles for OpenMP, OpenACC, and CUDA
 - ❑ -mp=gpu
 - ❑ -acc (can be used together)
 - ❑ -cuda
- If you run into linking problems (undefined refs.)
 - ❑ Avoid mixing .c and .cpp files (.cpp and .cu is fine)
 - ❑ Or use `extern "C" { ... }` appropriately


NVIDIA HPC compiler

- We provide a `Makefile` template for exercises!
- Most important compiler flags:
 - `-gpu=cc90`
 - Compile code for compute capability 9.0 (Hopper)
 - Default is cc. 1.0 (Tesla), latest is cc. 9.x (Hopper)
 - `-Minfo`
 - Set output comments from compiler to verbose
 - `-gpu=lineinfo`
 - Generate line-number information for device code (e.g., used in NsightTM profiler)

<https://docs.nvidia.com/hpc-sdk/compilers/hpc-compilers-user-guide>

GPU programming model

GPU programming model




■ Host – the CPU

- In charge, manages resources
- Runs main(), etc.

■ Device – the GPU


- Co-processor / accelerator
- Runs specific tasks

GPU programming model




- ① Data CPU → GPU

GPU programming model




- ① Data CPU → GPU
- ② Data GPU → CPU

GPU programming model



- ① Data CPU → GPU
- ② Data GPU → CPU
- ③ Allocate GPU memory

GPU programming model



- ① Data CPU → GPU
- ② Data GPU → CPU
- ③ Allocate GPU memory
- ④ Offload code to device

OpenMP offload basics

OpenMP offload basics

■ Syntax C/C++:

```
#pragma omp target [clause]
{
    ...
}
```

■ Clause can be

- ❑ private(list)
- ❑ firstprivate(list)
- ❑ if([target:] scalar_expr)
- ❑ nowait
- ❑ depend([modifier,] type: list)
- ❑ allocate([allocator:] list)
- ❑ in_reduction(op: list)
- ❑ thread_limit(int_expr)
- ❑ device([modifier:] int_expr)
- ❑ map([{alloc | to | from | tofrom}:] list)
- ❑ defaultmap(behavior[:var-category])
- ❑ is_device_ptr(list)
- ❑ has_device_addr(list)

OpenMP offload basics

■ Syntax C/C++:

```
#pragma omp target [clause]  
{  
    ...  
}
```

Not currently supported in nvcc++

■ Clause can be

- private(list)
- firstprivate(list)
- if([target:] scalar_expr)
- nowait
- depend([modifier,] type: list)
- ~~allocate([allocator.] list)~~
- ~~in_reduction(op: list)~~
- thread_limit(int_expr)
- device([modifier:] int_expr)
- map([{alloc | to | from | tofrom}:] list)
- defaultmap(behavior[:var-category])
- is_device_ptr(list)
- has_device_addr(list)

OpenMP offload basics

■ First OpenMP offload “Hello world”:

```
#include <stdio.h>

int main(int argc, char *argv[] ) {
    #pragma omp target
    {
        printf("Hello world from the device!\n");
    } // end target
    return(0);
}
```

OpenMP offload basics

```
$ nvcc++ -mp=gpu -Minfo -o hello hello.cpp
main:
    11, #omp target
        11, Generating "nvkernel_main_F1L11_2" GPU
kernel

$ ./hello
Hello world from the device!
```

- ❑ Note: The compiler generates a (CUDA) kernel to run on the GPU with a specific name starting with nvkernel_[function name]_
- ❑ The programmer need to explicitly create parallelism on the device – target is not enough

OpenMP offload basics

■ Syntax C/C++:

```
#pragma omp target teams [clause]
{
    ...
}
```

■ Extra clauses with teams are

- ❑ shared(list)
- ❑ default (behavior)
- ❑ reduction (op: list)
- ❑ num_teams (int_expr)

(implementation defined
limit = 65536 teams)

OpenMP offload basics

■ Second OpenMP offload “Hello world”:

```
#include <stdio.h>

int main(int argc, char *argv[]) {
    #pragma omp target teams
    {
        printf("Hello world from the device!\n");
    } // end target teams
    return(0);
}
```

OpenMP offload basics

```
$ ./hello
Hello world from the device!
Hello world from the device!
...
Hello world from the device!
```

```
$ ./hello | wc -l
114
```

- ❑ By default `construct target teams` creates as many teams as there are compute units (SMs) on the device
- ❑ The teams are scheduled to run independently
- ❑ Only a single thread in each team executes the code

OpenMP offload basics

■ Syntax C/C++:

```
#pragma omp target parallel [clause]
{
    ...
}
```

■ Extra clauses with parallel are

- proc_bind(master | close | spread)

OpenMP offload basics

■ Third OpenMP offload “Hello world”:

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[]) {
    #pragma omp target parallel
    {
        printf("Hello world from %d!\n",
               omp_get_thread_num());
    } // end target parallel
    return(0);
}
```

OpenMP offload basics

```
$ ./hello
Hello world from 4960!
Hello world from 4961!
...
Hello world from 10944!
```

```
$ ./hello | wc -l
14592
```

- ❑ Note: The order of execution will be different from run to run (but in “groups” of 32)!
- ❑ The default no. of threads depends on the OpenMP implementation and the hardware
- ❑ Here it starts $114 * 4 * 32 = 14592$ threads in 1 team
 - [CUDA] number of threads in a warp!

OpenMP offload basics

■ Syntax C/C++:

```
#pragma omp target teams parallel [clause]
{
    ...
}
```

■ No extra clauses

OpenMP offload basics

■ Final OpenMP offload “Hello world”:

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[]) {
    #pragma omp target teams parallel \
        num_teams(114) thread_limit(4*32)
    {
        printf("Hello world from (%d, %d) !\n",
               omp_get_team_num(),
               omp_get_thread_num());
    } // end target teams parallel
    return(0);
}
```

OpenMP offload basics

```
$ nvc++ -mp=gpu -Minfo -o hello hello.cpp
main:
  22, #omp target teams num_teams(114) thread_limit(128)
  22, Generating "nvkernel_main_F1L22_2" GPU kernel
        #omp parallel
```

```
$ ./hello
Hello world from (4, 64) !
Hello world from (4, 65) !...
Hello world from (74, 31) !
```

```
$ ./hello | wc -l
14592
```

- ❑ If we did not specify anything it would start 114 teams of 992 threads = $114 * 31 * 32 = 113088$ threads



[CUDA] just below the maximum number of blocks per SM!

OpenMP offload basics

Work-sharing constructs – loop:

- Syntax C/C++:

```
#pragma omp target teams loop [clause]  
for-loop
```

- **loop** asserts the ability of a loop to be run in any order, finding the available parallelism
- Clause can be any of the clauses accepted by the **target**, **teams** or **loop** directives with identical meanings and restrictions

OpenMP offload basics

Work-sharing constructs – loop:

- `omp target teams loop`
 - Recommended way
 - Can use `num_teams` and `thread_limit` clauses
- `omp target loop`
 - Fully automatic
 - No `num_teams` and `thread_limit` clauses
- `omp target parallel loop`
 - Uses only threads, and doesn't use teams
 - Might be useful for light computation

OpenMP offload basics

Work-sharing constructs – distribute:

- Syntax C/C++:

```
#pragma omp target teams distribute \
parallel for [clause]
```

for-loop

- The iterations are distributed and executed in parallel by all threads of the teams
- Clause can be any of the clauses accepted by the target, teams, distribute or parallel for directives with identical meanings and restrictions

OpenMP offload basics

Work-sharing constructs – distribute:

```
#define N 16

int main(int argc, char *argv[]) {
    double a[N], b[N], c[N];
    for (int i = 0; i < N; i++)
        a[i] = b[i] = i * 1.0;

#pragma omp target teams \
    distribute parallel for
for (int i = 0; i < N; i++)
    c[i] = a[i] + b[i];

for (int i = 0; i < N; i++)
    printf("%f\n", c[i]);
}
```

❑ Will this run?

- Yes! – OpenMP 6.0:
“implicit data-mapping rules”

❑ Is it efficient?

- No!

OpenMP offload basics

■ Synchronization – remember for last week:

- ❑ `#pragma omp single`
- ❑ `#pragma omp critical`
- ❑ `#pragma omp atomic`
- ❑ `#pragma omp ordered`
- ❑ `#pragma omp barrier`

■ Implied barriers

- ❑ exit from parallel region
- ❑ exit from `omp for/omp do/omp workshare`
- ❑ exit from sections
- ❑ exit from `single`

OpenMP 6.0 specification:
“The binding thread set for
a barrier region is the
current **team**.”

OpenMP offload basics

■ Synchronization – remember for last week:

- ❑ `#pragma omp single`
- ❑ ~~`#pragma omp critical`~~
- ❑ `#pragma omp atomic`
- ❑ ~~`#pragma omp ordered`~~
- ❑ `#pragma omp barrier`

Not currently supported in nvcc++

■ Implied barriers

- ❑ exit from parallel region
- ❑ exit from `omp for/omp do/omp workshare`
- ❑ ~~exit from sections~~
- ❑ exit from `single`

OpenMP 6.0 specification:
“The binding thread set for a barrier region is the current team.”

OpenMP offload basics

■ What does this mean in practice

working host version

```
int count = 0;  
#pragma omp parallel \  
    num_threads(16) \  
    shared(count)  
{  
    #pragma omp atomic  
    count += 1;  
  
    // Wait for all threads done  
    #pragma omp barrier  
  
    #pragma omp master  
    printf("# of threads is %d\n",  
          count);  
}
```

incorrect device version

```
int count = 0;  
#pragma omp target teams parallel \  
    num_teams(114) thread_limit(64) \  
    map(tofrom:count)  
{  
    #pragma omp atomic  
    count += 1;  
  
    // Wait for all threads done  
    #pragma omp barrier  
  
    #pragma omp master  
    printf("# of threads is %d\n",  
          count);  
}
```

OpenMP offload basics

■ What does this mean in practice

```
$ ./threads  
# of threads is 16
```

```
$ ./threads  
# of threads is 16
```

```
$ ./threads  
# of threads is 16
```

```
$ ./threads_offload  
# of threads is 5925  
# of threads is 5945  
[... 111 ...]  
# of threads is 6319
```

```
$ ./threads_offload  
# of threads is 5886  
# of threads is 5904  
[... 111 ...]  
# of threads is 6319
```

OpenMP offload basics

■ What does this mean in practice

working device version

```
int count = 0;  
#pragma omp target teams parallel \  
    num_teams(114) thread_limit(64)\ \  
    map(tofrom:count)  
{  
    #pragma omp atomic  
    count += 1;  
}  
// To wait for all threads done we  
use two separate offload regions  
#pragma omp target map(to:count)  
{  
    printf("# of threads is %d\n",  
        count);  
}
```

```
$ ./threads_offload  
# of threads is 7296
```

```
$ ./threads_offload  
# of threads is 7296
```

- ❑ Synchronization among all threads requires two separate offload regions


OpenMP offload basics summary




#pragma omp target ← Execute code on device



OpenMP offload basics summary



OpenMP offload basics summary



OpenMP runtime library

OpenMP offload runtime library



■ New library routines

Name

`int omp_get_num_teams(void)`

`int omp_get_team_num(void)`

`int omp_get_num_devices(void)`

`void omp_set_default_device(
int dev_num)`

`int omp_get_default_device(void)`

`int omp_get_initial_device(void)`

`int omp_is_initial_device(void)`

`int omp_get_num_procs(void)`

Functionality

get number of teams

get team number

get number of devices

set the default device

get the default device

get initial device (=host)

are we on the host?

get number of processors

for more details see the OpenMP 5.0 specifications (<https://www.openmp.org/specifications/>)

OpenMP offload runtime library



■ New library routines

Name

```
void* omp_target_alloc(size_t,  
int dev_num)
```

```
void omp_target_free(void*,  
int dev_num)
```

```
int omp_target_memcpy(...,  
int dev_num)
```

```
int omp_target_memcpy_rect(...,  
int dev_num)
```

```
omp_target_associate_ptr(...)
```

```
omp_target_disassociate_ptr(...)
```

```
omp_target_is_present(...)
```

Functionality

allocate memory on device

free memory on device

memcpy to and from device

memcpy to and from device of
a rectangular subvolume

combining device ptr with
host ptr to be used in map
clause

(we use `is_device_ptr` clause)

Offload environment variables

- `OMP_DEFAULT_DEVICE = n`
 - Sets the default device when “device(n)” clause is not specified (error if $n \geq$ no. of devices)
- `OMP_NUM_TEAMS = n`
 - Maximum number of `num_teams` allowed
- `NVCOMPILER_OMP_CUDA_GRID =teams,limit`
 - Force a specific `num_teams` and `thread_limit`
- `OMP_TARGET_OFFLOAD = [mandatory | disabled | default]`
 - Controls whether offload region runs on device or host

Warmup of GPUs

- It takes time to get a context on a device
 - Idle GPUs are in power saving mode
 - Just-in-time compilation (CUDA)
 - Transfer of kernel to GPU memory
 - Approx. 0.2 seconds (on our nodes)
- Warmup run
 - Required for accurate performance benchmarking in case of short runtimes
 - First offload that modifies the GPU context will initiate ‘warm up’ of the device
 - First transfer of data starts the device data environment

Exercises

- Do the first exercise
 - ex1_nvaccelinfo
 - Please note that nvcc v25.11 requires gcc version 15.x or older and CUDA 13.0 or older (so please do not load the newer versions even though they exist)
- Then start the second exercise
 - ex2_helloworld
 - Template Makefile available on DTU Learn
- Next lecture at 13.00 (Monday)!

Acknowledgements

- Some slides are from Jeff Larkin, NVIDIA's HPC Software team:
 - <https://developer.nvidia.com/blog/author/jlarkin/>
- Some slides are from Michael Klemm, OpenMP Architecture Review Board, AMD:
 - <https://www.openmp.org/about/our-team/officers-and-staff/>
- Some slides are from “OpenMP 4.5 target” by Tom Scogland and Oscar Hernandez:
 - ECP OpenMP tutorial 06-28-2017
- Some slides are from NVIDIA Developer
 - <https://developer.nvidia.com/>

End of lecture