

GPU Programming with OpenMP

Part 2: Data mapping



Hans Henrik Brandenborg Sørensen
DTU Computing Center

DTU Compute
<hhbs@dtu.dk>



$$f(x+\Delta x) = \sum_{i=0}^{\infty} \frac{(\Delta x)^i}{i!} f^{(i)}(x)$$

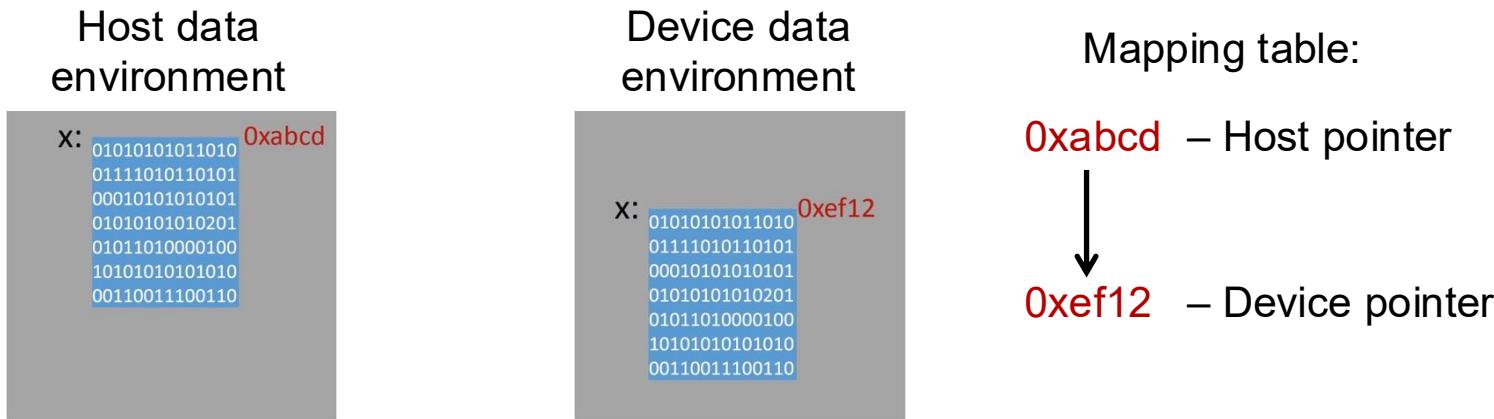
A colorful mathematical collage featuring various symbols: a purple theta, a yellow integral sign, a pink infinity symbol, a red sigma, a purple factorial exclamation mark, a red plus sign, a pink square root of 17, a red delta, and a pink e to the power of i pi. There are also red arrows and a red equals sign.

Overview

- OpenMP data mapping
 - map clause
 - target data
 - target enter / exit data
 - target update
 - declare target
- OpenMP array sections
- OpenMP runtime library
 - Device versions of malloc, free, memcpy
 - Matrix allocation on the device

OpenMP data environment

- OpenMP maintains a device data environment and a mapping table that records what memory pointers have been mapped from the host



- The table also maintains the translation between host memory pointers and device memory pointers
- Presence checks determine if data is already part of the device data environment before a transfer is made

OpenMP offload basics

■ target clauses related to data environment

- ❑ map
- ❑ is_device_ptr
- ❑ has_device_addr
- ❑ defaultmap
- ❑ ~~allocate~~
- ❑ ~~uses_allocator*~~

Not currently supported in nvcc++

OpenMP offload syntax

- Syntax C/C++:

```
map ( [ map-type-modifier [, ] map-type-
        modifier [, ] ... ] map-type : ] list)
```

- The **map** clause specifies how a list item is mapped from the host data environment to the corresponding list item in the device data environment

OpenMP offload syntax

- map-type can be
 - to
 - Allocates memory and moves data to the device
 - from
 - Allocates memory and moves data from the device
 - tofrom
 - Allocates memory and moves data to and from the device
 - alloc
 - Allocates memory on the device
 - release
 - The reference count is decreased by one
 - delete
 - Deletes data from the device (the reference count is set to 0)

OpenMP offload syntax

- map-type-modifier can be
 - always
 - Always copies the data to and from the device
 - ~~close~~
 - A hint to the runtime to allocate memory close to the target device
 - ~~mapper (mapper identifier)~~
 - Use a user-defined mapper
 - ~~present~~
 - This clause is ordered before other map clauses w/o present
 - ~~iterator (iterator definition)~~
 - Reference iterators defined by an iterators-definition

Not currently supported in nvc++

OpenMP offload syntax

■ Example

```
double a[N], b[N], c[N];  
/*  
   initialize arrays  
*/  
#pragma omp target teams loop map(to: a, b) map(from: c)  
for (int i = 0; i < N; i++)  
    c[i] = a[i] + b[i];
```

- ❑ Arrays a and b are transferred **to** the device before the offload
- ❑ Array c is **allocated** on the device before the offload
- ❑ Array c is transferred **from** the device after the offload

OpenMP offload syntax

■ Example – dynamic allocation

```
double *a, *b, *c;  
/*  
   allocate and initialize arrays of size N  
*/  
#pragma omp target teams loop \  
    map(to: a[0:N], b[0:N]) map(from: c[0:N])  
for (int i = 0; i < N; i++)  
    c[i] = a[i] + b[i];  
  
/* free arrays */
```

OpenMP array sections

- ❑ Arrays a and b are transferred **to** the device before the offload
- ❑ Array c is **allocated** on the device before the offload
- ❑ Array c is transferred **from** the device after the offload

OpenMP array sections

- An array section designates a subset of the elements in an array given by:

[lower-bound : length : stride]

- Must be a subset of the original array
 - Array sections are allowed on multidimensional arrays
- Must be integers or integer expressions
 - If lower-bound is left out it defaults to 0
 - length must evaluate to a non-negative integer and must be explicitly specified when the size of the array dimension is not known
 - stride must evaluate to a positive integer (default 1)

OpenMP array sections

■ Example – dynamic allocation

```
double *a, *b, *c;  
/*  
   allocate and initialize arrays of size N  
*/  
#pragma omp target teams loop \  
    map(to: a[0:N/2], b[0:N/2]) map(from: c[0:N/2])  
for (int i = 0; i < N/2; i++)  
    c[i] = a[i] + b[i];  
  
#pragma omp target teams loop \  
    map(to: a[N/2:N/2], b[N/2:N/2]) map(from: c[N/2:N/2])  
for (int i = N/2; i < N; i++)  
    c[i] = a[i] + b[i];
```

OpenMP offload syntax

■ Another example

```
double a[N], b[N], c[N], d[N];  
/*  
   initialize arrays  
*/  
#pragma omp target teams loop map(to: a, b) map(from: c)  
for (int i = 0; i < N; i++)  
    c[i] = a[i] + b[i];  
  
#pragma omp target teams loop map(to: a, b) map(from: d)  
for (int i = 0; i < N; i++)  
    d[i] = a[i] - b[i];
```

- ❑ Two offload regions that both use a and b (transferred 2 times)
- ❑ This works, but it is not the best we can do!

OpenMP offload syntax

■ Syntax C/C++:

```
#pragma omp target data [clause]
{
    ...
}
```

■ Clause can be

- if([target data :] scalar_expr)
- device(int_expr)
- map(...)
- use_device_ptr(ptr-list)
- use_device_addr(list)

■ Creates persistent data environment within { }

OpenMP offload syntax

■ So the better solution is

```
double a[N], b[N], c[N], d[N];
/*
    initialize arrays
*/
#pragma omp target data map(to: a, b) map(from: c, d)
{
    #pragma omp target teams loop map(to: a, b) map(from: c)
    for (int i = 0; i < N; i++)
        c[i] = a[i] + b[i];
    #pragma omp target teams loop map(to: a, b) map(from: d)
    for (int i = 0; i < N; i++)
        d[i] = a[i] - b[i];
}
```

Presence check: a and b already available!

- ❑ Arrays a and b are transferred once and used in both offloads regions

OpenMP measuring runtimes

■ Measuring the transfer time

```
double a[N], b[N], c[N];
/*
    initialize arrays and warmup device
*/
double t = omp_get_wtime();
#pragma omp target data map(to: a, b) map(from: c)
{
    double t = omp_get_wtime();
    #pragma omp target teams loop map(to: a, b) map(from: c)
    for (int i = 0; i < N; i++)
        c[i] = a[i] + b[i];
    printf("Runtime w/o transfer: %f\n", omp_get_wtime() - t);
}
printf("Runtime with transfer: %f\n", omp_get_wtime() - t);
```

```
$ ./vecadd
Runtime w/o transfer: 0.000565
Runtime with transfer: 0.002533
```

- ❑ For reliable runtimes run offload region several times and average

OpenMP offload syntax

■ Syntax C/C++:

```
#pragma omp target enter data [clause]  
#pragma omp target exit data [clause]
```

■ Clause can be

- if([target data :] scalar_expr)
- device(int_expr)
- map(...)
- depend([depend-modifier,] depen-type : list)
- nowait

■ Standalone directives that specifies mapping to the data environment of the default device

OpenMP offload syntax

■ Target enter / exit example

main.cpp

```
int main(int argc, char *argv[]) {  
    double a[N], b[N], c[N];  
  
    initialize(a, b, c);  
  
    #pragma omp target teams loop \  
        map(to: a, b) map(from: c)  
    for (int i = 0; i < N; i++)  
        c[i] = a[i] + b[i];  
  
    finalize(a, b, c);  
  
    return(0);  
}
```

functions.cpp

```
void initialize(double *a, double *b, double *c) {  
    for (int i = 0; i < N; i++)  
        a[i] = b[i] = i;  
    #pragma omp target enter data \  
        map(to: a[:N], b[:N]) map(alloc: c[:N])  
}  
  
void finalize(double *a, double *b, double *c) {  
    #pragma omp target exit data \  
        map(release: a[:N], b[:N]) map(from: c[:N])  
    for (int i = 0; i < N; i++)  
        printf("%f\n", c[i]);  
}
```

- Remember to release a and b in order to clean up data environment

OpenMP measuring runtime



■ H2D and D2H transfers

main.cpp

```
int main(int argc, char *argv[]) {
    double a[N], b[N], c[N];
    /*
        warm up device
    */
    initialize(a, b, c);

    #pragma omp target teams loop \
        map(to: a, b) map(from: c)
    for (int i = 0; i < N; i++)
        c[i] = a[i] + b[i];

    finalize(a, b, c);

    return(0);
}
```

functions.cpp

```
void initialize(double *a, double *b, double *c) {
    for (int i = 0; i < N; i++)
        a[i] = b[i] = i;
    double t = omp_get_wtime();
    #pragma omp target enter data \
        map(to: a[:N], b[:N]) map(alloca: c[:N])
    printf("H2D transfer time: %f\n",
           omp_get_wtime() - t);
}

void finalize(double *a, double *b, double *c) {
    double t = omp_get_wtime();
    #pragma omp target exit data \
        map(release: a[:N], b[:N]) map(from: c[:N])
    printf("D2H transfer time: %f\n",
           omp_get_wtime() - t); }
```

```
$ ./vecadd
H2D transfer time: 0.001761
D2H transfer time: 0.001139
```

- ❑ Remember to warm up the device or transfer times will be +0.2 secs

OpenMP offload syntax

■ Syntax C/C++:

```
#pragma omp target data update [clause]
```

■ Clause can be

- ❑ if([target data :] scalar_expr)
- ❑ device(int_expr)
- ❑ depend([depend-modifier,] depen-type : list)
- ❑ nowait
- ❑ to([motion-modifier[,] ...]:] list)
- ❑ from([motion-modifier[,] ...]:] list)

■ Standalone directive that makes the device data environment consistent with their original list items, according to the specified motion clauses

OpenMP offload syntax

■ Target update example

```
#pragma omp target data map(to: a, b) map(from: c, d)
{
    #pragma omp target teams loop map(to: a, b) map(from: c)
    for (int i = 0; i < N; i++)
        c[i] = a[i] + b[i];

    modify_b_on_host(b);
    #pragma omp target update to (b)

    #pragma omp target teams loop map(to: a, b) map(from: d)
    for (int i = 0; i < N; i++)
        d[i] = a[i] - b[i];
}
```

- ❑ Array b is modified in the host data environment and needs to be updated in the device data environment before the final offload

OpenMP offload syntax

■ Syntax C/C++:

```
#pragma omp declare target
```

declarations

```
#pragma omp end declare target
```

```
#pragma omp declare target [(list) | clause]
```

■ Clause can be

- to(list)
- ~~link(list)~~
- device_type(host | nohost | any)

■ This directive specifies that global variables and functions (C/C++) are mapped to a device for all device executions (or for a specific one via link)

OpenMP offload syntax

■ Using declare target

```
#pragma omp declare target
double a[N], b[N], c[N];

#pragma omp end declare target

void vecadd() {
    for (int i = 0; i < N; i++)
        c[i] = a[i] + b[i];
}

#pragma omp declare target (vecadd)

int main(int argc, char *argv[]) {
    /* initialize arrays on host */
    vecadd(); // Call host version
    #pragma omp target update to(a, b)
    #pragma omp target
    vecadd(); // Call device version
    #pragma omp target update from(c)
}
```

← no mapping – arrays are globally available

Implicit data mapping

■ Why did this run?

```
#define N 16

int main(int argc, char *argv[]) {
    double a[N], b[N], c[N];
    for (int i = 0; i < N; i++)
        a[i] = b[i] = i * 1.0;

#pragma omp target teams \
    distribute parallel for
for (int i = 0; i < N; i++)
    c[i] = a[i] + b[i];

for (int i = 0; i < N; i++)
    printf("%f\n", c[i]);
}
```

arrays are implicitly mapped to the device

arrays are implicitly mapped back to the host

Implicit data mapping

- If a variable is a scalar then it is implicitly `firstprivate` (i.e., not mapped)
- If a variable is not a scalar (i.e. an array or struct) then implicitly a `map (tofrom: . . .)` is added
- If a variable is a pointer its actual value has no meaning in the device memory and it “is treated as if it is the base pointer of a zero-length array section that had appeared in a `map clause`”
- However, if the `defaultmap` clause is present in the construct, it takes precedence

Summary of data mapping

- Map clause on a target construct
 - Maps variables for a single target region
 - Enclosed region executes on device and maps data
- Target data
 - Map variables across multiple target regions
 - Region does not execute on device, only maps data
- Declare target
 - Mapping variables and functions for the whole execution of the program (“globally mapped”)
- Target enter/exit/update data
 - Map variables in stand-alone clauses

OpenMP runtime library

OpenMP offload runtime library



■ Managing memory yourself

Name

```
void* omp_target_alloc(size_t,  
int dev_num)
```

```
void omp_target_free(void*,  
int dev_num)
```

```
int omp_target_memcpy(...,  
int dev_num)
```

```
int omp_target_memcpy_rect(...,  
int dev_num)
```

```
omp_target_associate_ptr(...)
```

```
omp_target_disassociate_ptr(...)
```

```
omp_target_is_present(...)
```

Functionality

allocate memory on device

free memory on device

memcpy to and from device

memcpy to and from device of
a rectangular subvolume

combining device ptr with
host ptr to be used in map
clause
(we use is_device_ptr clause)

OpenMP offload syntax

- Syntax C/C++:

```
void* omp_target_alloc(size_t size,  
                      int dev_num);
```

- This routine allocates memory of `size` bytes in the device data environment of device `dev_num` and returns a device pointer to that memory
- Syntax C/C++:

```
void omp_target_free(void *dev_ptr,  
                     int dev_num);
```

- This routine frees the memory at `dev_ptr` in the device data environment of device `dev_num`

OpenMP offload syntax

■ Syntax C/C++:

```
int omp_target_memcpy(void *dst,  
                      void *src,  
                      size_t length,  
                      size_t dst_offset,  
                      size_t src_offset,  
                      int dst_dev_num,  
                      int src_dev_num);
```

- This routine copies **length** bytes of memory at **offset src_offset** from **src** in the device data environment of device **src_dev_num** to **dst** starting at **offset dst_offset** in the device data environment of device **dst_dev_num**

OpenMP offload syntax

- The `is_device_ptr` clause indicates that its list items are device pointers
 - Inside the target construct, each list item is privatized and the new list item is initialized to the device address to which the original list item refers
- Support for device pointers created outside of OpenMP, specifically outside of any OpenMP mechanism that returns a device pointer, is implementation defined

CUDA device pointers are fully supported in nvc++

OpenMP offload syntax

■ Example – dynamic allocation

```
double *a, *b, *c, *a_d, *b_d, *c_d; // Notation _d is convention for device pointers.  
/* allocate and initialize arrays on host */  
  
int dev_num = omp_get_default_device();  
a_d = (double*)omp_target_alloc(N * sizeof(double), dev_num);  
b_d = (double*)omp_target_alloc(N * sizeof(double), dev_num);  
c_d = (double*)omp_target_alloc(N * sizeof(double), dev_num);  
  
omp_target_memcpy(a_d, a, N * sizeof(double), 0, 0, dev_num, omp_get_initial_device());  
omp_target_memcpy(b_d, b, N * sizeof(double), 0, 0, dev_num, omp_get_initial_device());  
  
#pragma omp target teams loop is_device_ptr(a_d, b_d, c_d)  
for (int i = 0; i < N; i++)  
    c_d[i] = a_d[i] + b_d[i];  
  
omp_target_memcpy(c, c_d, N * sizeof(double), 0, 0, omp_get_initial_device(), dev_num);  
...  
omp_target_free(a_d, dev_num);  
omp_target_free(b_d, dev_num);  
omp_target_free(c_d, dev_num);
```

clause `is_device_ptr` required here

OpenMP offload runtime library



■ Example matrix allocation – host version

```
/* Routine for allocating two-dimensional array */
double **malloc_2d(int m, int n) {
    if (m <= 0 || n <= 0)
        return NULL;

    double **A = (double**)malloc(m * sizeof(double *));
    if (A == NULL)
        return NULL;

    A[0] = (double*)malloc(m * n * sizeof(double));
    if (A[0] == NULL) {
        free(A);
        return NULL;
    }

    for (int i = 0; i < m; i++)
        A[i] = A[0] + i * n;

    return A;
}
```

```
void free_2d(double **A) {
    free(A[0]);
    free(A);
}
```

OpenMP offload runtime library



■ Example matrix allocation – device version

```
/* Routine for allocating two-dimensional array on the device */
double **malloc_2d_dev(int m, int n, double **data) {
    if (m <= 0 || n <= 0)
        return NULL;

    double **A = (double**)omp_target_alloc(m*sizeof(double *), omp_get_default_device());
    if (A == NULL)
        return NULL;

    double *a = (double*)omp_target_alloc(m*n * sizeof(double), omp_get_default_device());
    if (a == NULL) {
        omp_target_free(A, omp_get_default_device());
        return NULL;
    }

    #pragma omp target is_device_ptr(A, a)
    for (int i = 0; i < m; i++)
        A[i] = a + i * n;

    *data = a; return A;
}
```

```
void free_2d_dev(double **A,
                  double *data) {
    omp_target_free(data,
                    omp_get_default_device());
    omp_target_free(A,
                    omp_get_default_device());
}
```

OpenMP offload runtime library



■ Example matrix allocation

```
double **A = malloc_2d(N, N);                                // Allocate A on host

double *data;
double **A_d = malloc_2d_dev(N, N, &data);                  // Allocate A_d on device

/* initialize A on host */

omp_target_memcpy(data, A[0], N * N * sizeof(double), // Copy data from A to A_d
                  0, 0, omp_get_default_device(), omp_get_initial_device());

double sum = 0.0;
#pragma omp target teams distribute parallel for reduction(+:sum) is_device_ptr(A_d)
for (int i = 0; i < N; i++)
    for (int j = 0; j < N; j++)
        sum += A_d[i][j] * A_d[i][j];
printf("Frobenius norm: %f\n", sqrt(sum));

free_2d(A);                                              // Free A on host
free_2d_dev(A_d, data);                                  // Free A_d on device
```

Exercises

- Finish up the first two exercises
 - ex1_nvallocinfo
 - ex2_helloworld

- Begin the third exercise
 - ex3_mandelbrot

End of lecture