

Introduction to GPU performance tuning



Hans Henrik Brandenborg Sørensen

DTU Computing Center

DTU Compute

<hhbs@dtu.dk>

$$f(x+\Delta x) = \sum_{i=0}^{\infty} \frac{(\Delta x)^i}{i!} f^{(i)}(x)$$

A dense, abstract collage of mathematical symbols and numbers. It includes a large purple Greek letter Θ, a yellow integral sign with 'a' and 'b' as limits, a red infinity symbol, a blue summation symbol with a red exclamation mark, a green delta symbol, a red e^{i\pi} term, a blue square root of 17, a red plus sign, a blue Omega symbol, a red equals sign, a blue brace over the number 2.7182818284, a red chi-squared symbol, and various other mathematical characters like curly braces and arrows.



Overview

- Recap from week 1 (now with GPUs)
 - Performance metrics
 - Assessing your performance
- Speed-up GPU vs. CPU (and what is ‘fair’?)
- GPU performance tuning process
 - Transpose example
 - How many threads should I launch?
 - Latency hiding

GPU performance metrics

Performance tuning terminology

- **Execution time** [seconds]
 - Time to run the application (wall or cpu/gpu)
- **Performance** [Gflops]
 - How many floating point operations per second
- **Latency** [cycles or seconds]
 - Time from initiating a memory access or other action until the result is available
- **Bandwidth** [Gb/s]
 - The rate at which data can be transferred
- **Blocking** [blocksize]
 - Dividing matrices into tiles to fit memory hierarchy

You know
these from
week 1+2 of
this course!

Performance tuning terminology

- **Throughput** [#/s or Gb/s]
 - Sustained rate for instructions executed or data reads + writes achieved in practice
 - **Occupancy** [%]
 - Ratio of active threads to max possible in hardware
 - **Instruction level parallelism (ILP)** [#]
 - How many independent instructions can be executed (=pipelining)
 - **Thread level parallelism (TLP)** [#]
 - How many independent threads can be launched
 - **Coalescing**
 - 32 neighbor threads in team are reading data from a contiguous, aligned, region of global memory
- 
- Common GPU optimization terminology.

Which performance metric to use?

- Compute bound
 - Limited by # flops * time per flop

Gflops = # floating point operations / 10^9 / runtime

Which performance metric to use?

■ Compute bound

- Limited by # flops * time per flop

$$\text{Gflops} = \# \text{ floating point operations} / 10^9 / \text{runtime}$$

■ Memory bound

- Limited by # bytes moved / bandwidth

$$\text{Bandwidth} = (\text{Bytes_read} + \text{Bytes_written}) / 10^9 / \text{runtime}$$

How to assess your performance?



■ Compute bound

- Compare with the theoretical peak performance
 - Run nvaccelinfo to find specs and calculate, e.g.



SP: $114 \text{ SM} * 128 \text{ cores/SM} * 1.755 \text{ GHz} * 2 \text{ flops} = 51.2 \text{ Tflops}$

DP: $(1/2) * 51.2 \text{ Tflops} = 25.6 \text{ Tflops}$

How to assess your performance?



■ Compute bound

- Compare with the theoretical peak performance
 - Run nvaccelinfo to find specs and calculate, e.g.



SP: $114 \text{ SM} * 128 \text{ cores/SM} * 1.755 \text{ GHz} * 2 \text{ flops} = 51.2 \text{ Tflops}$

DP: $(1/2) * 51.2 \text{ Tflops} = 25.6 \text{ Tflops}$

■ Memory bound

- Compare with the theoretical peak bandwidth
 - Run nvaccelinfo to find specs and calculate, e.g.

Peak bandwidth = $1.593 \text{ GHz} * (5120 / 8) \text{ bytes} * 2 = 2039 \text{ GB/s}$

How to assess your performance?



■ Compute bound

- Compare with the theoretical peak performance
 - Run nvaccelinfo to find specs and calculate, e.g.

SP: 114 SM * 12

40-60%: okay

ops = 51.2 Tflops

DP:

60-75%: good

ops = 25.6 Tflops

■ Memory bound

>75%: excellent

- Compare with the theoretical peak bandwidth
 - Run nvaccelinfo to find specs and calculate, e.g.

$$\text{Peak bandwidth} = 1.593 \text{ GHz} * (5120 / 8) \text{ bytes} * 2 = 2039 \text{ GB/s}$$

How to assess CPU performance?



- Find the CPU specs online:

<https://www.amd.com/en/products/cpu/amd-epyc-9354>



How to assess CPU performance?



- Find the CPU specs online:

<https://www.amd.com/en/products/cpu/amd-epyc-9354>

- Compute bound

SP: $32 \text{ cores} * 3.75 \text{ GHz (AVX-512)} * 16 \text{ (AVX-512)} * 2 \text{ (FMA units)} * 2 \text{ flops per core} = 7.68 \text{ Tflops}$

DP: $(1/2) * 7.68 \text{ Tflops} = 3.84 \text{ Tflops}$

- Memory bound

Peak bandwidth = $4.8 \text{ GHz} * (64 / 8) \text{ bytes} * 12 = 460.8 \text{ GB/s}$

Speed-up

Speed-up (now with GPUs)

- GPU definition of speed-up is traditionally

$$\text{Speedup} = \frac{\text{CPUTime[s]}}{\text{GPUtime[s]}}$$

and usually written in times (×) manner, e.g., 3.2 ×

- Useful for indicating performance without telling what the performance actually is(!)

Speed-up (now with GPUs)

- GPU definition of speed-up is traditionally

$$\text{Speedup} = \frac{\text{CPUtime[s]}}{\text{GPUtime[s]}}$$

and usually written in times (×) manner, e.g., $3.2 \times$

- Useful for indicating performance without telling what the performance actually is(!)

But be fair when comparing to CPU times – speed-ups of $100 \times$ - $1000 \times$ (see http://www.nvidia.com/object/cuda_showcase_html.html) are unrealistic when the hardware specs are taken into account

Speed-up (now with GPUs)

- Expected speed-up on our nodes

- Compute bound: $25.6 / 3.84 = 6.7x$
 - Memory bound: $2039 / 460.8 = 4.4x$

Speed-up (now with GPUs)

- Expected speed-up on our nodes
 - Compute bound: $25.6 / 3.84 = 6.7x$
 - Memory bound: $2039 / 460.8 = 4.4x$
- Better speed-ups may be seen in practice
 - Execution model (SIMD / SIMT) – see next slide
 - Compiler成熟度 (vectorization is difficult)
 - ...but not by orders of magnitude(!)

Executing of the GPU

- Execution model maps code to instructions
 - SIMD – Single Instruction, Multiple Data
 - SIMT – Single Instruction, Multiple Threads
- CPU (SIMD mapping)
 - The compiler takes care of mapping code to the most efficient instructions **at compile time**
- GPU (SIMT mapping)
 - Kernels look like serial functions for a **SINGLE** thread
 - CUDA will automatically launch on **MANY** threads
 - Code is mapped to instructions **at runtime!**

Example of speed-up calculation



$$\pi = \int_0^1 \frac{4}{1+x^2} dx \approx \frac{1}{N} \sum_{i=1}^N \frac{4}{1 + \left(\frac{i-0.5}{N}\right)^2}.$$

```
double pi(long N)
{
    double sum = 0.0;
    double h = 1.0/N;
    #pragma omp parallel for \
        reduction(+: sum)
    for(long i=1; i<=N; i++) {
        double x = h*(i-0.5);
        sum += 4.0/(1.0+x*x);
    }
    return h*sum;
}
```

Example of speed-up calculation



$$\pi = \int_0^1 \frac{4}{1+x^2} dx \approx \frac{1}{N} \sum_{i=1}^N \frac{4}{1 + \left(\frac{i-0.5}{N}\right)^2}.$$



```
double pi(long N)
{
    double sum = 0.0;
    double h = 1.0/N;
    #pragma omp parallel for \
        reduction(+: sum)
    for(long i=1; i<=N; i++) {
        double x = h*(i-0.5);
        sum += 4.0/(1.0+x*x);
    }
    return h*sum;
}
```

```
icx -O3 -fopenmp pi_cpu.cpp
```

OMP_NUM_THREADS	-O3 (s)
1	99.5
2	50.2
4	25.6
8	12.9
16	6.5
32	3.3
64	3.1

Example of speed-up calculation



$$\pi = \int_0^1 \frac{4}{1+x^2} dx \approx \frac{1}{N} \sum_{i=1}^N \frac{4}{1 + \left(\frac{i-0.5}{N}\right)^2}.$$



```
double pi(long N)
{
    double sum = 0.0;
    double h = 1.0/N;
    #pragma omp parallel for \
        reduction(+: sum)
    for(long i=1; i<=N; i++) {
        double x = h*(i-0.5);
        sum += 4.0/(1.0+x*x);
    }
    return h*sum;
}
```

```
icx -O3 -fopenmp pi_cpu.cpp
icx -xcore-avx512 -O3 ...
```

OMP_NUM_THREADS	-O3 (s)	avx (s)
1	99.5	30.3
2	50.2	15.1
4	25.6	7.6
8	12.9	3.8
16	6.5	1.9
32	3.3	1.0
64	3.1	0.9

Example of speed-up calculation



$$\pi = \int_0^1 \frac{4}{1+x^2} dx \approx \frac{1}{N} \sum_{i=1}^N \frac{4}{1 + \left(\frac{i-0.5}{N}\right)^2}.$$



```
double pi(long N)
remark #15305: vectorization support: vector length 2
remark #15399: vectorization support: unroll factor set to 4
remark #15300: LOOP WAS VECTORIZED
remark #15486: divides: 1

#pragma omp parallel for
remark #15305: vectorization support: vector length 8
remark #15399: vectorization support: unroll factor set to 4
remark #15300: LOOP WAS VECTORIZED
remark #15486: divides: 1

    sum += 4.0 / (1.0+x*x) ;
}
return h*sum;
}
```

```
icx -O3 -fopenmp pi_cpu.cpp
icx -xcore-avx512 -O3 ...
```

OMP_NUM_THREADS	-O3 (s)	avx (s)
1	99.5	30.3
2	50.2	15.1
4	25.6	7.6
8	12.9	3.8
16	6.5	1.9
32	3.3	1.0
64	3.1	0.9

Example of speed-up calculation



$$\pi = \int_0^1 \frac{4}{1+x^2} dx \approx \frac{1}{N} \sum_{i=1}^N \frac{4}{1 + \left(\frac{i-0.5}{N}\right)^2}.$$

```
double pi(long N)
{
    double sum = 0.0;
    double h = 1.0/N;
    #pragma omp target teams \
        num_teams(16384) \
        thread_limit(128) \
        distribute parallel for \
        reduction(+: sum)
    for(long i=1; i<=N; i++) {
        double x = h*(i-0.5);
        sum += 4.0/(1.0+x*x);
    }
    return h*sum;
}
```

Example of speed-up calculation

$$\pi = \int_0^1 \frac{4}{1+x^2} dx \approx \frac{1}{N} \sum_{i=1}^N \frac{4}{1 + \left(\frac{i-0.5}{N}\right)^2}.$$



```
double pi(long N)
{
    double sum = 0.0;
    double h = 1.0/N;
    #pragma omp target teams \
        num_teams(16384) \
        thread_limit(128) \
        distribute parallel for \
        reduction(+: sum)
    for(long i=1; i<=N; i++) {
        double x = h*(i-0.5);
        sum += 4.0/(1.0+x*x);
    }
    return h*sum;
}
```

```
nvc++ -fast -mp=gpu pi_gpu.cpp
```

num_teams	thread_limit	nvc (s)
16384	1	6.17
16384	2	3.08
16384	4	1.54
16384	8	0.77
16384	32	0.19
16384	64	0.18
16384	128	0.18

Example of speed-up calculation

$$\pi = \int_0^1 \frac{4}{1+x^2} dx \approx \frac{1}{N} \sum_{i=1}^N \frac{4}{1 + \left(\frac{i-0.5}{N}\right)^2}.$$



1 CPU / # threads = # cores / avx512

Speed-up = 1.0 seconds / 0.18 seconds = 5.6 ×

1 GPU / best teams parallel configuration / deduct warm up



GPU performance tuning process

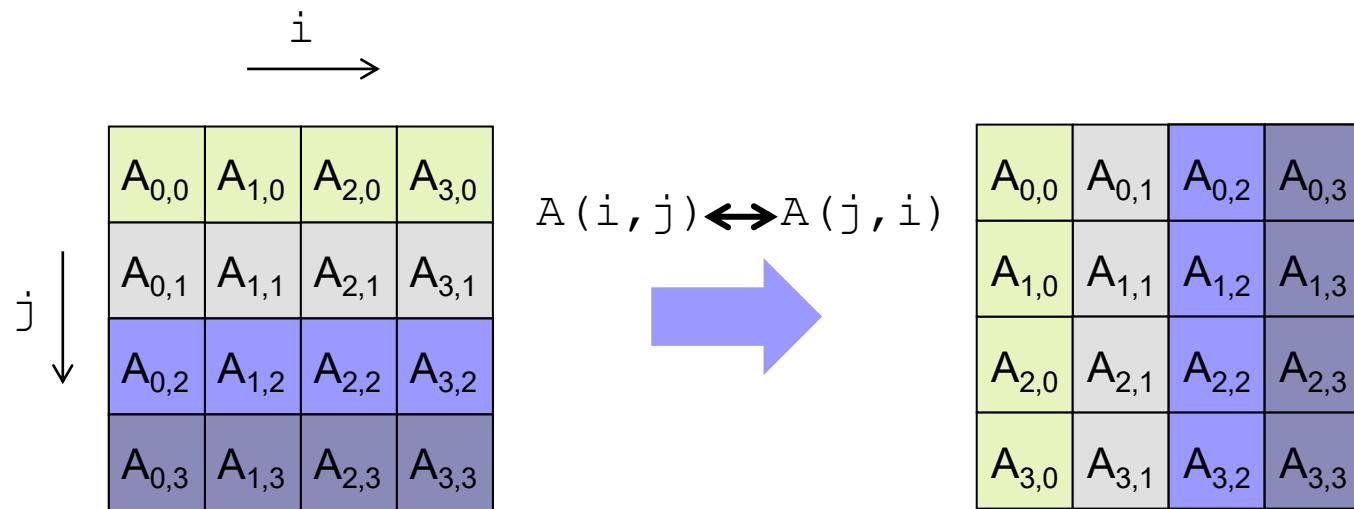
GPU performance tuning process

- Determine what limits the offload performance
 - Parallelism (concurrency bound)
 - Memory accesses (memory bound) ... or a combination
 - Floating point operations (compute bound)
- Use appropriate performance metric for the offload
 - Or use speed-up between kernel modifications
- Address the limiters in the order of importance
 - Determine how close you are to the theoretical peaks
 - Analyze
 - Apply optimizations ... and iterate with small steps

Transpose example

Row major storage

A _{0,0}	A _{1,0}	A _{2,0}	A _{3,0}	A _{0,1}	A _{1,1}	A _{2,1}	A _{3,1}	A _{0,2}	A _{1,2}	A _{2,2}	A _{3,2}	A _{0,3}	A _{1,3}	A _{2,3}	A _{3,3}
------------------	------------------	------------------	------------------	------------------	------------------	------------------	------------------	------------------	------------------	------------------	------------------	------------------	------------------	------------------	------------------



- We will use this example to illustrate the process of performance tuning an OpenMP offload code “step-by-step”

Transpose example (v1 seq)

```
// Reference sequential CPU transpose
void transpose(double **A, double **At)
{
    for(int i=0; i < N; i++)
        for(int j=0; j < N; j++)
            At[i][j] = A[j][i];
}
```

```
// Baseline sequential GPU transpose
#pragma omp declare target(transpose)

void transpose_seq(double **A, double **At)
{
    #pragma omp target
    transpose(A, At);
}
```

Transpose example (v1 seq)

```
#define N 7296
...
// CPU reference transpose for checking result
transpose(A, At_CPU);

#pragma omp target enter data \
    map(to: A[0:N][0:N]) map(alloca: At[0:N][0:N])

// GPU sequential version
transpose_seq(A, At);

#pragma omp target exit data \
    map(release: A[0:N][0:N]) map(from: At[0:N][0:N])

check(At, At_CPU); // Check result
...
```

Transpose example (v1 seq)

```
$ nvcc -fast -Msafeptr -Minfo -mp=gpu -gpu=cc90 -acc -c -o transpose.o transpose.cpp
...
transpose_seq(double**, double**, long):
    17, Loop not vectorized/parallelized: not countable
27, #omp target
    27, Generating "nvkernel__Z13transpose_seqPPdS0_1_F1L27_3" GPU kernel
...
& nsys profile --trace=cuda --stats=true ./transpose
...
[4/6] Executing 'gpukernsum' stats report
```

Time (%)	Total Time (ns)	Instances	Avg (ns)	...	Name
100.0	12,305,926,981	1	12,305,926,981.0		nvkernel__Z13transpose_seqPPdS0_1_F1L27_3
0.0	144,226	2	72,113.0		nvkernel__Z13malloc_2d_deviiPPd_F1L45_2

nsys command line
profiling tool

Version	v1 seq
Time [ms]	12305

Transpose example (v2 per col)



```
// OpenMP offload transpose using one thread per col of A
void transpose_per_col(double **A, double **At)
{
    #pragma omp target teams distribute parallel for \
        num_teams(114) thread_limit(64)
    for(int i = 0; i < N; i++)
        for(int j = 0; j < N; j++)
            At[i][j] = A[j][i];
}
```

Transpose example (v2 per col)



```
$ nvc++ -fast -Msafeptr -Minfo -mp=gpu -gpu=cc90 -acc -c -o transpose.o transpose.cpp
...
transpose_per_col(double **, double **):
    52, #omp target teams distribute parallel for num_teams(114) thread_limit(64)
    52, Generating "nvkernel_Z17transpose_per_colPPds0_1_F1L52_13" GPU kernel
    58, Loop parallelized across teams and threads(128), schedule(static)
    58, Loop not vectorized/parallelized: not countable
    59, Loop not vectorized: unprofitable for target
        Loop unrolled 4 times

& nsys profile --trace=cuda --stats=true ./transpose
...
[4/6] Executing 'gpukernsum' stats report
```

Time (%)	Total Time (ns)	Instances	Avg (ns)	...	Name
100.0	419,244,419	100	4,192,444.2		nvkernel_Z17transpose_per_colPPds0_1_F1L52_13
0.0	144,226	2	72,113.0		nvkernel_Z13malloc_2d_deviiPPd_F1L45_2

Version	v1 seq	v2 per col
Time [ms]	12305	4.19

Transpose example (v3 per elm)



```
// OpenMP offload transpose using one thread per element
void transpose_per_elm(double **A, double **At)
{
    #pragma omp target teams distribute parallel for \
        collapse(2) num_teams(N*N/64) thread_limit(64)
    for(int i = 0; i < N; i++)
        for(int j = 0; j < N; j++)
            At[i][j] = A[j][i];
}
```

Transpose example (v3 per elm)



```
$ nvc++ -fast -Msafeptr -Minfo -mp=gpu -gpu=cc90 -acc -c -o transpose.o transpose.cpp
...
transpose_per_elm(double **, double **):
    63, #omp target teams distribute parallel for num_teams(831744) thread_limit(64)
    63, Generating "nvkernel__Z17transpose_per_elmPPdS0_1_F1L63_17" GPU kernel
    72, Loop parallelized across teams and threads(128), schedule(static)
    73, Loop not vectorized/parallelized: not countable

& nsys profile --trace=cuda --stats=true ./transpose
...
[4/6] Executing 'gpukernsum' stats report
```

Time (%)	Total Time (ns)	Instances	Avg (ns)	..	Name
99.8	78,072,972	100	780,729.7	..	nvkernel__Z17transpose_per_elmPPdS0_1_F1L63_17
0.2	152,064	2	76,032.0	..	nvkernel__Z13malloc_2d_deviiPPd_F1L45_2

Version	v1 seq	v2 per col	v3 per elm
Time [ms]	12305	4.19	0.78

How many threads should I run?



- Why is one thread per core not enough?
 - E.g. `num_teams(114) thread_limit(64)`

How many threads should I run?

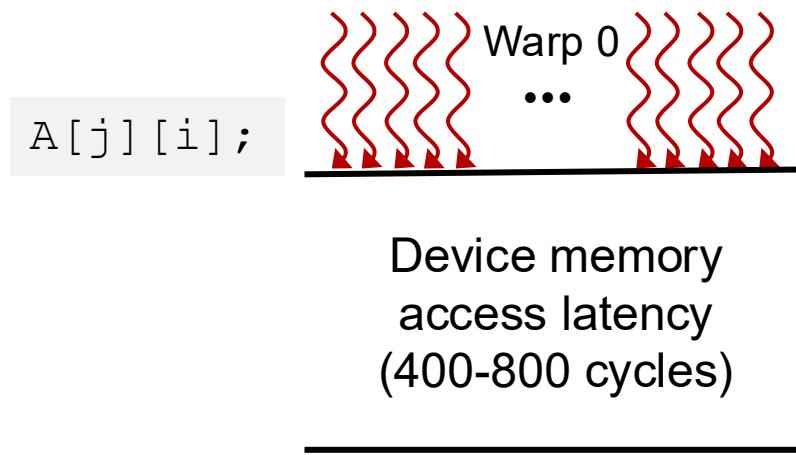
- Why is one thread per core not enough?
 - E.g. `num_teams(114) thread_limit(64)`



How many threads should I run?

- Why is one thread per core not enough?

- E.g. `num_teams(114) thread_limit(64)`



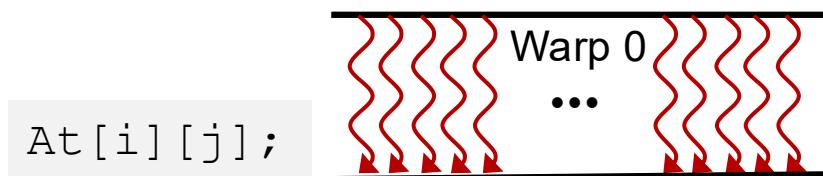
How many threads should I run?

- Why is one thread per core not enough?

- E.g. `num_teams(114) thread_limit(64)`



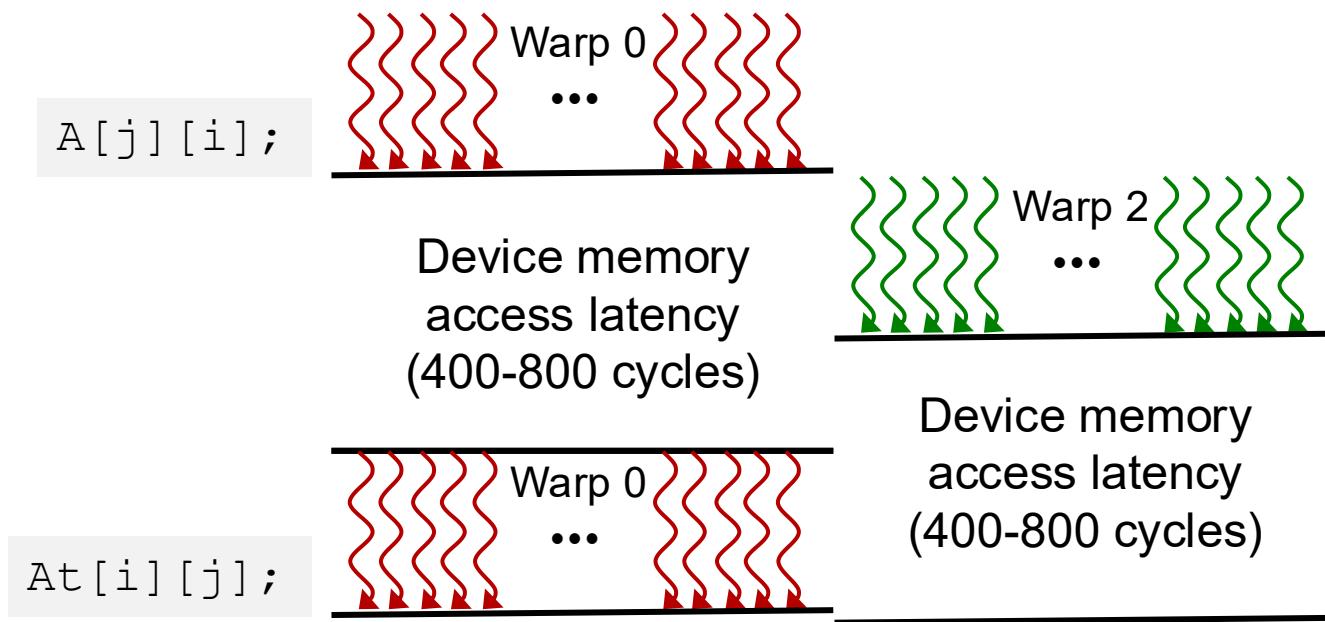
Device memory
access latency
(400-800 cycles)



How many threads should I run?

- Why is one thread per core not enough?

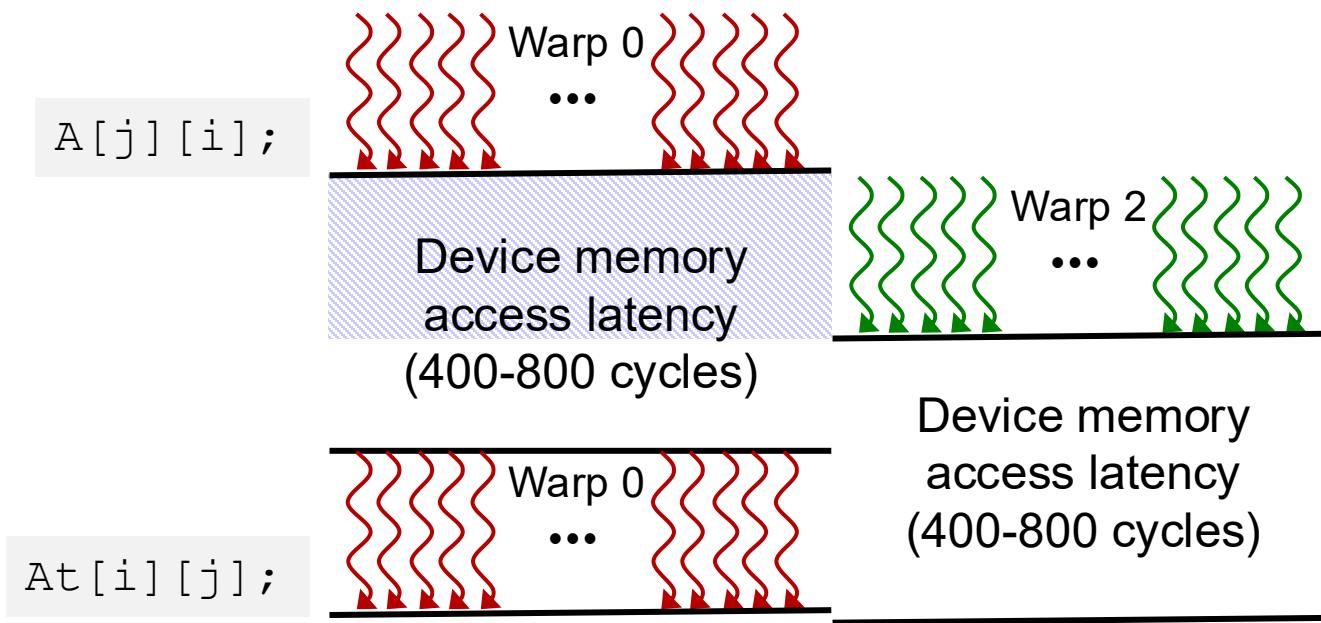
- E.g. `num_teams(114) thread_limit(64)`



How many threads should I run?

- Why is one thread per core not enough?

- E.g. `num_teams(114) thread_limit(64)`



- Reason: We can hide memory latency by having idle warps to schedule while waiting for data

Transpose example (v4 warp)

```
// OpenMP offload transpose using one warp per col
void transpose_warp_per_col(double **A, double **At)
{
    #pragma omp target teams loop \
        num_teams(N) thread_limit(32)
    for(int i = 0; i < N; i++) {
        #pragma omp loop bind(parallel)
        for(int j = 0; j < N; j++)
            At[i][j] = A[j][i];
    }
}
```

Transpose example (v4 warp)

```
$ nvc++ -fast -Msafeptr -Minfo -mp=gpu -gpu=cc90 -acc -c -o transpose.o transpose.cpp
...
transpose_warp_per_col(double **, double **, long):
    80, #omp target teams loop num_teams(7296) thread_limit(32)
    80, Generating "nvkernel_Z22transpose_warp_per_colPPdS0_1_F1L80_19" GPU kernel
        Generating NVIDIA GPU code
    86, Loop parallelized across teams(7296) /* blockIdx.x */
    88, Loop parallelized across threads(32) /* threadIdx.x */
    80, Generating Multicore code
    86, Loop parallelized across threads
88, Loop is parallelizable
Loop not vectorized: unprofitable for target
Loop unrolled 4 times
```

```
& nsys profile --trace=cuda --stats=true ./transpose
```

```
...
```

```
[4/6] Executing 'gpukernsum' stats report
```

Time (%)	Total Time (ns)	Instances	Avg (ns)	..	Name
99.8	67,454,944	100	674,549.4		nvkernel_Z22transpose_warp_per_rowPPdS0_1_F1L78_21
0.2	152,096	2	76,048.0.		nvkernel_Z13malloc_2d_deviiPPd_F1L45_2

Version	v1 seq	v2 per col	v3 per elm	v4 warp
Time [ms]	12305	4.19	0.78	0.67

End of lecture